



# Open vSwitch

**From: OpenStack: OVS Deep Dive**

*Justin Pettit Eric Lopez 07 November 2013*

# Features

- Open vSwitch supports the following features:
- Visibility into inter-VM communication via NetFlow, sFlow(R), IPFIX, SPAN, RSPAN, and GRE-tunneled mirrors
- LACP (IEEE 802.1AX-2008)
- Standard 802.1Q VLAN model with trunking
- Multicast snooping
- IETF Auto-Attach SPBM and rudimentary required LLDP support
- BFD and 802.1ag link monitoring
- STP (IEEE 802.1D-1998) and RSTP (IEEE 802.1D-2004)
- Fine-grained QoS control
- Support for HFSC qdisc
- Per VM interface traffic policing
- NIC bonding with source-MAC load balancing, active backup, and L4 hashing
- OpenFlow protocol support (including many extensions for virtualization)
- IPv6 support
- Multiple tunneling protocols (GRE, VXLAN, STT, and Geneve, with IPsec support)
- Remote configuration protocol with C and Python bindings
- Kernel and user-space forwarding engine options
- Multi-table forwarding pipeline with flow-caching engine
- Forwarding layer abstraction to ease porting to new software and hardware platforms

# Status: check with Claude (4.5)

- Open vSwitch (OVS) remains widely used and relevant, particularly in:
  - **OpenStack deployments** - still the default virtual switch
  - **Kubernetes networking** - used by CNI plugins like OVN-Kubernetes, Antrea
  - **NFV/SDN environments** - for programmable networking
- **Current landscape includes:**
- **eBPF-based alternatives gaining traction:**
  - **Cilium** - increasingly popular for Kubernetes, offers better performance
  - **Calico** (eBPF mode) - competing in container networking
- **Cloud provider solutions:**
  - AWS VPC networking (proprietary)
  - Azure VFP (Virtual Filtering Platform)
  - GCP Andromeda
- **Hardware offload:**
  - DPDK-accelerated OVS
  - SmartNICs with OVS offload
- **Status:** OVS isn't being replaced wholesale, but:
  - For new Kubernetes deployments, many choose Cilium or Calico for better observability and performance
  - Legacy virtualization and OpenStack still heavily use OVS
  - OVN (Open Virtual Network) extends OVS with L3 capabilities and remains active
  - The shift is toward eBPF-based solutions for containers while OVS maintains its position in traditional virtualization and specific SDN use cases.

# Status: check with Meta AI (llama-4)

- Open vSwitch (OVS) remains a popular and widely-used open-source virtual switch in the world of networking and cloud networking. Its status can be summarized as follows:
  - **Still widely used and maintained:** OVS is actively maintained by a large community of developers and is widely used in various cloud and networking projects, such as OpenStack, Kubernetes, and NFV (Network Function Virtualization) environments.
- **Key features and strengths:**
  - **High-performance switching:** OVS provides high-performance switching capabilities, with support for various hardware acceleration technologies like DPDK (Data Plane Development Kit).
  - **Flexibility and programmability:** OVS is highly configurable and programmable, with support for OpenFlow, OVSDB (Open vSwitch Database), and other protocols.
  - **Virtualization and containerization:** OVS is widely used in virtualized and containerized environments, providing connectivity and networking services for VMs and containers.
- **Other popular alternatives:**
  - While OVS remains a popular choice, other alternatives have gained traction in recent years:
  - **FD.io (Vector Packet Processing):** FD.io is an open-source project providing a high-performance, scalable, and modular packet processing framework. It's used in various projects, including Cisco's VPP (Vector Packet Processing).
  - **Contiv-VPP:** Contiv-VPP is a Kubernetes networking plugin that leverages VPP (Vector Packet Processing) technology to provide high-performance networking for containerized workloads.
  - **Calico:** Calico is a popular open-source networking project that provides a scalable and secure networking solution for cloud-native applications. While not a traditional virtual switch, Calico provides a networking fabric that can be used in conjunction with OVS or other switching technologies.
  - **Cilium:** Cilium is another popular open-source project that provides a networking and security solution for containerized workloads. It uses eBPF (Extended Berkeley Packet Filter) technology to provide high-performance networking and security features.

## Cloud and networking trends:

The networking landscape is evolving, with trends like cloud-native networking, NFV, and SDN (Software-Defined Networking) driving the development of new networking technologies. OVS remains a key player in this ecosystem, but other technologies like those mentioned above are gaining traction.

In summary, Open vSwitch is still a widely-used and maintained open-source virtual switch, but other alternatives are gaining popularity, driven by the evolving needs of cloud-native networking, NFV, and SDN.

## (Partial) List of Contributors

---




vmware®



## Main Components

---



# OVS Components

- **Open vSwitch Database (OVSDB)** contains switch configuration and keeps track of created and modified interfaces. All the configuration is stored on persistent storage and survives a reboot. The OVSDB-server communicates with ovs-vswitchd and the controller using the OVSDB management protocol.
- **Open vSwitch Daemon:** The Open vSwitch daemon (ovs-vswitchd) is one of the major components of Open vSwitch. It communicates with the controller using OpenFlow and with the OVSDB-server through the OVSDB management protocol. The ovs-vswitchd communicates with the kernel module over netlink (a Linux kernel interface used for creating a connection between userspace processes and the kernel). It supports multiple independent data paths, known as bridges.

---

# **CONFIGURATION DATABASE**

## **ovsdb-server**

---

### **Database that holds switch-level configuration**

- Bridge, interface, tunnel definitions
- OVSDB and OpenFlow controller addresses

**Configuration is stored on disk and survives a reboot** Custom database with nice properties:

- Value constraints
- Weak references
- Garbage collection


**Log-based**

**Speaks OVSDB protocol to manager and ovs-vswitchd**

**Tools: [ovs-vsctl](#), [ovsdb-tool](#), [ovsdb-client](#), [ovs-appctl](#)**

## Core Tables

---



“Open\_vSwitch” is the root table and there is always only a single row. The tables here are the ones most commonly used; a full entity- relationship diagram is available in the ovs-vswitchd.conf.db man page.

## Debugging the Database

---

**ovs-vsctl: Configures ovs-vswitchd, but really a high-level interface for database**

ovs-vsctl add-br <bridge>

ovs-vsctl list-br ovs-vsctl add-port <bridge> <port>

ovs-vsctl list-ports <bridge>

ovs-vsctl get-manager <bridge>

ovs-vsctl get-controller <bridge>

ovs-vsctl list <table>

**ovsdb-tool: Command-line tool for managing database file**

- ovsdb-tool show-log[-mmm]<file>

## ovsdb-tool show-log

Record number	Time of Change	Caller's comment
root@vm-vswitch:~# ovsdb-tool show-log -m! ...!	record 3: 2011-04-13 16:03:52	"ovs-vsctl: /usr/bin/ovs-vsctl --timeout=20 -- --with-iface --if-exists del-port eth0 -- --may-exist add-br xenbr0 -- -- may-exist add-port xenbr0 eth0 -- set Bridge xenbr0 "other-config:hwaddr=\\"00:0c:29:ab:f1:e9\\" -- set Bridge xenbr0 fail_mode=standalone -- remove Bridge xenbr0 other_config disable-in-band -- br-set-external-id xenbr0 xs-network-uuids 9ae8bc91-cfb8-b873-1947-b9c4098e4f4b"! !table Port insert row "xenbr0":! !table Port insert row "eth0":! !table Interface insert row "eth0":! !table Interface insert row "xenbr0":! !table Open_vSwitch row a1863ada:! !table Bridge insert row "xenbr0":! ...!"

Database changes

---


# **FORWARDING PATH**

# Linux Bridge Design

**Simple forwarding**

**Matches destination MAC address and forwards**


**Packet never leaves kernel**



## Open vSwitch Design


---

- Decision about how to process packet made in userspace
- First packet of new flow goes to ovs-vswitchd, following packets hit cached entry in kernel



Flow entry	match field	counter	Action (Instruction)	priority	Timeout
1					
n	...	...	...		

# Inside OVS



## ovs-vswitchd

---

### Core component in the system:

- Communicates with outside world using OpenFlow
- Communicates with ovsdb-server using OVSDB protocol
- Communicates with kernel module over netlink
- Communicates with the system through netdev abstract interface

- Supports multiple independent datapaths (bridges)
- Packet classifier supports efficient flow lookup with wildcards and “explodes” these (possibly) wildcard rules for fast processing by the datapath
  - Implements mirroring, bonding, and VLANs through modifications of the same flow table exposed through OpenFlow
  - Checks datapath flow counters to handle flow expiration and stats requests

Tools: [ovs-ofctl](#), [ovs-appctl](#)

# OVS Kernel Module

- **Kernel module that handles switching and tunneling**
- **Fast cache of non-overlapping flows**
- **Designed to be fast and simple**
  - Packet comes in, if found, associated actions executed and counters updated. Otherwise, sent to userspace
  - Does no flow expiration
  - Knows nothing of OpenFlow
- **Implements tunnels**

Tools: `ovs-dpctl`

# Userspace Processing

- Packet received from kernel
- Given to the classifier to look for matching flows accumulates
- If “normal” action included, accumulates actions from “normal” processing, such as L2 forwarding and bonding
- Actions accumulated from configured modules, such as mirroring Prior, an exact match flow is generated with the accumulated actions and pushed down to the kernel module (along with the packet)

# Kernel Processing

- Packet arrives and header fields extracted Header fields are hashed and used as an index into a set of large hash tables
- If entry found, actions applied to packet and counters are updated
- If entry is not found, packet sent to userspace and miss counter incremented

# Megaflows

- Version 1.11 added support for wildcarding in the datapath
- ovs-vswitchd dynamically determines how much wildcarding can do
  - Flow table
  - Actions from matching flow
  - General switch configuration (e.g. bonding)
- With megaflows, “normal” performance close to Linux bridge

# Tunnels

- Tunnels in OVS are just virtual ports with own OpenFlow port number
- Keys set statically at creation time or dynamically through OpenFlow action
- Types:
  - GRE
  - VxLAN
  - LISP
- Visible in Kernel datapath:
  - `ovs-dpctl show`

---

# **UTILITIES**

## **ovs-ofctl speaks to OpenFlow module**

```
ovs-ofctl show<bridge>
ovs-ofctl dump-flows <bridge>
ovs-ofctl add-flow <bridge> <flow>
ovs-ofctl del-flows <bridge> [flow]
ovs-ofctl snoop <bridge>
```

## **OpenFlow plus extensions**

Resubmit Action: Simulate multiple tables in a single table

NXM: Extensible match

Registers: Eight 32-bit metadata registers

Fine-grained control over multiple controllers

## **See “hidden” flows (in-band, fail-open, etc):**

```
ovs-appctl bridge/dump-flows<bridge>
```

## ovs-ofctl show <br>

---

```
root@vm-vswitch:~# ovs-ofctl show br0! OFPT_FEATURES_REPLY (xid=0x2):
dpid:0000505400000005! n_tables:254, n_buffers:256! capabilities: FLOW_STATS
TABLE_STATS PORT_STATS QUEUE_STATS ARP_MATCH_IP!
actions: OUTPUT SET_VLAN_VID SET_VLAN_PCP STRIP_VLAN SET_DL_SRC SET_DL_DST SET_NW_SRC
SET_NW_DST SET_NW_TOS SET_TP_SRC SET_TP_DST ENQUEUE!
1 (eth0): addr:50:54:00:00:00:05!

    config:      0!
    state:       0!
    current:     1GB-FD COPPER AUTO_NEG!
    advertised:  10MB-HD 10MB-FD 100MB-HD 100MB-FD 1GB-FD COPPER AUTO_NEG!
    supported:   10MB-HD 10MB-FD 100MB-HD 100MB-FD 1GB-FD COPPER AUTO_NEG!
    speed:       1000 Mbps now, 1000 Mbps max!
2 (eth1): addr:50:54:00:00:00:06!
    config:      0!
    state:       0!
    current:     1GB-FD COPPER AUTO_NEG!
    advertised:  10MB-HD 10MB-FD 100MB-HD 100MB-FD 1GB-FD COPPER AUTO_NEG!
    supported:   10MB-HD 10MB-FD 100MB-HD 100MB-FD 1GB-FD COPPER AUTO_NEG!
    speed:       1000 Mbps now, 1000 Mbps max!
LOCAL(br0): addr:50:54:00:00:00:05!
    config:      0! 0!
    state:
    speed:       0 Mbps now, 0 Mbps max!
OFPT_GET_CONFIG_REPLY (xid=0x4): frags=normal miss_send_len=0!
```

OpenFlow  
port  
number

Interface  
name

## **ovs-ofctl dump-flows <br>**

---

**The default flow table includes a single entry that does “normal” processing:**

```
root@vm-vswitch:~# ovs-ofctl dump-flows br0! NXST_FLOW reply  
(xid=0x4) :! cookie=0x0, duration=4.05s, table=0, n_packets=8,  
n_bytes=784, idle_age=0, priority=0 actions=NORMAL! !
```

## Hidden Flows

---

**There are flows that OVS uses for its own purpose that are higher priority than can be configured from outside**

### Types

In-bandcontrol (priority $\geq$ 180000): Allow controller traffic to pass regardless of configured flows

Fail-open (priority = 0xf0f0f0): Allow all traffic to pass when a connection to the controller fails

**They are hidden from controllers and commands like “ovs-ofctl dump-flows” due to being higher priority than OpenFlow allows (>65535)**

**Only visible with “ovs-appctl bridge/dump-flows <bridge>”**

## **Kernel Datapath**

---

**ovs-dpctl speaks to kernel module See datapaths and their attached interfaces:**

`ovs-dpctlshow`

**See flows cached in datapath:**

`ovs-dpctldump-flows`

## ovs-dpctl show

---

hit: Packets hit existing entry

missed: Packets sent to userspace

lost: Dropped before getting to userspace

```
root@vm-vswitch:~# ovs-dpctl show!
system@ovs-system:~!
```

lookups: hit:188056 missed:7722 lost:0!

flows: 2!

masks: hit:199268 total:1 hit/pkt:1.02!

port 0: ovs-system (internal)!

port 1: br0 (internal)!

port 2: eth0!

port 3: eth1!

Interface name

Datapath port number

Interface type

## ovs-dpctl dump-flows



### Flows used to be exact-match:

```
in_port(2),eth(src=50:54:00:00:00:01, dst=50:54:00:00:00:03), eth_type(0x0800), ipv4(src=192.168.0.1, dst=192.168.0.2, proto=1, tos=0, ttl=64, frag=no), icmp(type=8, code=0), packets:3, bytes:294, used:0.185s, actions:3
```

```
in_port(3),eth(src=50:54:00:00:00:03, dst=50:54:00:00:00:01), eth_type(0x0800), ipv4(src=192.168.0.2, dst=192.168.0.1, proto=1, tos=0, ttl=64, frag=no), icmp(type=0, code=0), packets:3, bytes:294, used:0.205s, actions:2
```

### Starting in OVS 1.11, may contain wildcards:

```
in_port(3),eth(src=50:54:00:00:00:03, dst=50:54:00:00:00:01), eth_type(0x0800), ipv4(src=192.168.0.2//0.0.0.0, dst=192.168.0.1/0.0.0.0, proto=1/0, tos=0/0, ttl=64/0, frag=no/0x2), icmp(type=0/0, code=0/0), packets:95, bytes:9310, used:0.425s, actions:2
```

```
in_port(2),eth(src=50:54:00:00:00:01, dst=50:54:00:00:00:03), eth_type(0x0800), ipv4(src=192.168.0.1//0.0.0.0, dst=192.168.0.2/0.0.0.0, proto=1/0, tos=0/0, ttl=64/0, frag=no/0x2), icmp(type=8/0, code=0/0), packets:95, bytes:9310, used:0.525s, actions:3
```

## **ovs-appctl**

---

**Utility to invoke runtime control and query facilities in most OVS daemons**

**The “-t <target>” option specifies the daemon name (default is ovs-vswitchd)**

**All daemons support the following commands:**

help – Lists the commands supported by the target

version – Displays the version and compilation date of the target

vlog/list – List the known logging modules and their current levels

vlog/set [spec] – Sets logging levels

**Many interesting features supported, which are defined  
in the targets’ man pages**

## Flow Debugging

---

**Flow tables can become incredibly complex, but OVS has tools to make it easier to debug**

**Here is a set of rules to (poorly) implement a firewall (with an unnecessary resubmit) to block all TCP traffic except port 80:**

```
# Move TCP traffic arriving on port 1 to next stage of "pipeline"!
priority=100,tcp,in_port=1 actions=resubmit:4000! !
```

```
# Allow port TCP port 80 traffic (and implicitly drop all others)!
priority=100,tcp,in_port=4000,tp_dst=80 actions=NORMAL! !
```

```
# Allow all non-TCP traffic arriving on port 1!
priority=90,in_port=1 actions=NORMAL! !
```

```
# Allow all traffic arriving on port 2!
priority=100,in_port=2 actions=NORMAL! ! !
```

# Tracing Flow (ICMP Allowed)

```
root@vm-vswitch:~# ovs-appctl ofproto/trace
"skb_priority(0),in_port(2),skb_mark(0),eth(src=50:54:00:00:00:01, dst=50:54
:00:00:00:03),eth_type(0x0800),ipv4(src=192.168.0.1, dst=192.168.0.2, proto=1
,tos=0,ttl=64,frag=no),icmp(type=8,code=0)"
Bridge: br0
Flow:
icmp,metadata=0,in_port=1,vlan_tci=0x0000,d1_src=50:54:00:00:00:01,d1_dst=5
0:54:00:00:00:03,nw_src=192.168.0.1,nw_dst=192.168.0.2,nw_tos=0,nw_ecn=0,nw
_ttl=64,icmp_type=8,icmp_code=0
Rule: table=0 cookie=0 priority=90,in_port=1
OpenFlow actions=NORMAL
forwarding to learned port
Applied OpenFlow rule

Final flow: unchanged
Relevant fields:
skb_priority=0,icmp,in_port=1,vlan_tci=0x0000/0x1fff,d1_src=50:54:00:00:00:
01,d1_dst=50:54:00:00:00:03,nw_frag=no,icmp_code=0
Datapath actions: 3
Datapath action
```

# Tracing Flow (TCP allowed)

```
root@vm-vswitch:~# ovs-appctl ofproto/trace
"skb_priority(0),in_port(2),skb_mark(0),eth(src=50:54:00:00:00:01, dst=50:5
4:00:00:00:03),eth_type(0x0800),ipv4(src=192.168.0.1, dst=192.168.0.2, proto
=6,tos=0x10,ttl=64,frag=no),tcp(src=56176, dst=80),tcp_flags(0x002)"
Bridge: br0
Flow:
tcp,metadata=0,in_port=1,vlan_tci=0x0000,d1_src=50:54:00:00:00:01,d1_dst=5
0:54:00:00:00:03,nw_src=192.168.0.1,nw_dst=192.168.0.2,nw_tos=16,nw_ecn=0,
nw_ttl=64,tp_src=56176,tp_dst=80,tcp_flags=0x002
Rule: table=0 cookie=0 priority=100,tcp,in_port=1
OpenFlow actions=resubmit:4000
    Resubmitted flow: unchanged
    Resubmitted regs: reg0=0x0 reg1=0x0 reg2=0x0 reg3=0x0 reg4=0x0
reg5=0x0 reg6=0x0 reg7=0x0
    Resubmitted odp: drop
    Rule: table=0 cookie=0 priority=100,tcp,in_port=4000,tp_dst=80
    OpenFlow actions=NORMAL
        forwarding to learned port
Final flow: unchanged
Relevant fields:
skb_priority=0,tcp,in_port=1,vlan_tci=0x0000/0xffff,d1_src=50:54:00:00:00:
01,d1_dst=50:54:00:00:00:03,nw_frag=no,tp_dst=80
Datapath actions: 3
```

First applied OpenFlow rule

Second applied OpenFlow rule

Datapath flow description

Datapath action

# Tracing Flow (TCP denied)

```
root@vm-vswitch:~# ovs-appctl ofproto/trace
"skb_priority(0),in_port(2),skb_mark(0),eth(src=50:54:00:00:00:01, dst=50:54:00:00:00:03),eth_type(0x0800),ipv4(src=192.168.0.1, dst=192.168.0.2, proto=6, tos=0x10, ttl=64, frag=no),tcp(src=56177, dst=100),tcp_flags(0x002)"
Bridge: br0
Flow:
tcp,metadata=0,in_port=1,vlan_tci=0x0000,d1_src=50:54:00:00:00:01,d1_dst=50:54:00:00:00:03,nw_src=192.168.0.1,nw_dst=192.168.0.2,nw_tos=16,nw_ecn=0,nw_ttl=64,
tp_src=56177,tp_dst=100,tcp_flags=0x002
Rule: table=0 cookie=0 priority=100,tcp,in_port=1 ← First applied OpenFlow Rule
OpenFlow actions=resubmit:4000

    Resubmitted flow: unchanged
    Resubmitted regs: reg0=0x0 reg1=0x0 reg2=0x0 reg3=0x0 reg4=0x0 reg5=0x0
    reg6=0x0 reg7=0x0
    Resubmitted odp: drop
    No match
    No matching second flow,
    so implicit drop
    Datapath flow description
Final flow: unchanged
Relevant fields: skb_priority=0,tcp,in_port=1,nw_frag=no,tp_dst=100
Datapath actions: drop ← Datapath action
```

# Logging

- **ovs-appctl** configures running OVS daemons
- Most common use is to modify logging levels
- By default configures **ovs-vswitchd**, but “-t” option changes target
- Default level for log files is “info”, only thing lower is “dbg”

```
root@vm-vswitch:~# ovs-appctl vlog/list
      console    syslog    file
      -----  -----
bridge          EMER        ERR       INFO
vswitchd        EMER        ERR       INFO
...
root@vm-vswitch:~# ovs-appctl vlog/set ofproto:file:dbg
```

# Log Files

- **Open vSwitch logs: /var/log/openvswitch/\***
  - ovs-vswitchd.log
  - ovsdb-server.log
- **System: /var/log/messages**
- **Configuration Database: /etc/openvswitch/conf.d**

# Install OVS

- Install openVswitch

```
#sudo apt-get install openvswitch-switch
```

- Install the ovs-docker utility

- #cd /usr/bin
- wget <https://raw.githubusercontent.com/openvswitch/ovs/master/utilities/ovs-docker>
- #chmod a+rwx ovs-docker

# Config bridge

- Config bridge

```
#ovs-vsctl add-br ovs-br1
```

```
#ifconfig ovs-br1 173.16.1.1 netmask 255.255.255.0 up
```

(if no ifconfig, install it )

```
#apt install net-tools
```

- Show bridge

```
#ovs-vsctl show
```

# Add docker

- Add container

```
#docker run -dt --name c3 alpine sleep 1d
```

```
#docker run -dt --name c4 alpine sleep 1d
```

- Connect the container to the OVS

```
#ovs-docker add-port ovs-br1 eth1 c3 --ipaddress=173.16.1.20/24
```

```
#ovs-docker add-port ovs-br1 eth1 c4 --ipaddress=173.16.1.30/24
```

- Check container

```
#docker exec c3 ip link list
```

```
#docker exec c4 ip link list
```

```
root@ubuntu:/usr/bin# docker exec c3 ip link list
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000  
link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```
21: eth0@if22: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP link/ether 02:42:ac:11:00:03 brd ff:ff:ff:ff:ff:ff
```

```
root@ubuntu:/usr/bin# docker exec c3 ip link list  
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000  
link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00  
21: eth0@if22: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP  
link/ether 02:42:ac:11:00:03 brd ff:ff:ff:ff:ff:ff
```

# Set IP Address

```
root@ubuntu:/usr/bin# ovs-docker add-port ovs-br1 eth1 c3 --ipaddress=173.16.1.20/24
```

```
root@ubuntu:/usr/bin# docker exec c3 ip address list
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
```

```
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```
    inet 127.0.0.1/8 scope host lo
```

```
        valid_lft forever preferred_lft forever
```

```
21: eth0@if22: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
```

```
    link/ether 02:42:ac:11:00:03 brd ff:ff:ff:ff:ff:ff
```

```
    inet 172.17.0.3/16 brd 172.17.255.255 scope global eth0
```

```
        valid_lft forever preferred_lft forever
```

```
23: eth1@if24: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP qlen 1000
```

```
    link/ether 8e:52:06:9a:7c:9d brd ff:ff:ff:ff:ff:ff
```

```
    inet 173.16.1.20/24 scope global eth1
```

```
        valid_lft forever preferred_lft forever
```

```
root@ubuntu:/usr/bin# docker exec c3 ip link list
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN  
qlen 1000
```

```
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
```

```
21: eth0@if22: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500  
qdisc noqueue state UP
```

```
    link/ether 02:42:ac:11:00:03 brd ff:ff:ff:ff:ff:ff
```

```
23: eth1@if24: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500  
qdisc noqueue state UP qlen 1000
```

```
    link/ether 8e:52:06:9a:7c:9d brd ff:ff:ff:ff:ff:ff
```

```
root@ubuntu:/usr/bin# docker exec c4 ping 173.16.1.20
```

```
PING 173.16.1.20 (173.16.1.20): 56 data bytes
```

```
64 bytes from 173.16.1.20: seq=0 ttl=64 time=0.551 ms
```


```
64 bytes from 173.16.1.20: seq=1 ttl=64 time=0.102 ms
```

```
64 bytes from 173.16.1.20: seq=2 ttl=64 time=0.059 ms
```

```
root@ubuntu:/usr/bin# ovs-vsctl show
cd18c9a3-49a1-4608-999a-5fd3da3c9e1b
Bridge "ovs-br1"
  Port "fcf399d5e1cb4_l"
    Interface "fcf399d5e1cb4_l"
  Port "ovs-br1"
    Interface "ovs-br1"
      type: internal
  Port "76fec545fdf74_l"
    Interface "76fec545fdf74_l"
  Port "960d648383a84_l"
    Interface "960d648383a84_l"
  Port "d5607674b29d4_l"
    Interface "d5607674b29d4_l"
  Port "cc09edcb865d4_l"
    Interface "cc09edcb865d4_l"
ovs_version: "2.9.5"
root@ubuntu:/usr/bin#
```


# Network Using Namespace

- Very Simple Network using Full System Virtualization



# Network using Namespace

- Problems
  - Too much work even for creating such a simple network topology
  - Not programmable




```
sudo bash
# Create host namespaces
ip netns add h1
ip netns add h2
# Create switch
ovs-vsctl add-br s1
# Create links
ip link add h1-eth0 type veth peer name s1-eth1
ip link add h2-eth0 type veth peer name s1-eth2
ip link show
# Move host ports into namespaces
ip link set h1-eth0 netns h1
ip link set h2-eth0 netns h2
ip netns exec h1 ip link show
ip netns exec h2 ip link show
```

```
# Connect switch ports to OVS
ovs-vsctl add-port s1 s1-eth1
ovs-vsctl add-port s1 s1-eth2
ovs-vsctl show
# Set up OpenFlow controller
ovs-vsctl set-controller s1 tcp:127.0.0.1
ovs-controller ptcp: &
ovs-vsctl show
# Configure network
ip netns exec h1 ifconfig h1-eth0 10.1
ip netns exec h1 ifconfig lo up
ip netns exec h2 ifconfig h2-eth0 10.2
ip netns exec h1 ifconfig lo up
ifconfig s1-eth1 up
ifconfig s1-eth2 up
# Test network
ip netns exec h1 ping -c1 10.2
```

# Mininet

- A simple command-line tool / API which can ease the work
- The solution should allow us to easily create topologies for varying size, up to hundreds and thousand of nodes



# Lab

- ทดสอบหัวข้อ 7.6 การทดสอบการทำงานของโอลูเคนวีสวิตซ์เบื้องต้น
  - Docker 4 ตัว พร้อม VLAN
- ทดสอบการใช้ OVS บน Mininet (บทที่ 8)
  - mn --topo tree,2 --mac --arp --switch ovsk
    - ทดสอบการ Ping สามารถ ping ได้ทั้งหมดหรือไม่
  - mn --topo tree, 2 --mac --arp --switch ovsk --controller
    - ทดสอบการ Ping สามารถ ping ได้ทั้งหมดหรือไม่ แก้ไข อย่างไร