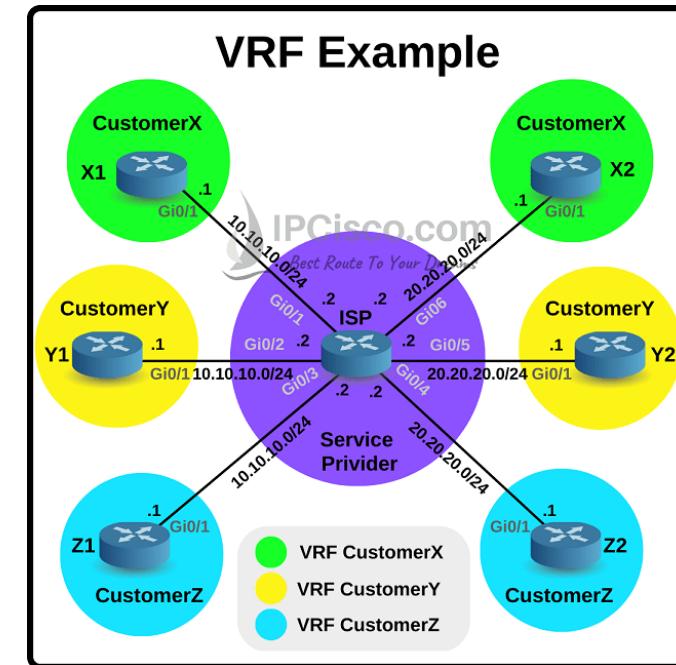
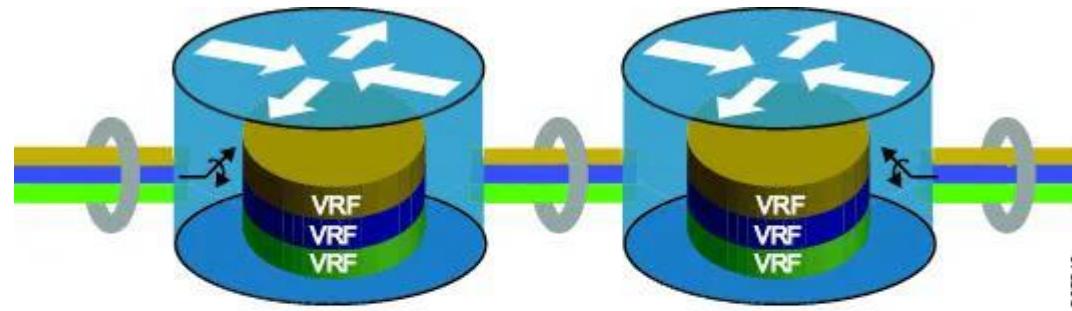
The background of the slide features a series of vertical, glowing lines of varying colors and intensities. These lines are primarily yellow and green, creating a sense of depth and motion against a dark, textured background.

# VxLAN Basic and FRR

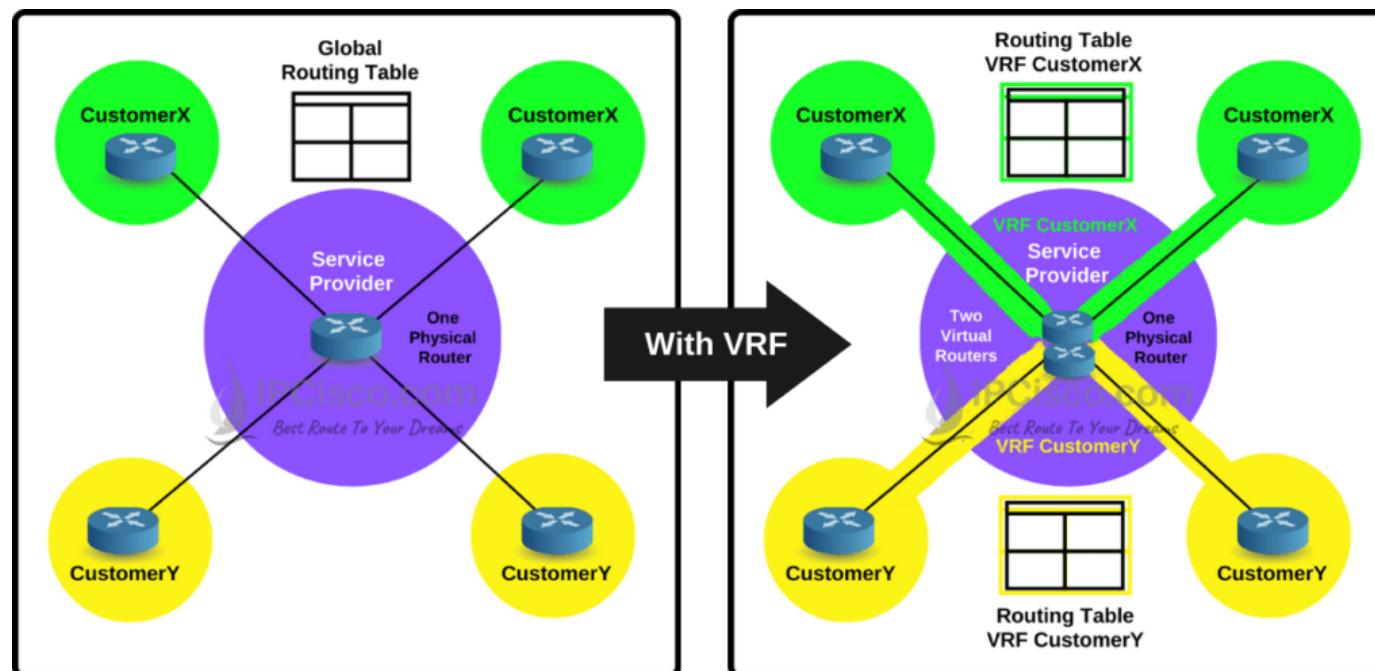
# Virtual Routing and Forwarding (VRF)

- A technique which creates multiple virtual networks within a single network entity ([Figure 1](#)). In a single network component, multiple VRF resources create the isolation between virtual networks.



# VRF and VLAN

- **VRFs** are similar to **VLANs** used in switches at layer 2 of the [OSI Reference Model](#). But VRF technology works at **Layer 3**. As you know, with VLANs, we can create multiple virtual switches in a physical switch. With VRFs, we can create multiple routers in a physical router.



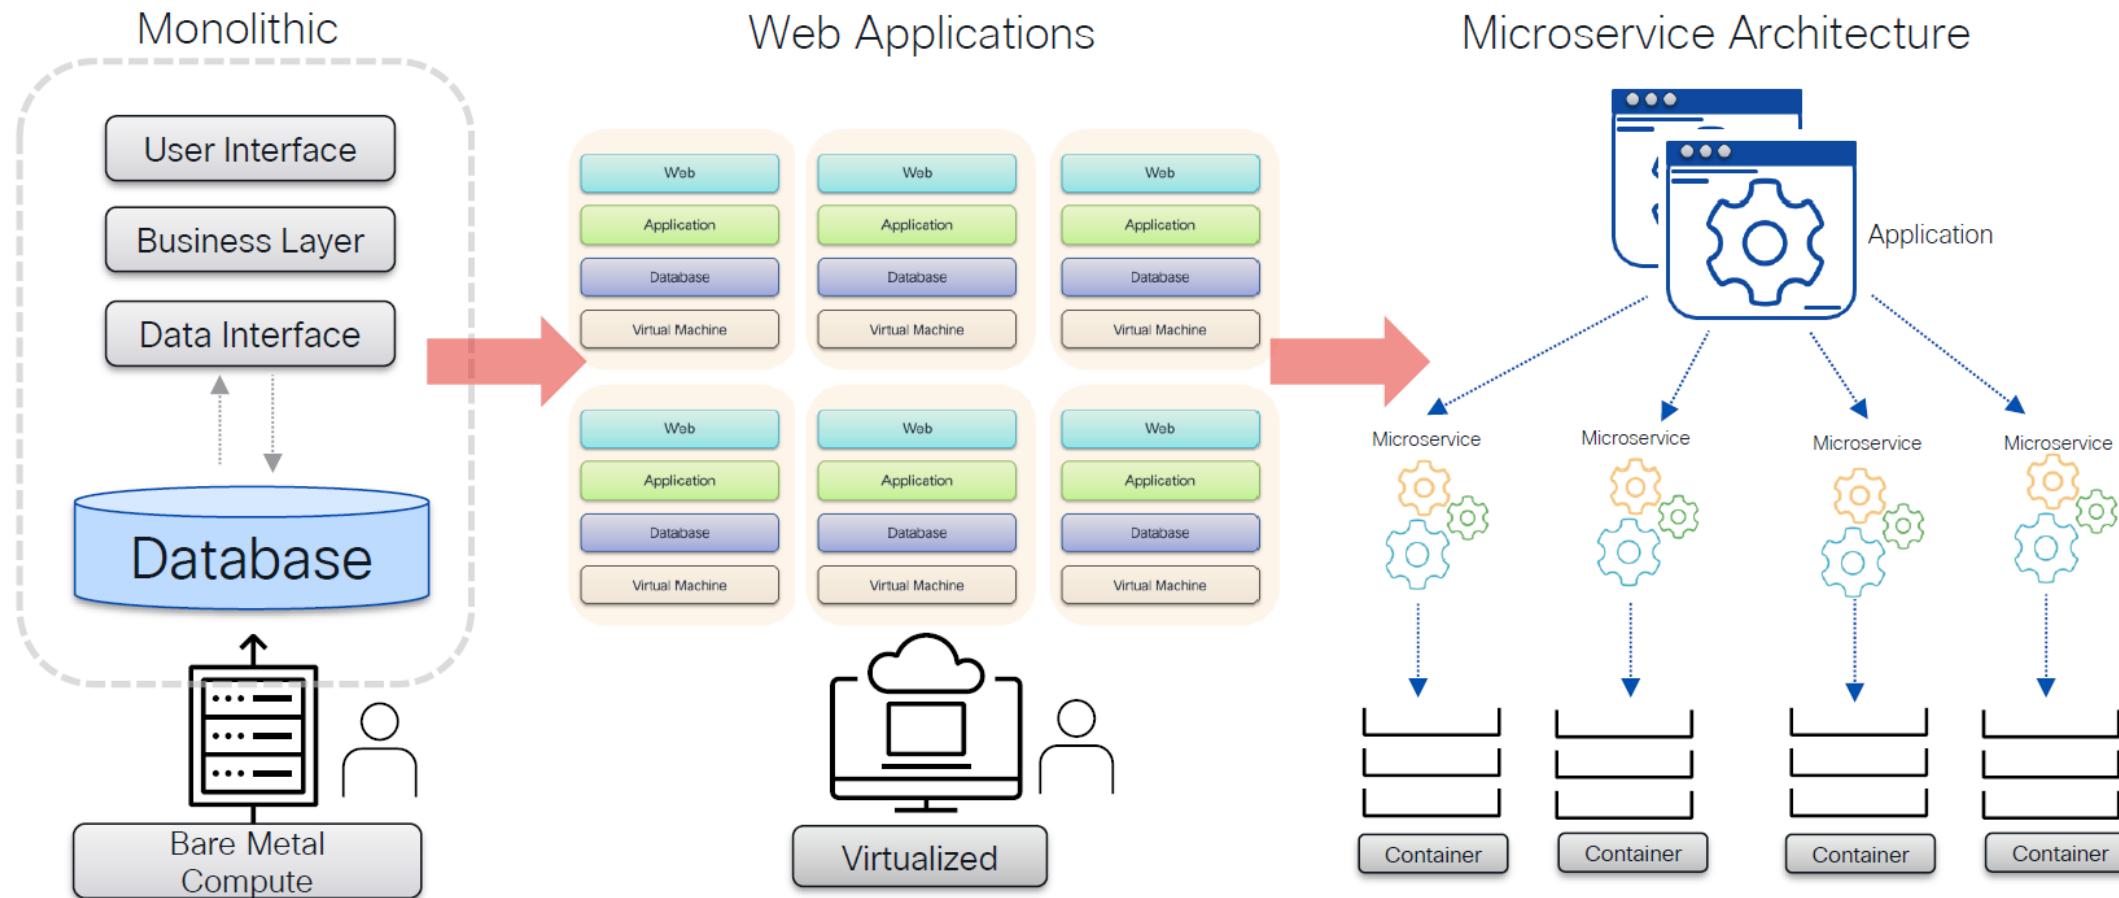
# Benefits of VRFs

- With VRFs, we can create multiple virtual routers in a single physical router with isolated routing and forwarding.
- With VRFs, we can use overlapping IP addresses on created separate virtual routers.
- With separate and small routing tables, VRFs provide easy troubleshooting.
- Useful for MP-BGP and MPLS deployments.

# VxLAN

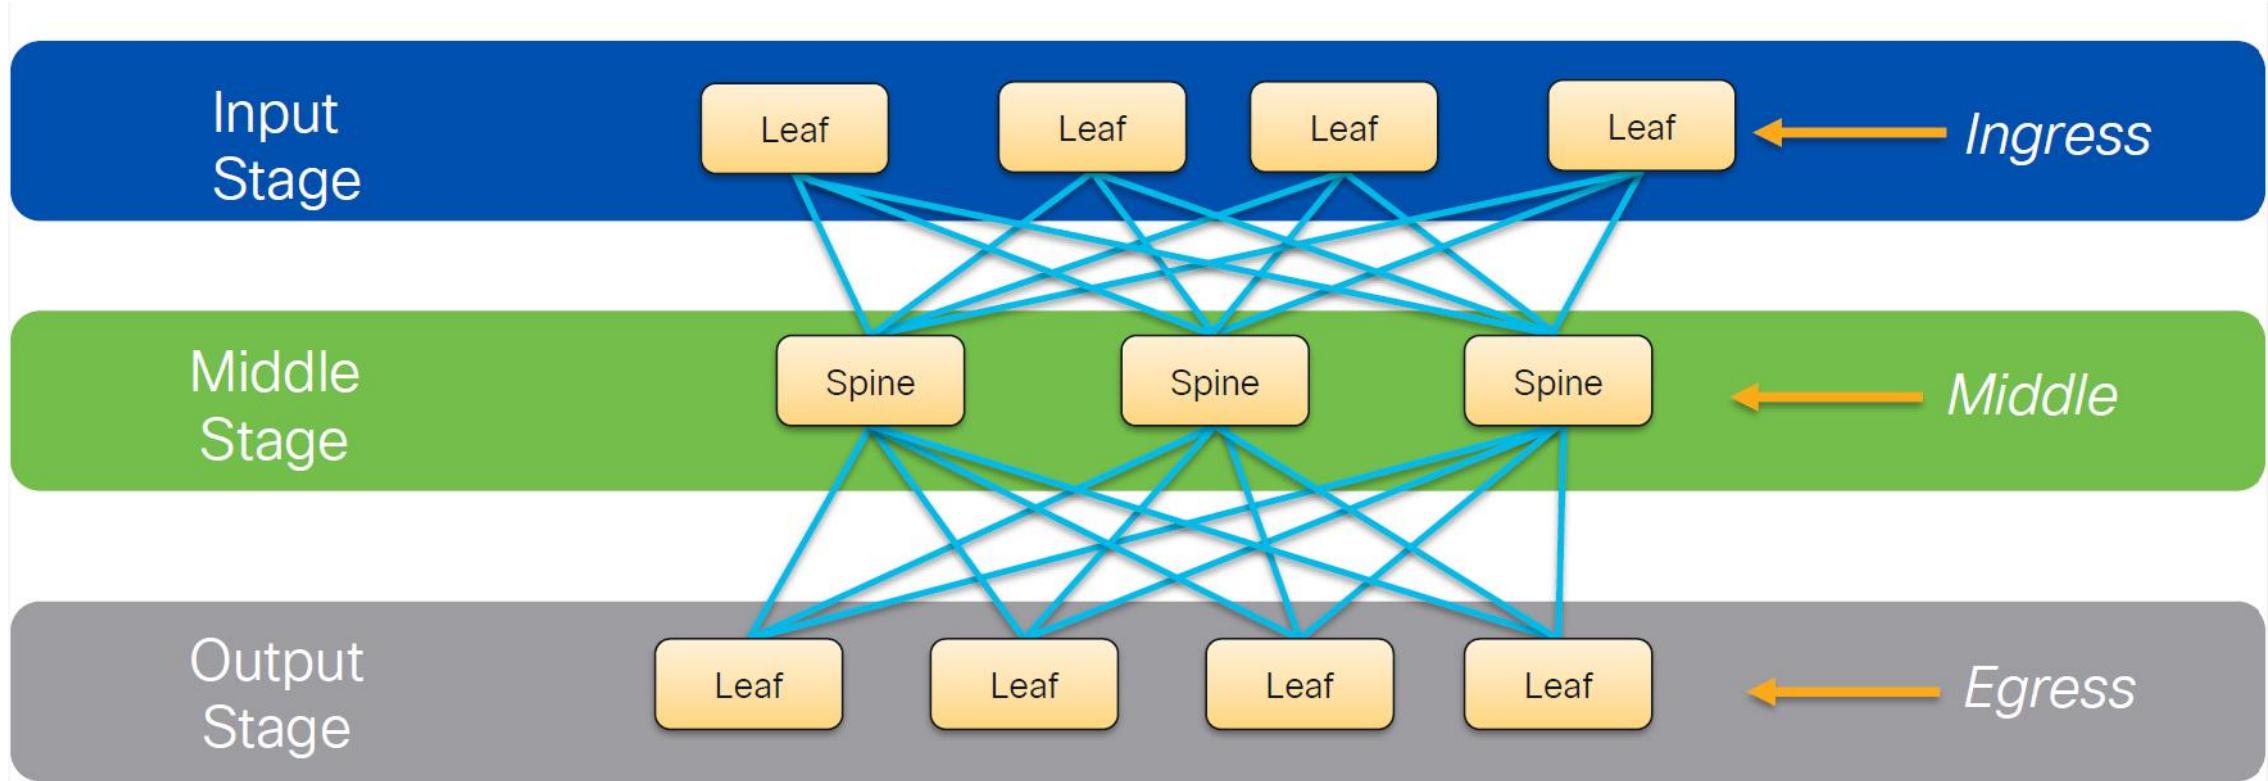
BRKDCN-1621

# Evolution of the Application Architectures



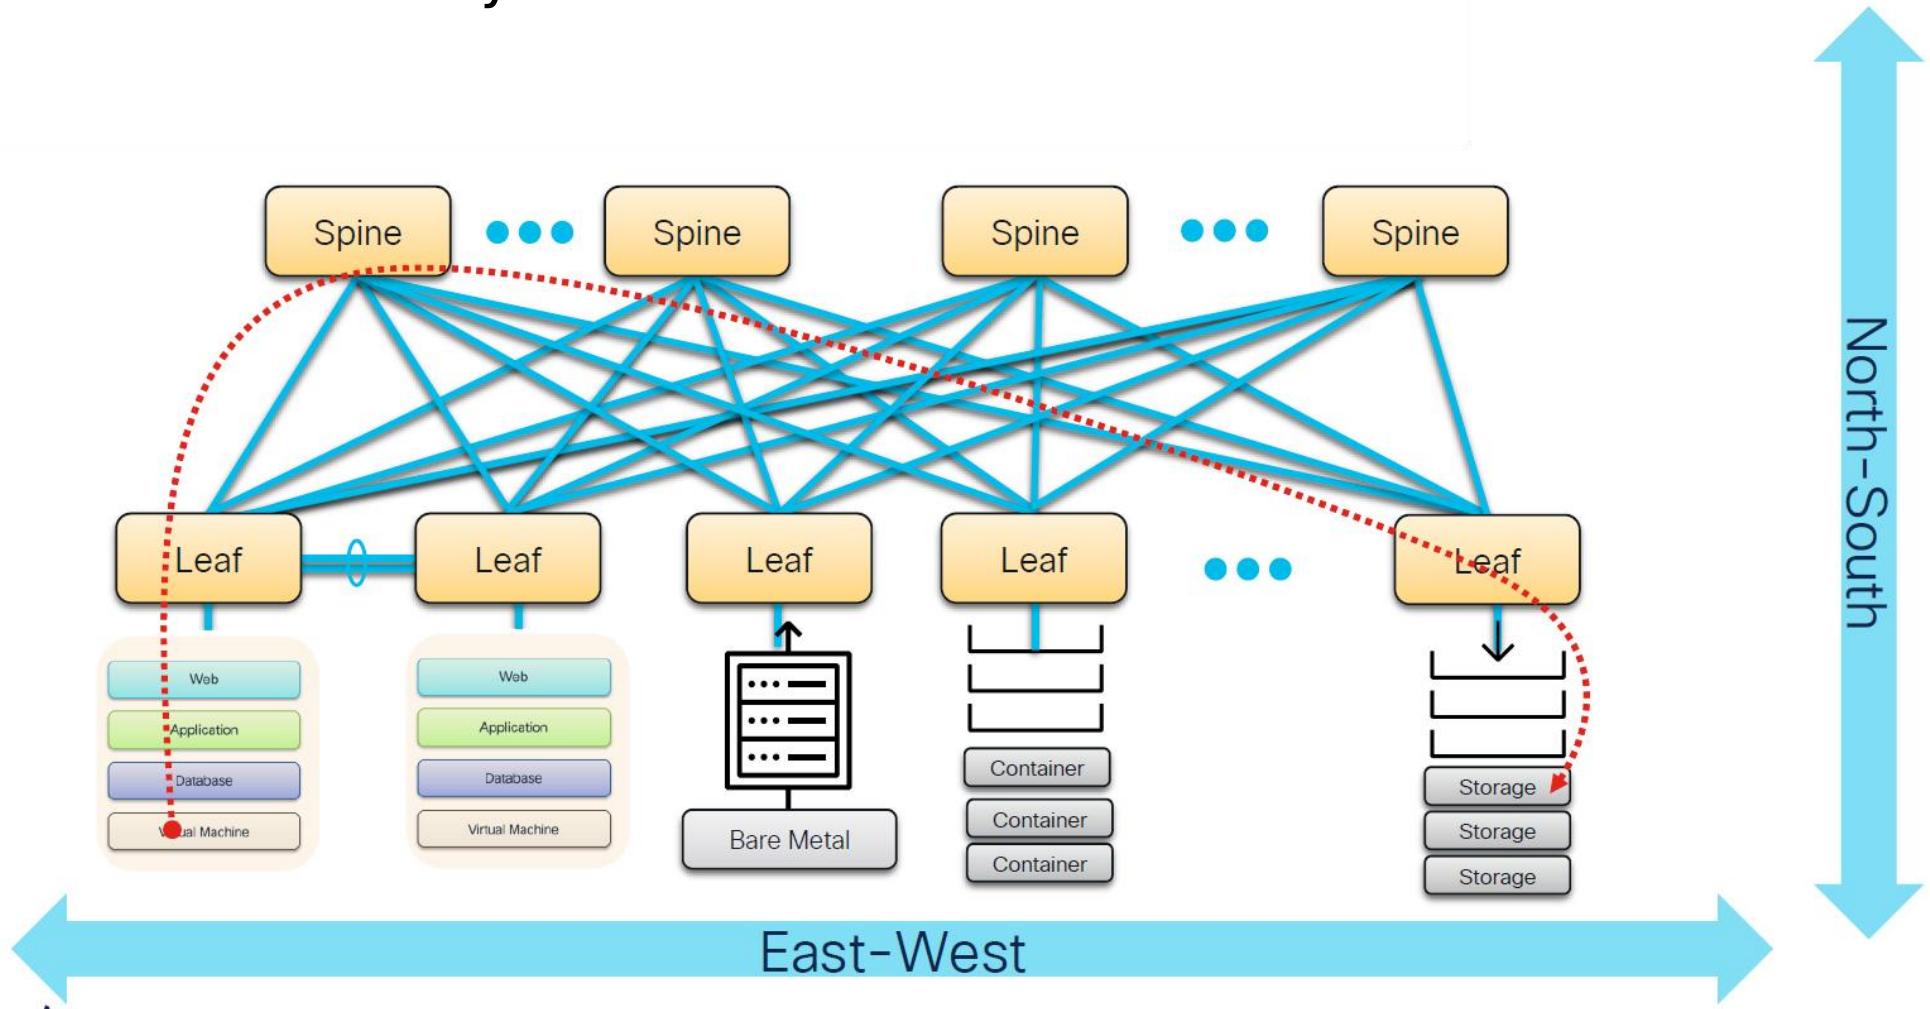
# Evolution of the Data Center Network

## 3-Stage Clos Fabric

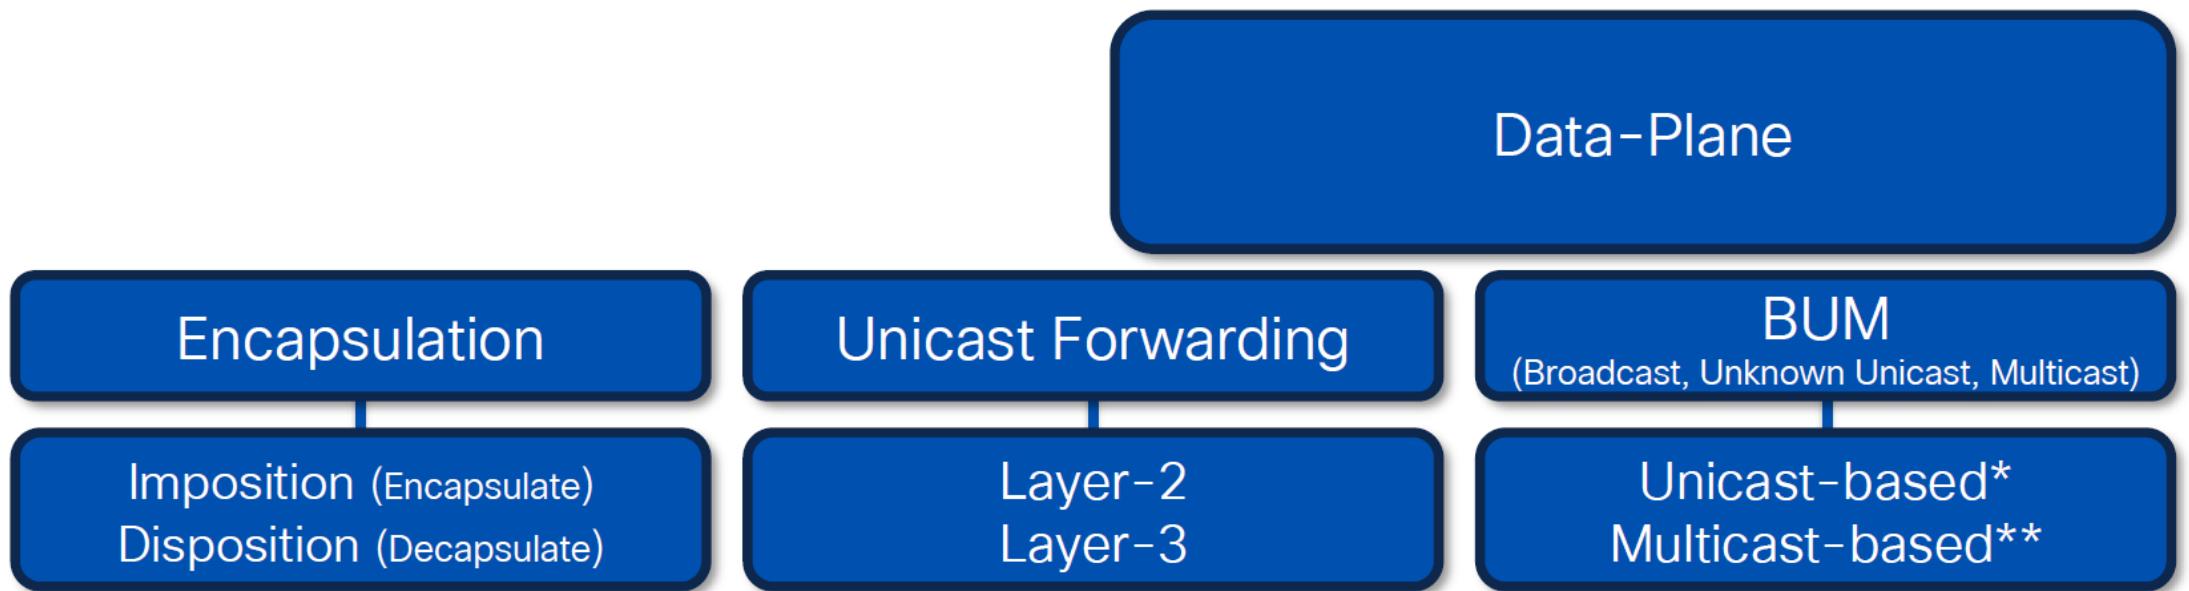


# Evolution of the Data Center Network

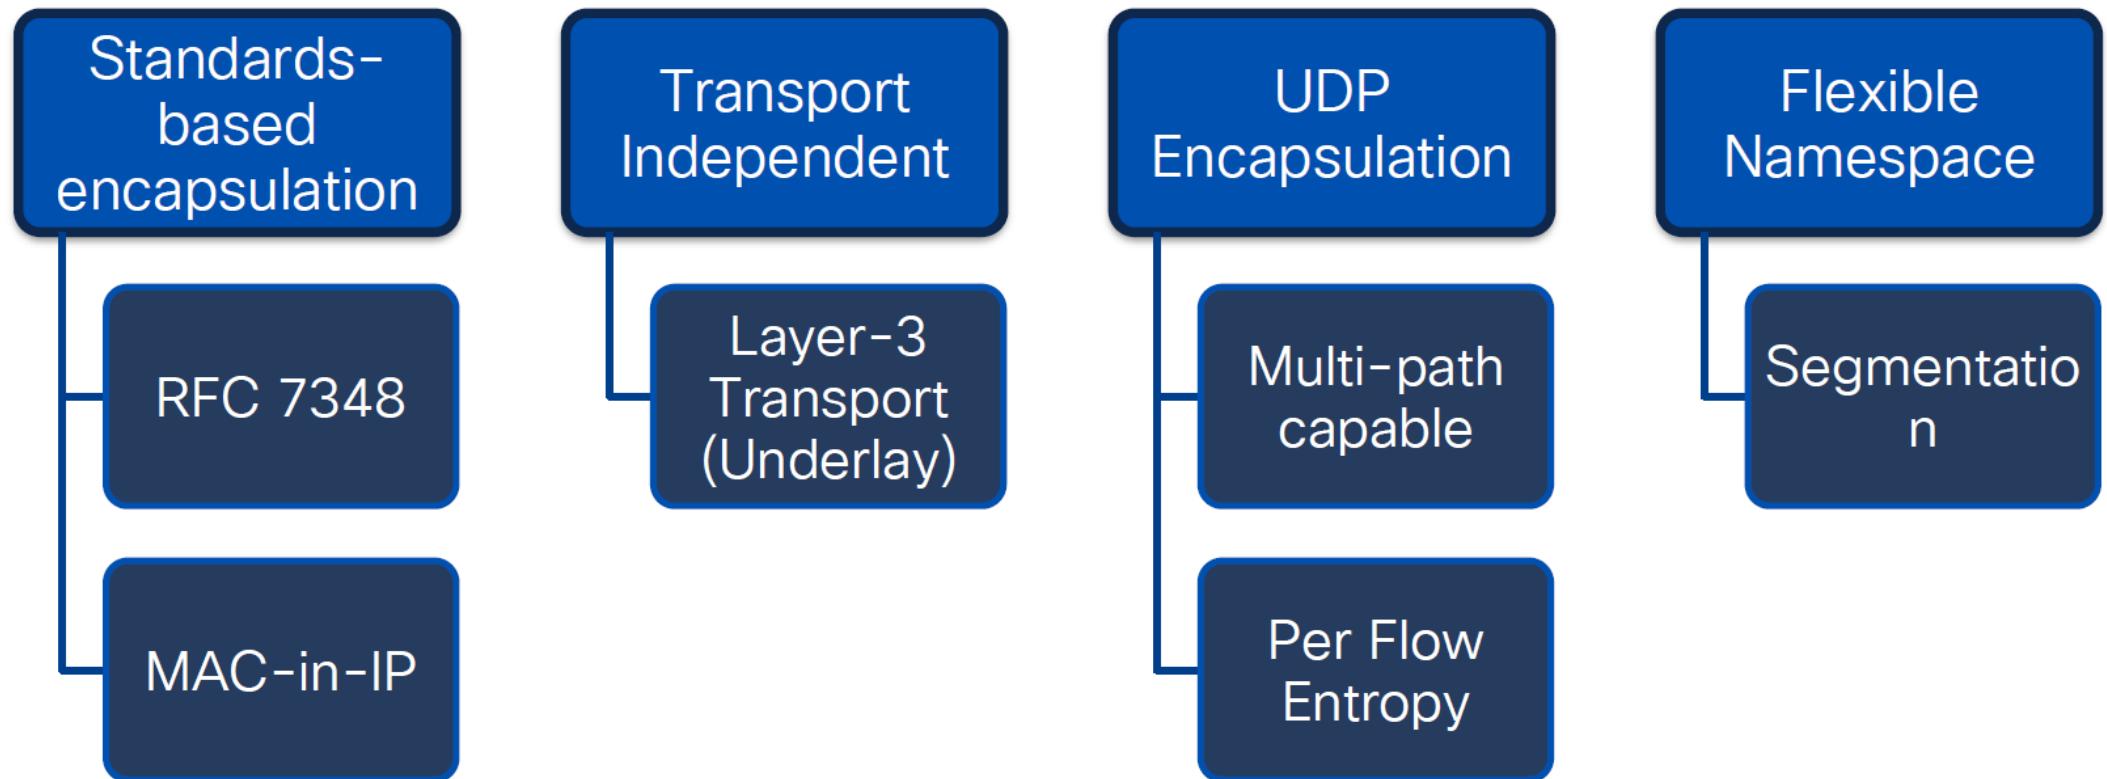
The need for network overlays



# Data Plane



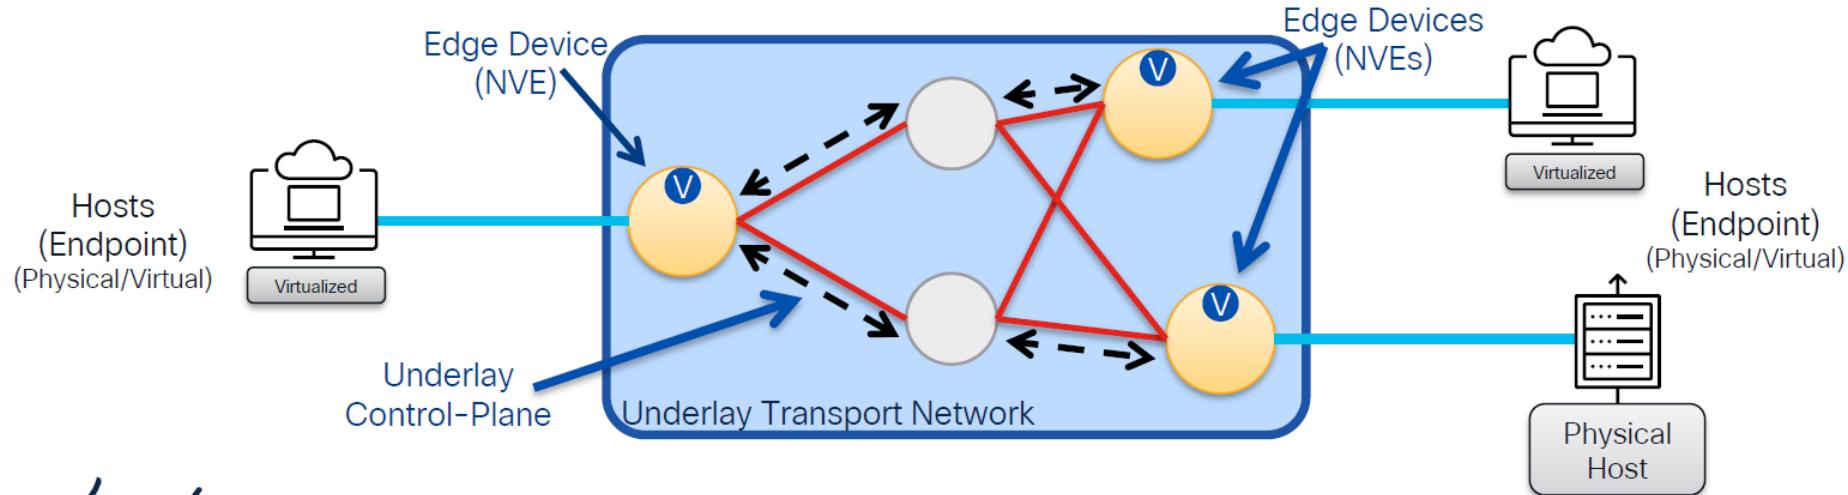
# What is VXLAN?



# Underlay Taxonomy

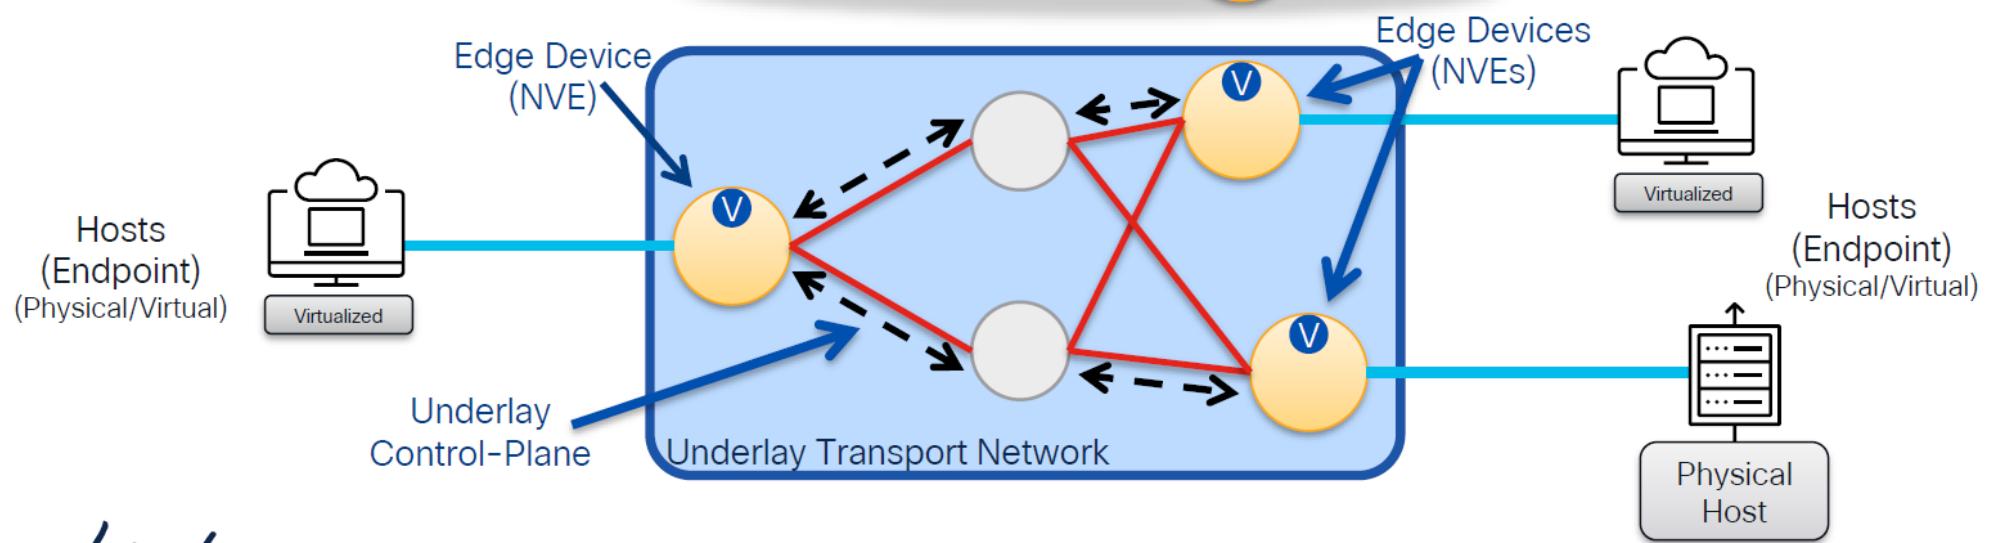
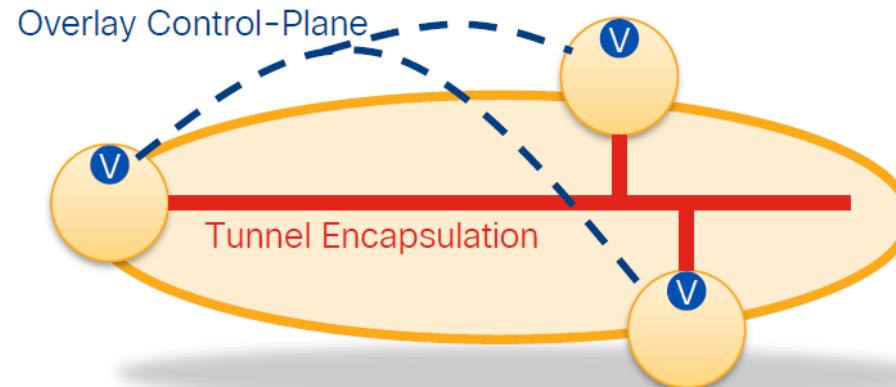
- Edge Devices host the VTEP
- Responsible for the encapsulation and decapsulation of the VXLAN Header

V VTEP: VXLAN Tunnel Endpoint  
NVE: Network Virtualisation Edge

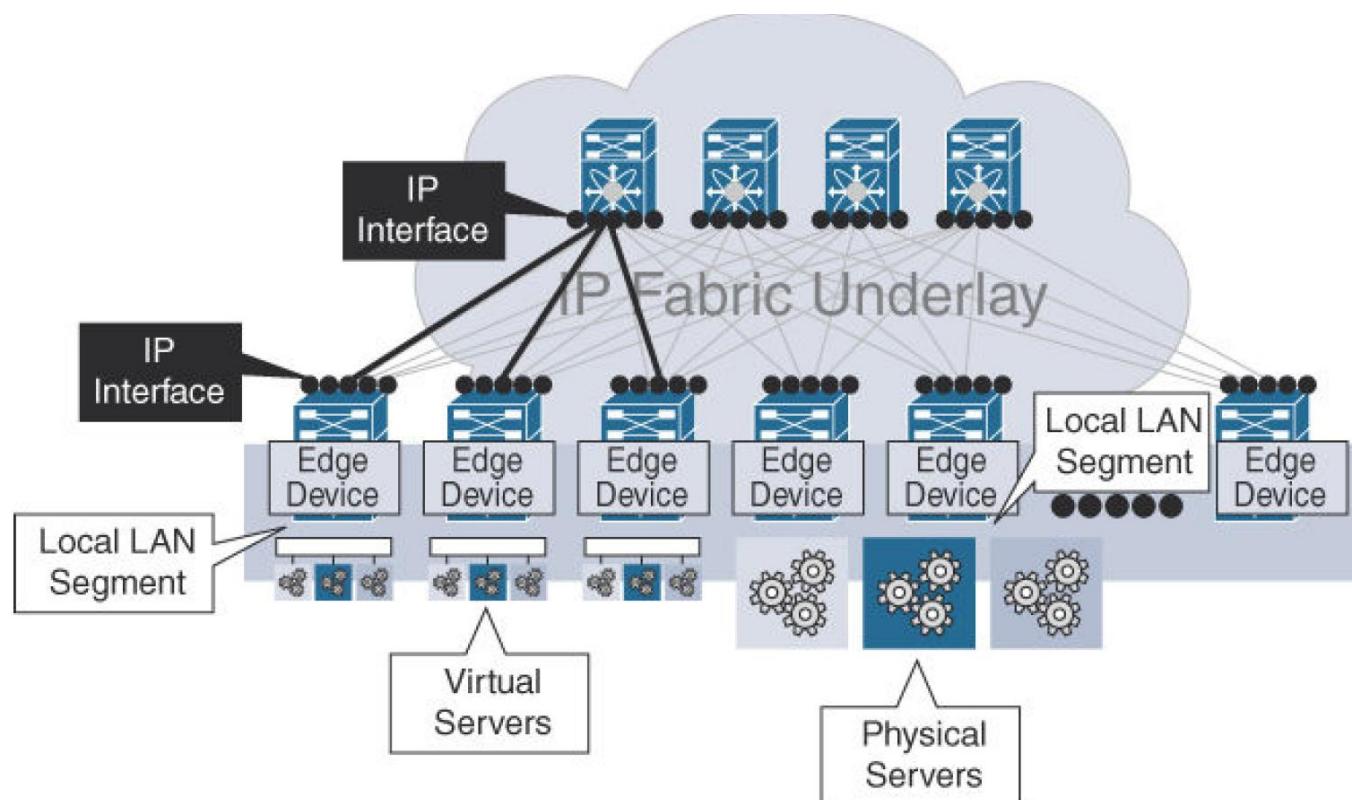


# Overlay Taxonomy

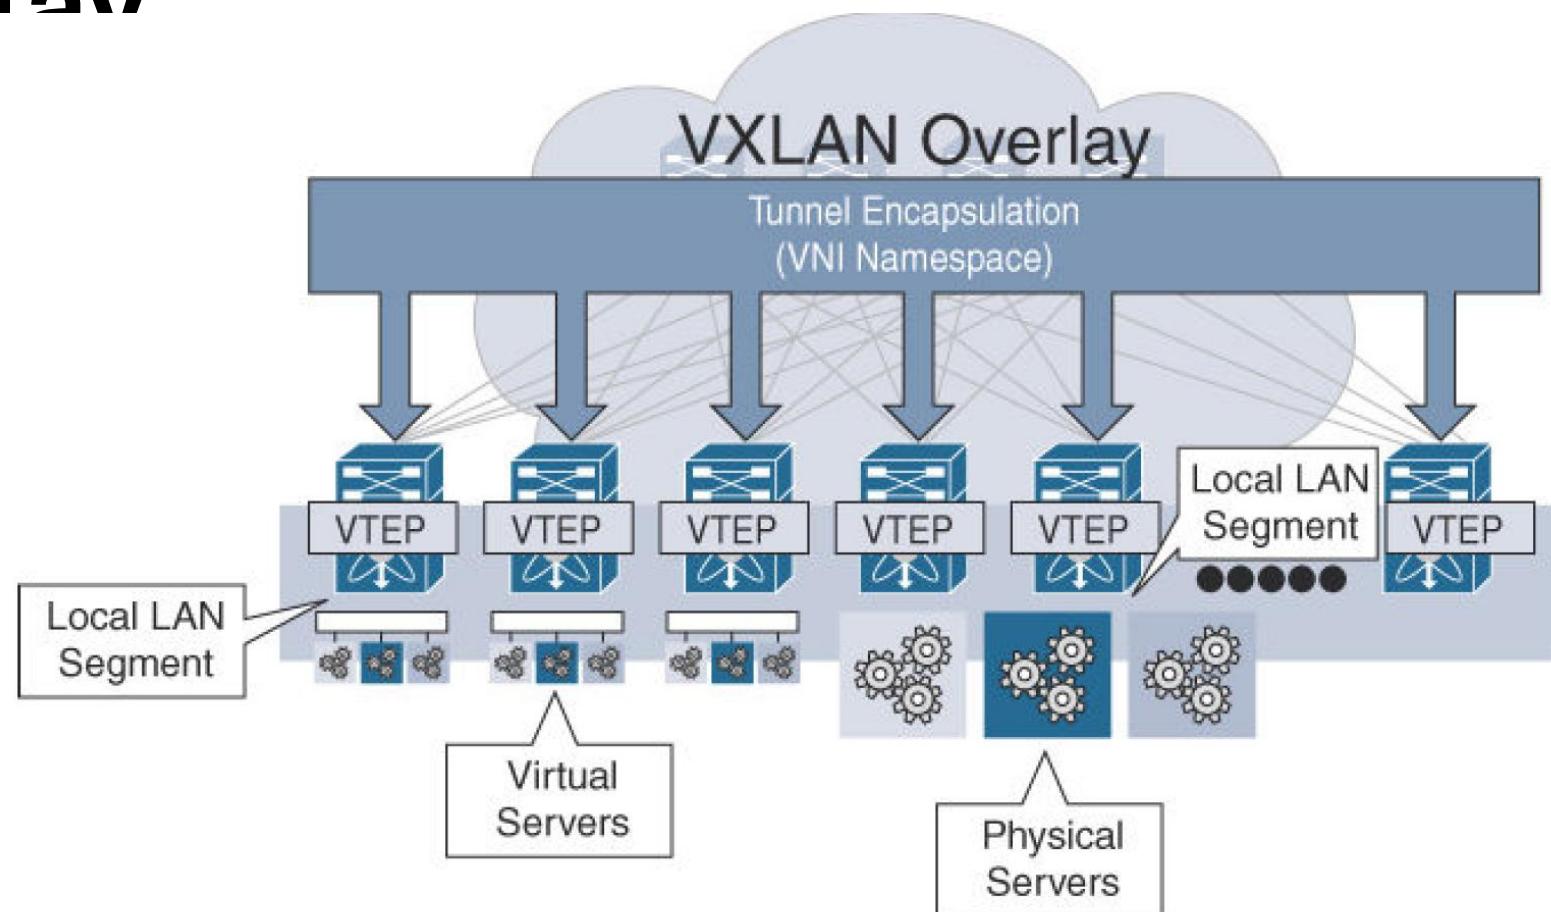
Service = Virtual Network  
Identifier = VN Identifier (VNI/VNID)  
VVTEP: VXLAN Tunnel Endpoint  
NVE: Network Virtualisation Edge



# Overview in Data Center: Underlay



# VxLAN Overlay

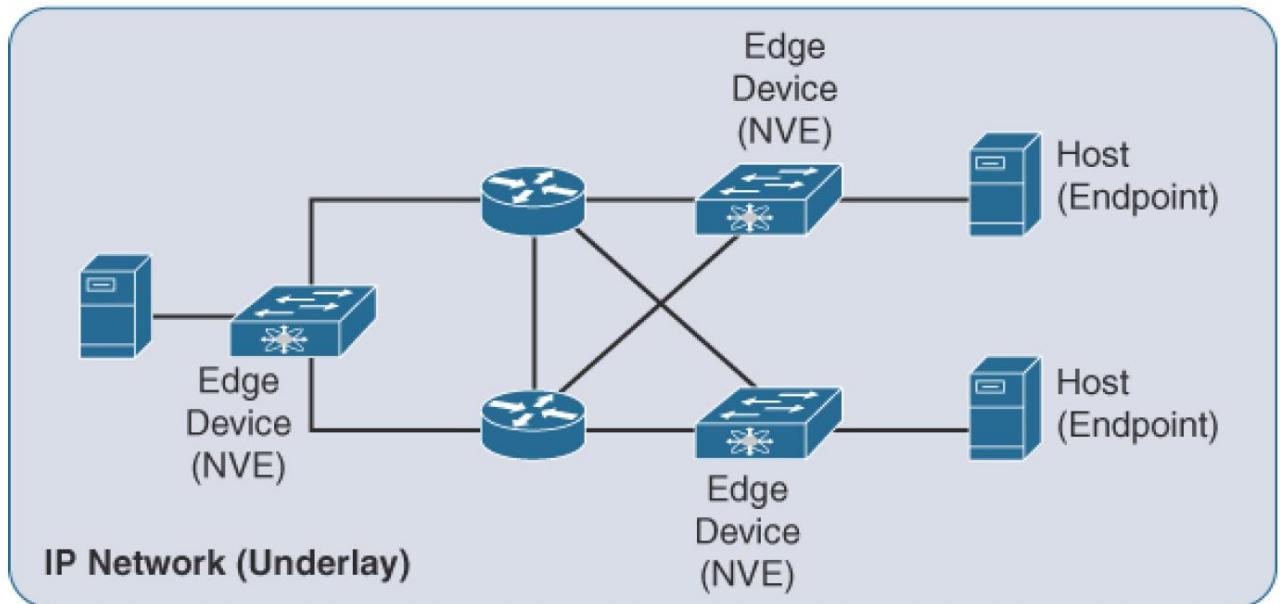
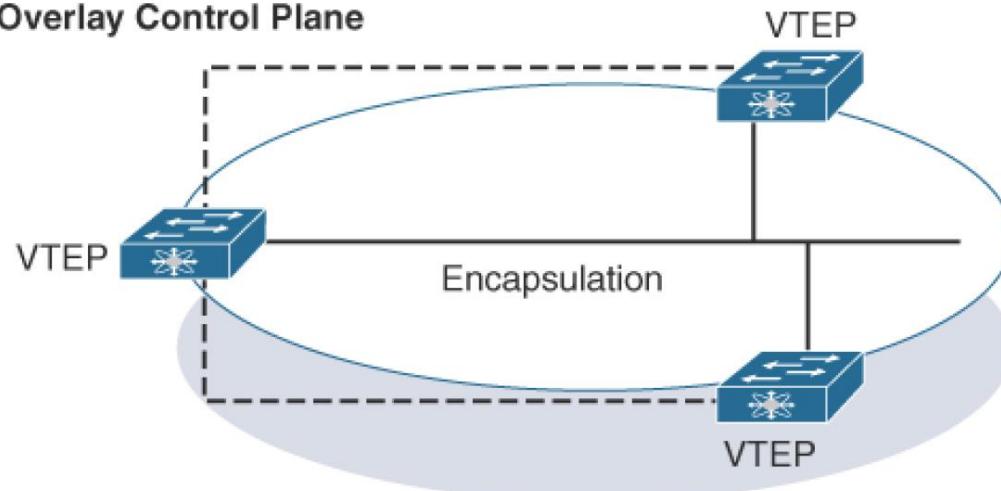


VTEP: VxLAN Tunnel Endpoint

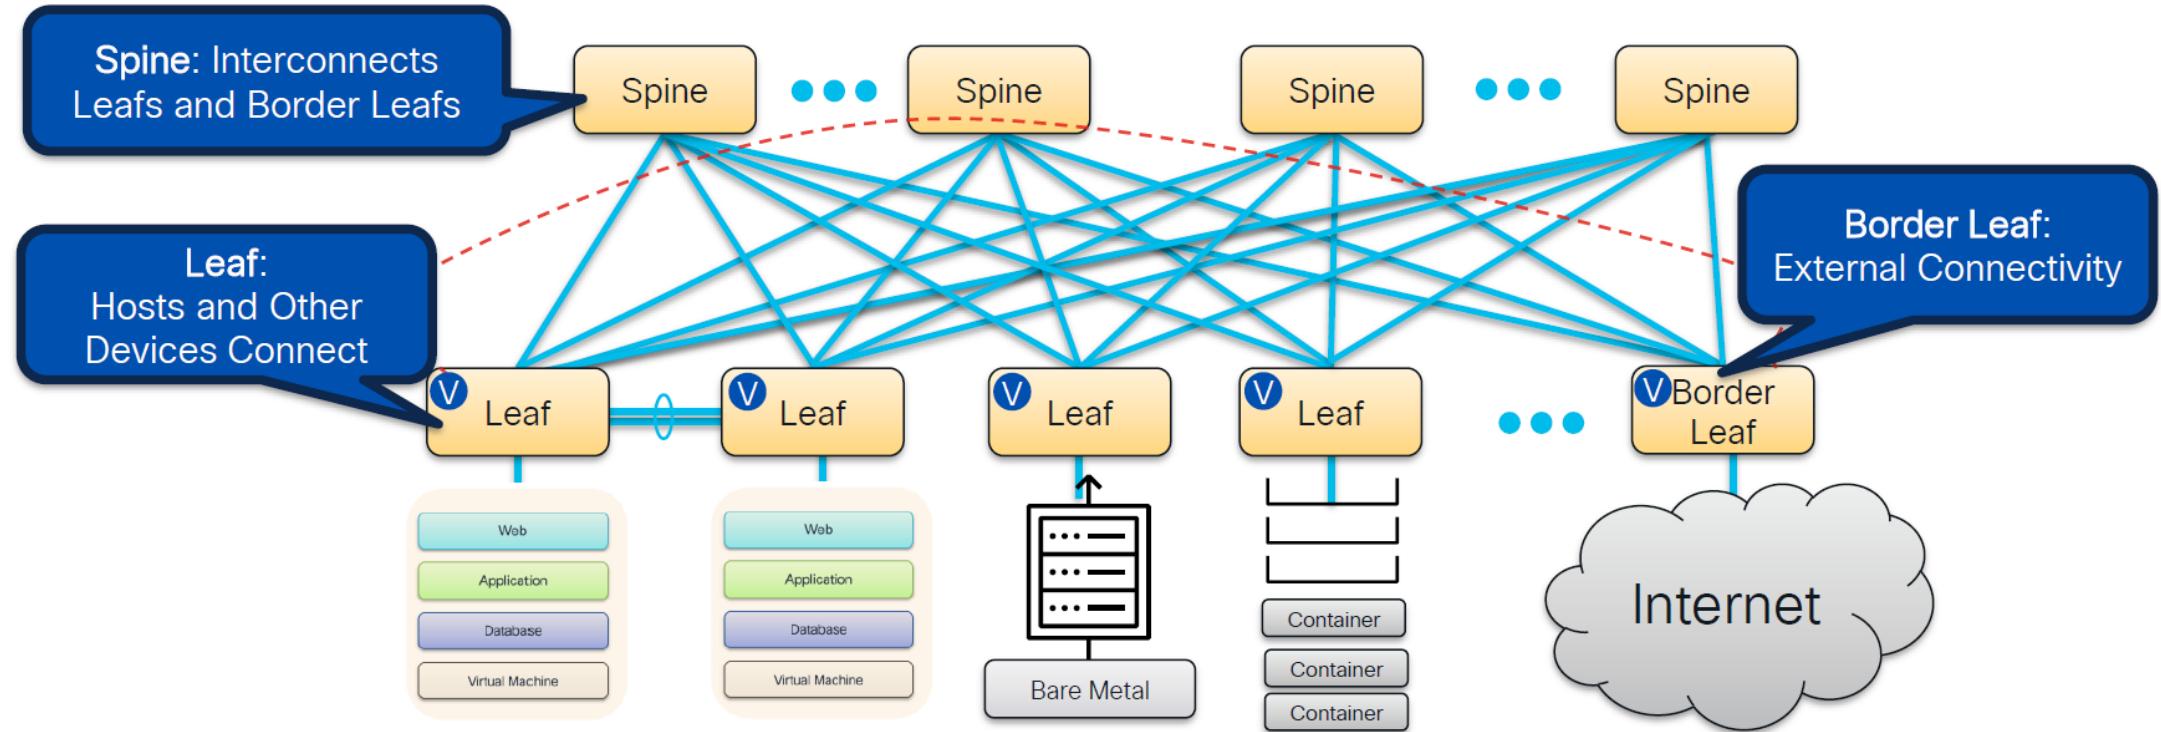
VNI/VNID: VxLAN Network Identifier

Building data centers with VxLAN BPG EVPN

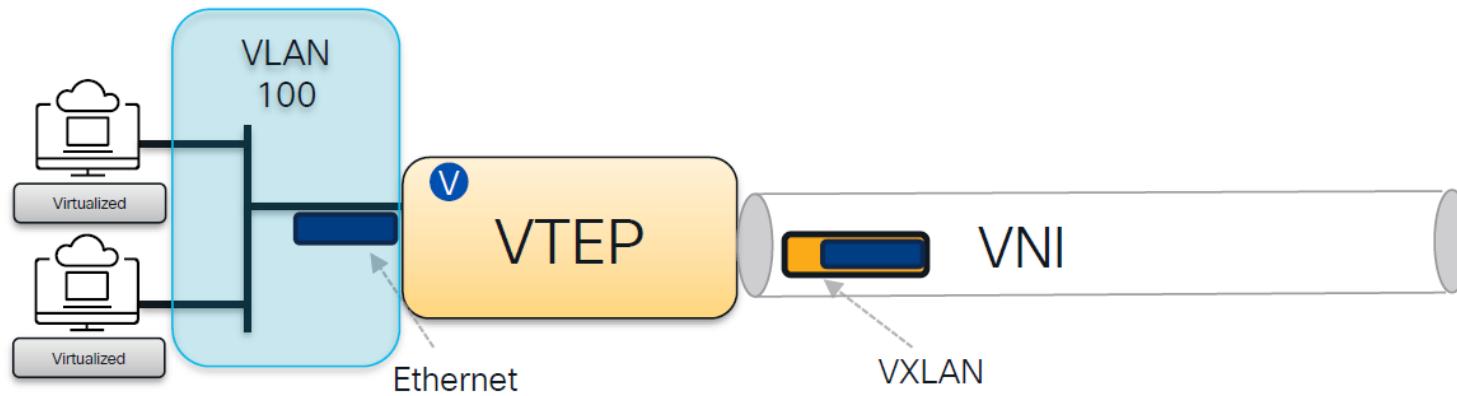
### Overlay Control Plane



# VTEP Device Roles

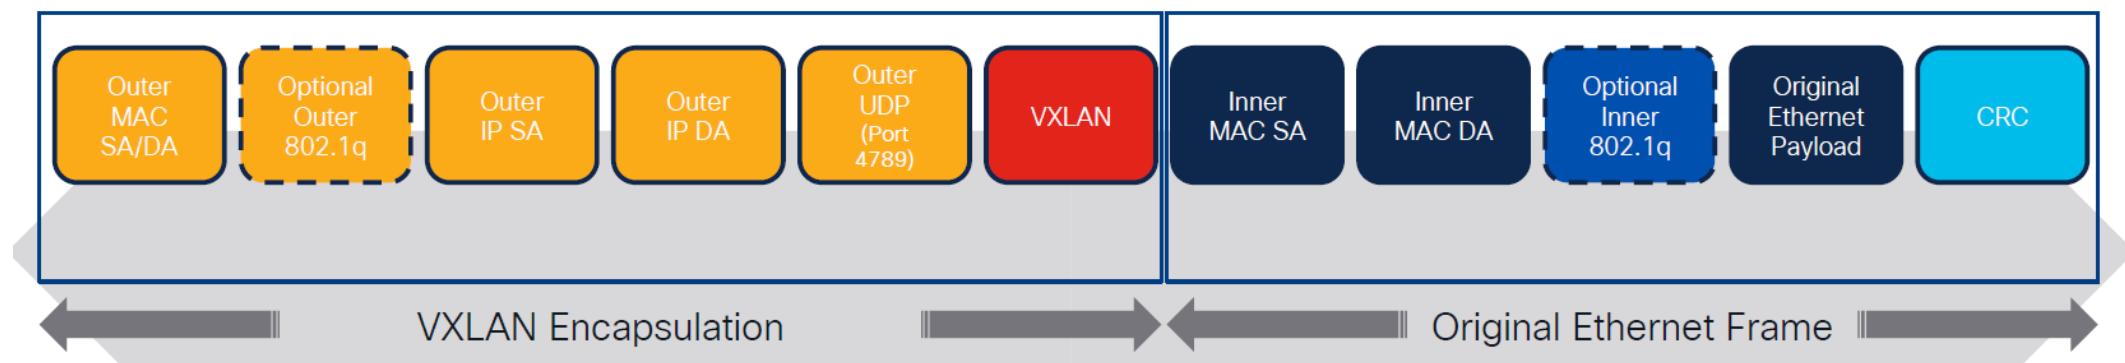
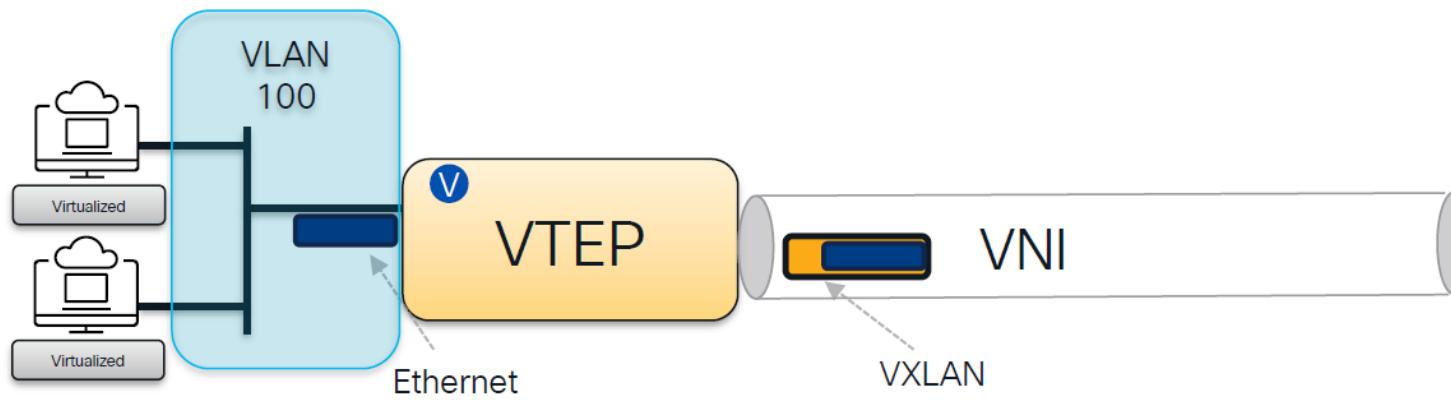


# VXLAN (or Virtual) Tunnel Endpoint (VTEP)



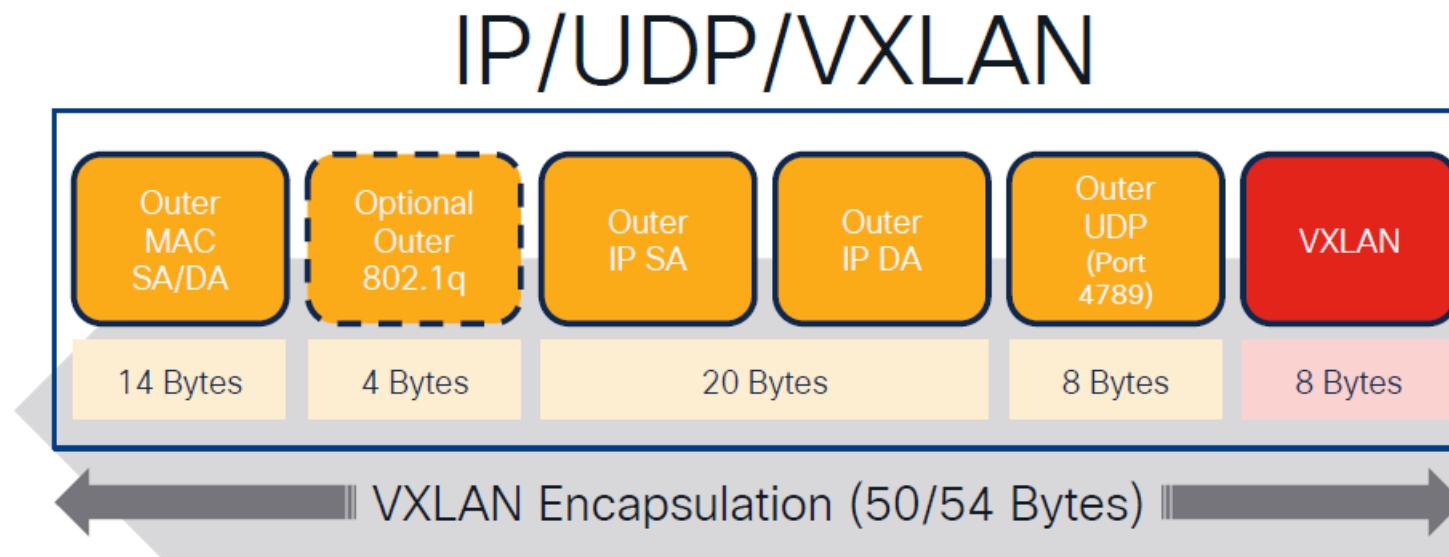
- VXLAN Tunnel Endpoint - Network Virtualisation Edge
- Each VTEP is uniquely identified by an IP Address
- VTEP Discovers or learns remote VTEPs, and end hosts attached to them
- VTEP bridges when forwarding packets within the same VNI and Routes for Inter-VNI traffic

# VXLAN Encapsulation / Packet Format



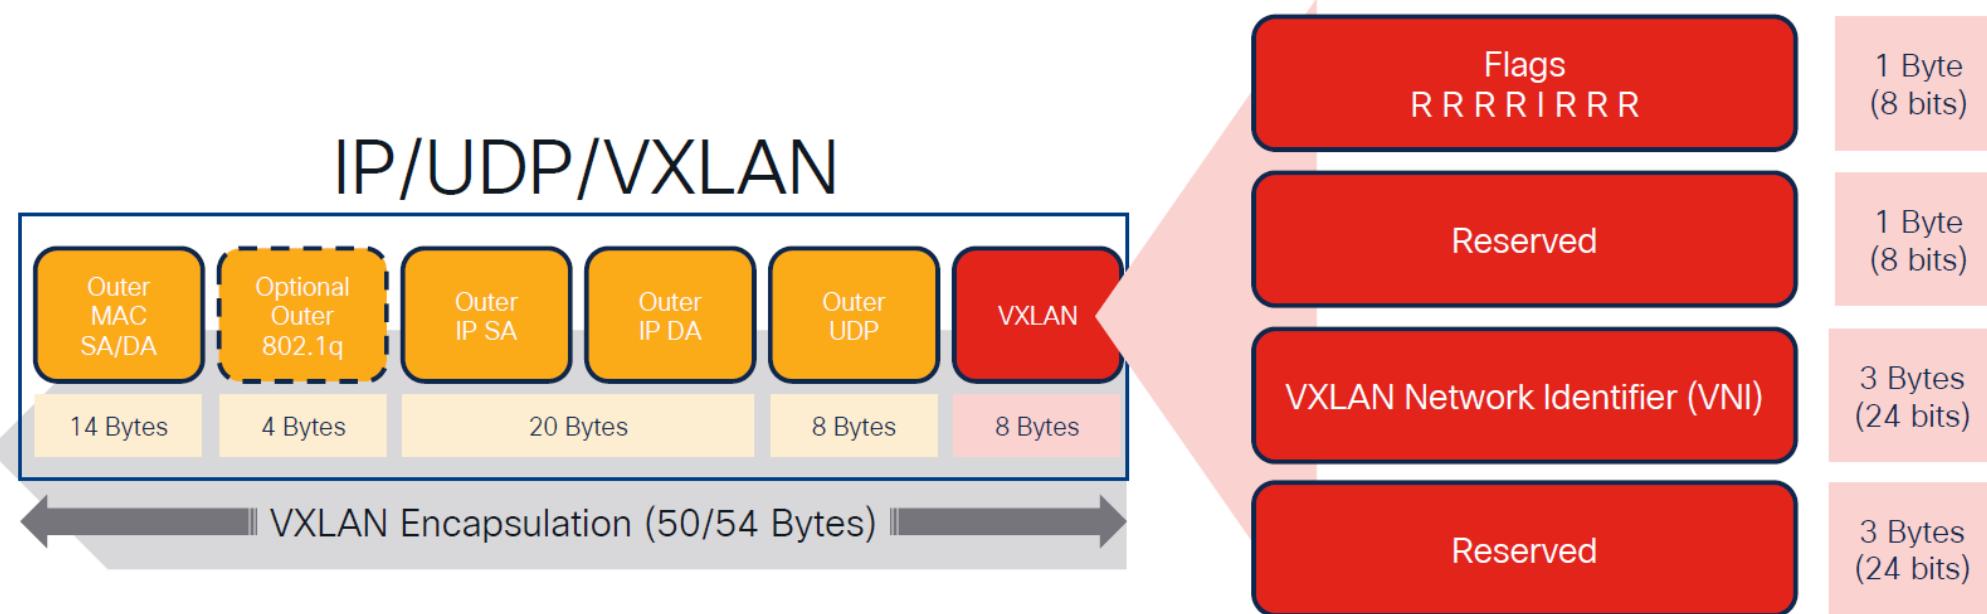
# VXLAN Packet Format

- VXLAN uses MAC in UDP encapsulation



- Adds UDP and VXLAN Header before original Ethernet Frame

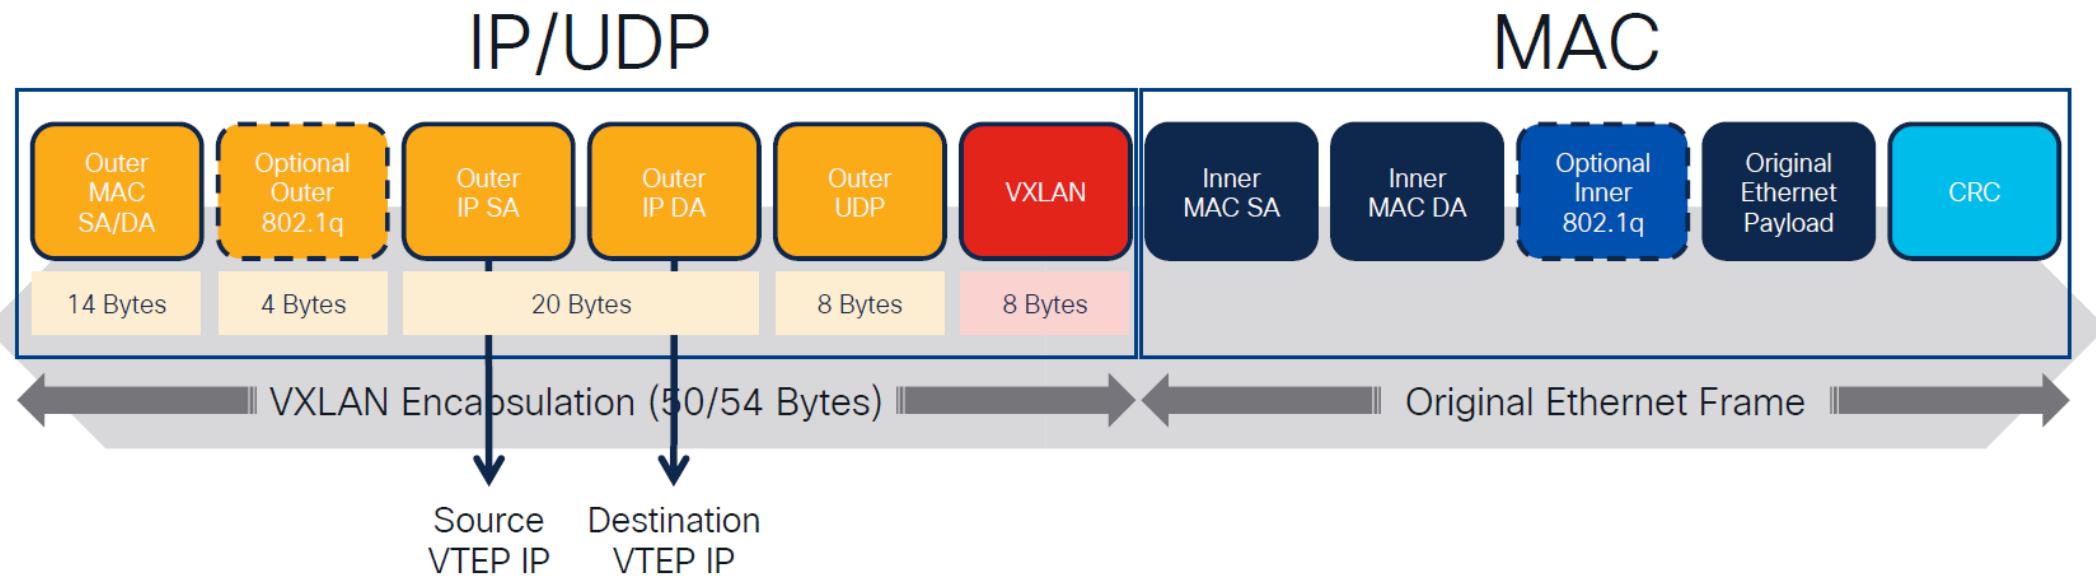
# VXLAN Header Details



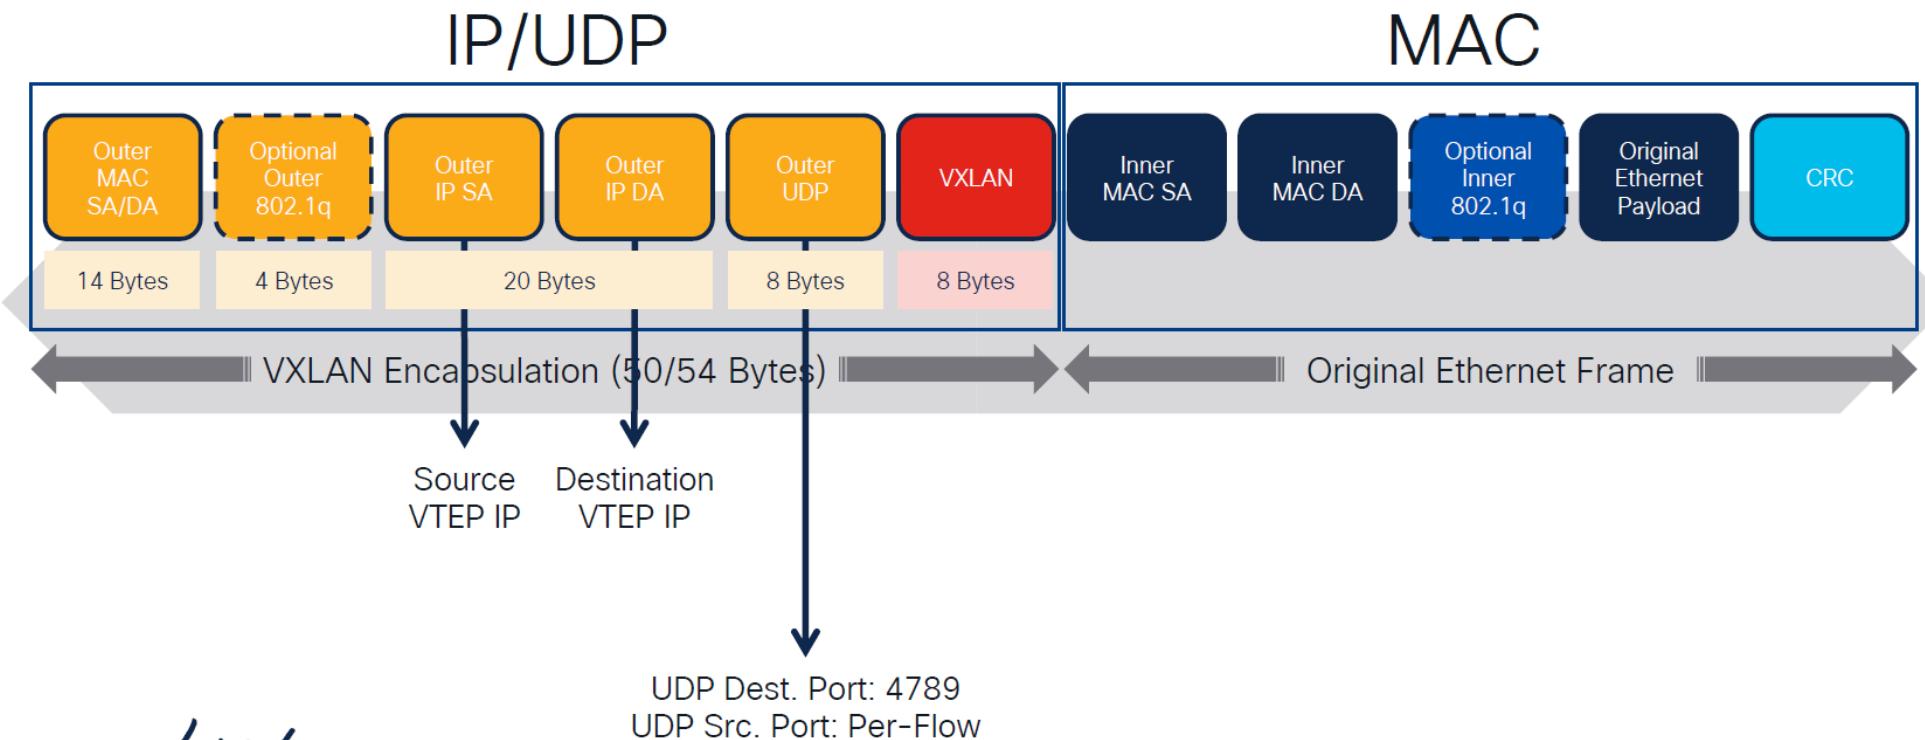
- Flags Field: I-flag (set to 1) for valid VNI. Other flags remain as R (set to 0)
- VNI Field: Allows VNI 1-16,777,215 (some implementation only 4096-16,777,215)

```
vlan 100  
vn-segment 50000
```

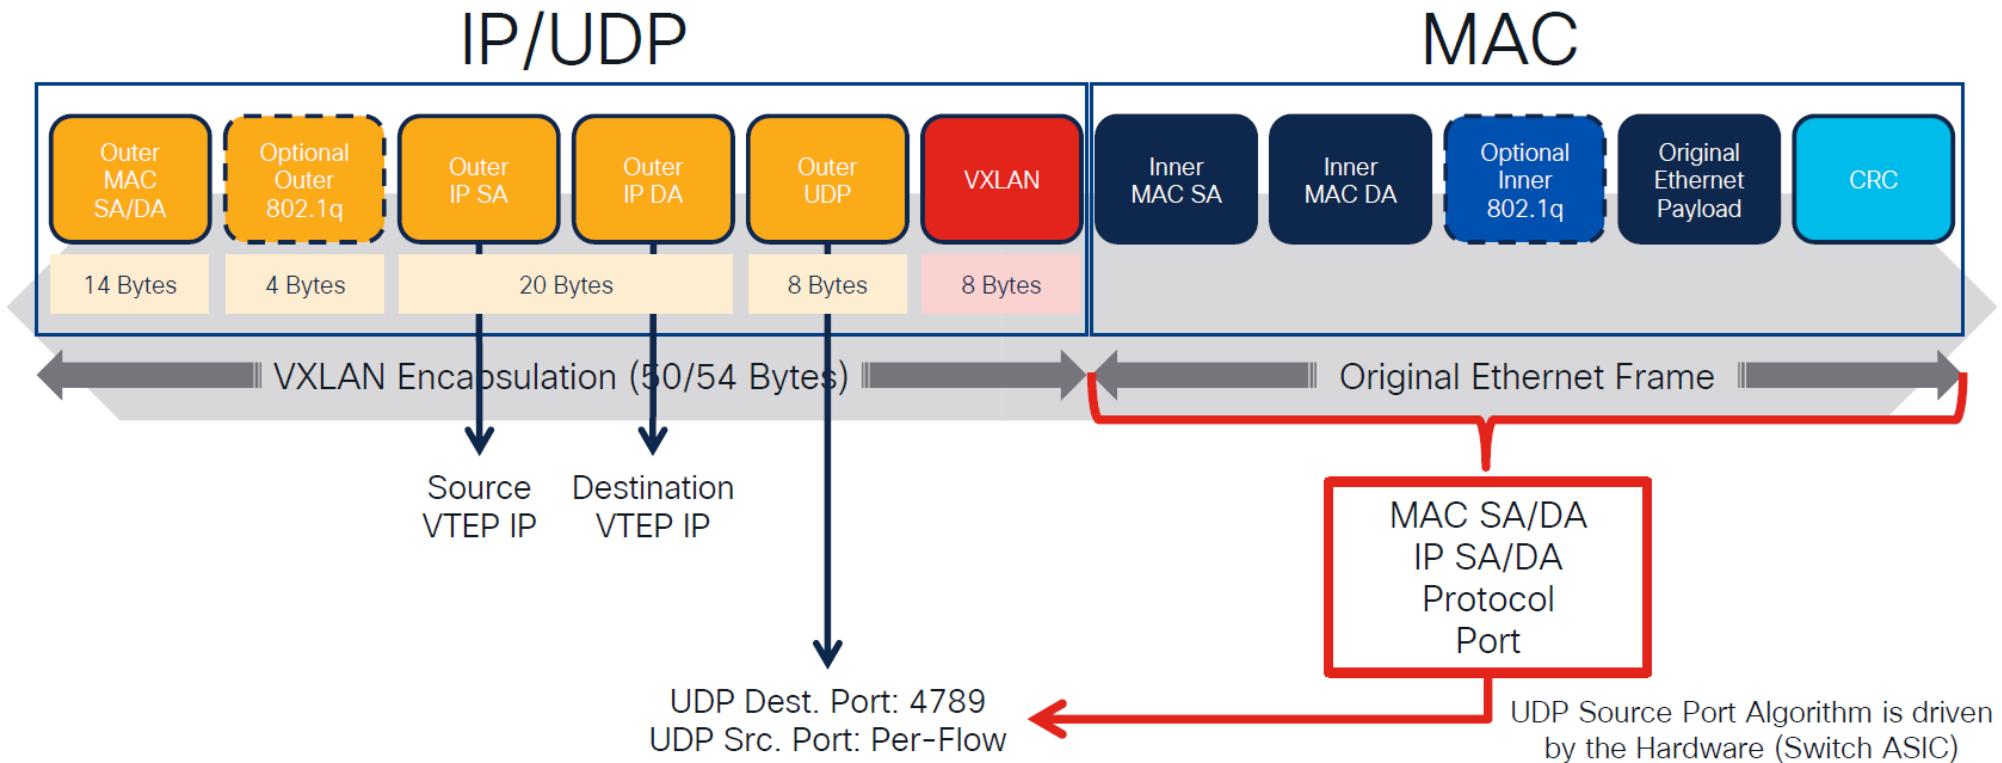
# Transport Independence



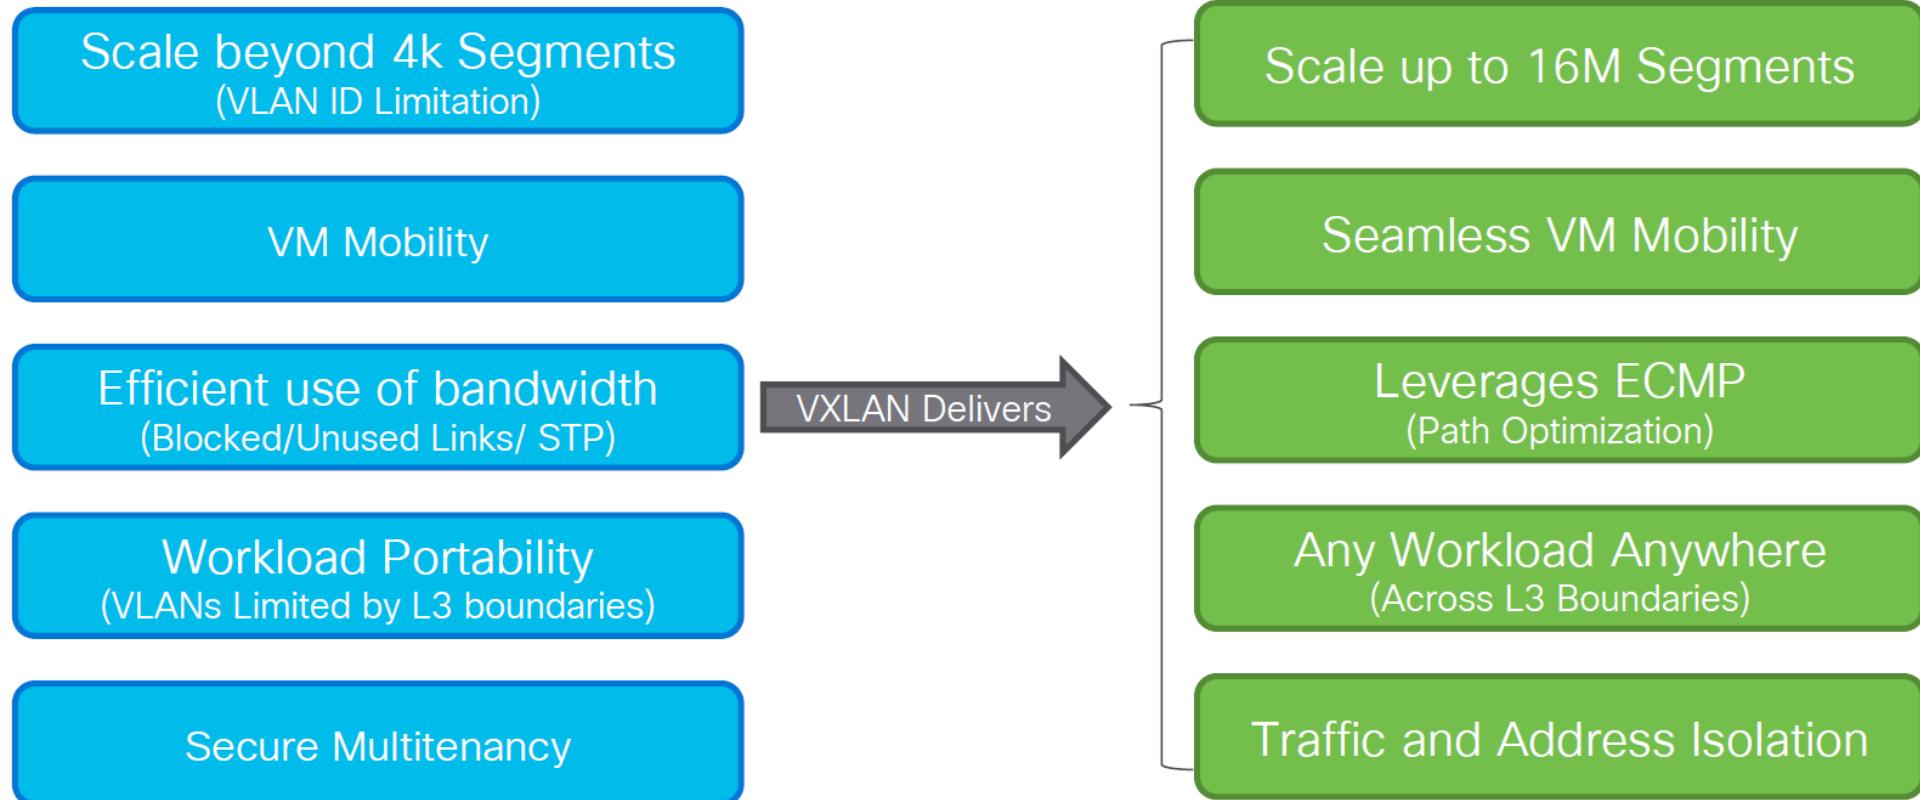
# Multipath Capable



# Multipath Capable

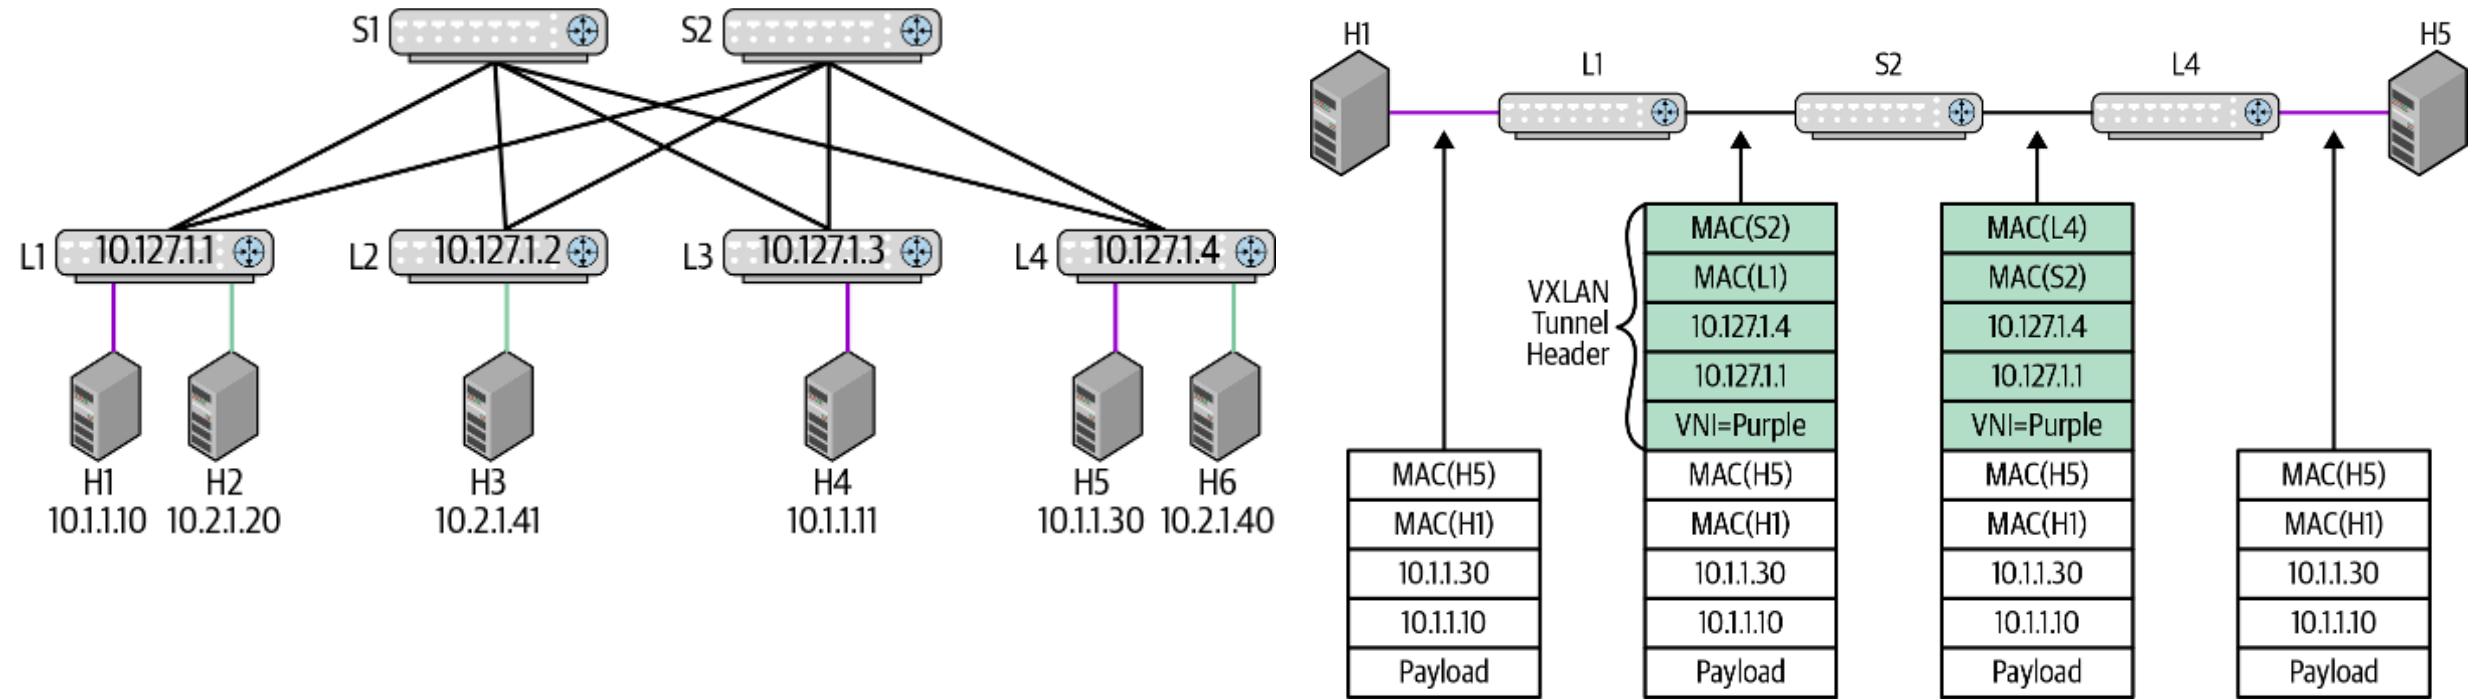


# Why VXLAN? - How did we get here?

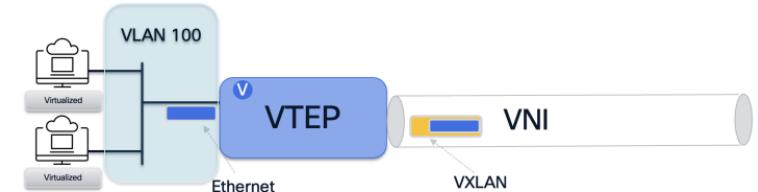


# VxLAN Bridging H1 to H5

- Same subnet forward



# Packet Capture - Host to VTEP



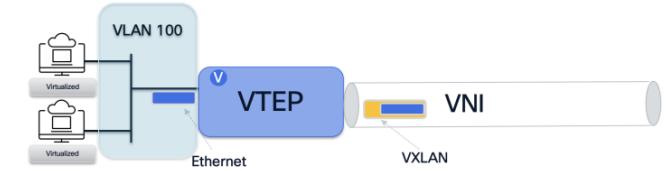
```
> Frame 2: 102 bytes on wire (816 bits), 102 bytes captured (816 bits)
  ✓ Ethernet II, Src: 52:0d:cf:cf:1b:08 (52:0d:cf:cf:1b:08), Dst: 20:20:00:00:aa (20:20:00:00:aa)
    > Destination: 20:20:00:00:aa (20:20:00:00:aa)
    > Source: 52:0d:cf:cf:1b:08 (52:0d:cf:cf:1b:08)
    Type: 802.1Q Virtual LAN (0x8100)
    [Stream index: 1]
  ✓ 802.1Q Virtual LAN, PRI: 0, DEI: 0, ID: 100
    000. .... .... = Priority: Best Effort (default) (0)
    ...0 .... .... = DEI: Ineligible
    .... 0000 0110 0100 = ID: 100 ← VLAN 100
    Type: IPv4 (0x0800)
  ✓ Internet Protocol Version 4, Src: 100.100.100.2, Dst: 200.200.200.2
    0100 .... = Version: 4
    .... 0101 = Header Length: 20 bytes (5)
    > Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
      Total Length: 84
      Identification: 0xbd56 (48470)
    > 000. .... = Flags: 0x0
      ...0 0000 0000 0000 = Fragment Offset: 0
      Time to Live: 255
      Protocol: ICMP (1)
      Header Checksum: 0xa520 [validation disabled]
      [Header checksum status: Unverified]
      Source Address: 100.100.100.2
      Destination Address: 200.200.200.2
      [Stream index: 0]
    Internet Control Message Protocol
      Type: 8 (Echo (ping) request)
      Code: 0
      Checksum: 0x15f4 [correct]
      [Checksum Status: Good]
      Identifier (BE): 13605 (0x3525)
      Identifier (LE): 9525 (0x2535)
      Sequence Number (BE): 52225 (0xcc01)
      Sequence Number (LE): 460 (0x01cc)
      [Response frame: 3]
      Timestamp from icmp data: Nov 7, 2024 08:14:13.468214000 IST
      [Timestamp from icmp data (relative): 58158.750333000 seconds]
    > Data (40 bytes)
```

Ethernet Header

IP Header

Payload

# Packet Capture - VTEP to Spine



```
> Frame 2: 148 bytes on wire (1184 bits), 148 bytes captured (1184 bits)
> Ethernet II, Src: 52:0d:18:e0:1b:08 (52:0d:18:e0:1b:08), Dst: 52:03:7e:03:1b:08 (52:03:7e:03:1b:08)
  Internet Protocol Version 4, Src: 10.3.0.1, Dst: 10.3.0.2
    0100 .... = Version: 4
    .... 0101 = Header Length: 20 bytes (5)
    Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
    Total Length: 134
    Identification: 0x0000 (0)
    000. .... = Flags: 0x0
    ...0 0000 0000 0000 = Fragment Offset: 0
    Time to Live: 255
    Protocol: UDP (17)
    Header Checksum: 0xa75e [validation disabled]
      [Header checksum status: Unverified]
    Source Address: 10.3.0.1
    Destination Address: 10.3.0.2
    [Stream index: 0]
    User Datagram Protocol, Src Port: 53936, Dst Port: 4789
    Virtual extensible Local Area Network
      Flags: 0x0800, VXLAN Network ID (VNI)
      Group Policy ID: 0
      VXLAN Network Identifier (VNI): 50000
      RESERVED: 0
    Ethernet II, Src: 52:0d:18:e0:1b:08 (52:0d:18:e0:1b:08), Dst: 52:11:5f:78:1b:08 (52:11:5f:78:1b:08)
      Destination: 52:11:5f:78:1b:08 (52:11:5f:78:1b:08)
      Source: 52:0d:18:e0:1b:08 (52:0d:18:e0:1b:08)
      Type: IPv4 (0x0800)
      [Stream index: 1]
    Internet Protocol Version 4, Src: 100.100.100.2, Dst: 200.200.200.2
      0100 .... = Version: 4
      .... 0101 = Header Length: 20 bytes (5)
      Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
      Total Length: 84
      Identification: 0xc086 (49286)
      000. .... = Flags: 0x0
      ...0 0000 0000 0000 = Fragment Offset: 0
      Time to Live: 254
      Protocol: ICMP (1)
      Header Checksum: 0xa2f0 [validation disabled]
        [Header checksum status: Unverified]
      Source Address: 100.100.100.2
      Destination Address: 200.200.200.2
      [Stream index: 1]
    Internet Control Message Protocol
```

IP Header

VTEP IP Addresses

UDP Header

VNI 50000

VXLAN Encapsulation

Payload

# VLAN vs VxLAN Setup

- **VLAN setup:**
- Create VLAN (e.g., VLAN 10)
- Assign switch ports to VLAN
- Configure trunk ports between switches
- Devices communicate within layer-2 domain

- **VXLAN setup:**
- Configure VTEP on each device/switch (assign IP address)
- Create VNI (e.g., VNI 5000)
- Map VNI to local VLAN
- Configure peer discovery (multicast group, static peers, or BGP EVPN)
- VXLAN tunnels form automatically between VTEPs

## Key differences:

- VLAN: Layer-2 only, limited to 4096 IDs, direct switch connectivity
- VXLAN: Layer-3 transport, 16M IDs, works across routed networks
- Creation order for VXLAN: VTEP → VNI → VNI-to-VLAN mapping → Peer discovery
- The VTEP is the foundational element - without it, VXLAN cannot function.

# Testing...

- [https://github.com/faysalmehedi/vxlan-ovs-docker-lab?fbclid=IwAR0OW7MTpkOc0nfHGeY5otzmPo\\_ng5YRQm0ITF6JsjbAbFAVAjf0Svz7sd4](https://github.com/faysalmehedi/vxlan-ovs-docker-lab?fbclid=IwAR0OW7MTpkOc0nfHGeY5otzmPo_ng5YRQm0ITF6JsjbAbFAVAjf0Svz7sd4)

# Docker Networks

## Network Drivers

Bridge

Host

None

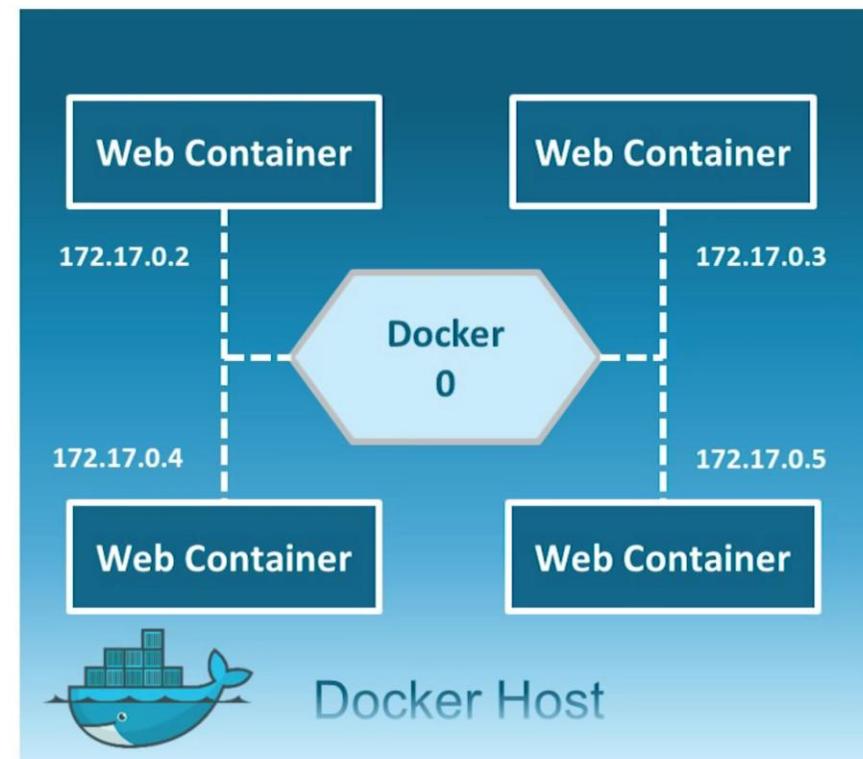
Macvlan

Overlay



# Network Drivers: Bridge

The default network driver. If you don't specify a driver, this is the type of network you are creating. Bridge networks are usually used when your applications run in standalone containers that need to communicate.



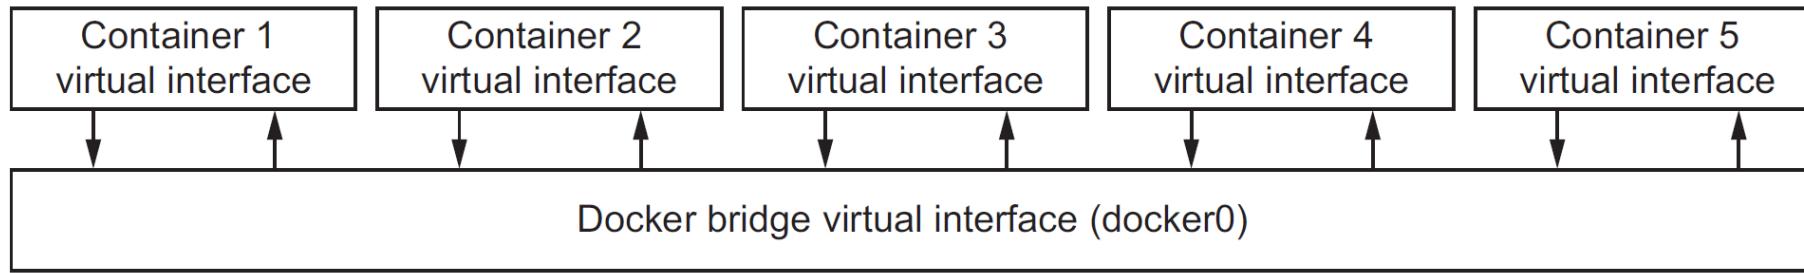
Learn **DevOps** from Industry Experts

edureka!

DevOps Certification Training

[www.edureka.co/devops](http://www.edureka.co/devops)

# Bridged Container: Inter-container communication



User can

- Define the address and subnet of the bridge
- Define the range of IP addresses that can be assigned to containers
- Define the maximum transmission unit (MTU)

# Docker network

```
modern@ubuntu:~$ docker network
```

```
Usage: docker network COMMAND
```

```
Manage networks
```

```
Commands:
```

connect	Connect a container to a network
create	Create a network
disconnect	Disconnect a container from a network
inspect	Display detailed information on one or more networks
ls	List networks
prune	Remove all unused networks
rm	Remove one or more networks

```
Run 'docker network COMMAND --help' for more information on a command.
```

# Check for Network installed (from install docker)

- Name
  - Just Name ..
- Scope
  - Local = single host
  - Swarm = multi-host

**'bridge'** is the default one where all new containers will connect to unless specified otherwise, and it maps to docker0 virtual bridge (the one you see when running ifconfig in the host)  
**'host'** maps the container directly to the host, so it is not recommended due to security concerns  
**'none'** provides no connectivity at all

```
root@ubuntu:/home/modern# docker network ls
NETWORK ID      NAME      DRIVER      SCOPE
03e9693f71df   bridge    bridge      local
51c819f4deb6   host      host       local
acf2bf319f06   none     null       local
root@ubuntu:/home/modern# █
```

# Inspect

```
root@ubuntu:/home/modern# docker network inspect bridge
```

```
[  
 {  
   "Name": "bridge",  
   "Id": "03e9693f71df12e37b7ca9c8606f8f3c22d44e9943ac95432751ffa94475e4c4"  
,  
   "Created": "2019-12-24T08:16:41.135287876-08:00",  
   "Scope": "local",  
   "Driver": "bridge",  
   "EnableIPv6": false,  
   "IPAM": {  
     "Driver": "default",  
     "Options": null,  
     "Config": [  
       {  
         "Subnet": "172.17.0.0/16",  
         "Gateway": "172.17.0.1"  
       }  
     ]  
   },  
   "Internal": false,  
   "Attachable": false,  
   "Ingress": false  
 },  
 {"  
   "Internal": false,  
   "Attachable": false,  
   "Ingress": false  
 }]
```

- No IPv6
- IPAM = IP Address Management
- Default subnet = 172.17.0.0 and gateway = 172.17.0.1

```
,  
   "Internal": false,  
   "Attachable": false,  
   "Ingress": false,  
   "ConfigFrom": {  
     "Network": ""  
   },  
   "ConfigOnly": false,  
   "Containers": {},  
   "Options": {  
     "com.docker.network.bridge.default_bridge": "true",  
     "com.docker.network.bridge.enable_icc": "true",  
     "com.docker.network.bridge.enable_ip_masquerade": "true",  
     "com.docker.network.bridge.host_binding_ipv4": "0.0.0.0",  
     "com.docker.network.bridge.name": "docker0",  
     "com.docker.network.driver.mtu": "1500"  
   },  
   "Labels": {}  
 }
```

No Container Connected Yet

# Create a bridge

- Bridge name ps-bridge

```
modern@ubuntu:~$ sudo docker network create -d bridge --subnet 10.0.0.1/24 ps-bridge  
[sudo] password for modern:  
913ba9ee873a476c82b39814dce1c8fa080ea8a1c73131c8c7bf72529a6d27fe
```

- Check it

```
modern@ubuntu:~$ sudo docker network ls  
NETWORK ID      NAME      DRIVER      SCOPE  
03e9693f71df    bridge    bridge      local  
51c819f4deb6    host      host       local  
acf2bf319f06    none     null       local  
913ba9ee873a    ps-bridge bridge      local  
modern@ubuntu:~$
```

# ตรวจสอบที่สร้างขึ้น

```
root@ubuntu:/home/modern# docker network inspect ps-bridge
[
    {
        "Name": "ps-bridge",
        "Id": "913ba9ee873a476c82b39814dce1c8fa080ea8a1c73131c8c7bf72529a6d27fe",
        "Created": "2019-12-24T08:51:58.481215154-08:00",
        "Scope": "local",
        "Driver": "bridge",
        "EnableIPv6": false,
        "IPAM": {
            "Driver": "default",
            "Options": {},
            "Config": [
                {
                    "Subnet": "10.0.0.1/24"
                }
            ]
        },
        "Internal": false,
        "Attachable": false,
        "Ingress": false,
        "ConfigFrom": {
            "Network": ""
        },
        "ConfigOnly": false,
        "Containers": {},
        "Options": {},
        "Labels": {}
    }
]
```

# Install bridge function module

- Install

```
root@ubuntu:/home/modern# apt-get install bridge-utils
Reading package lists... Done
Building dependency tree
Reading state information... Done
bridge-utils is already the newest version (1.5-15ubuntu1).
bridge-utils set to manually installed.
0 upgraded, 0 newly installed, 0 to remove and 128 not upgraded.
```

- Show

```
root@ubuntu:/home/modern# brctl show
bridge name      bridge id          STP enabled     interfaces
br-913ba9ee873a    8000.02420439398f    no
docker0          8000.02428309553e    no
root@ubuntu:/home/modern#
```

- The first one -> one we just create
- The second one -> relate to default bridge

# Check

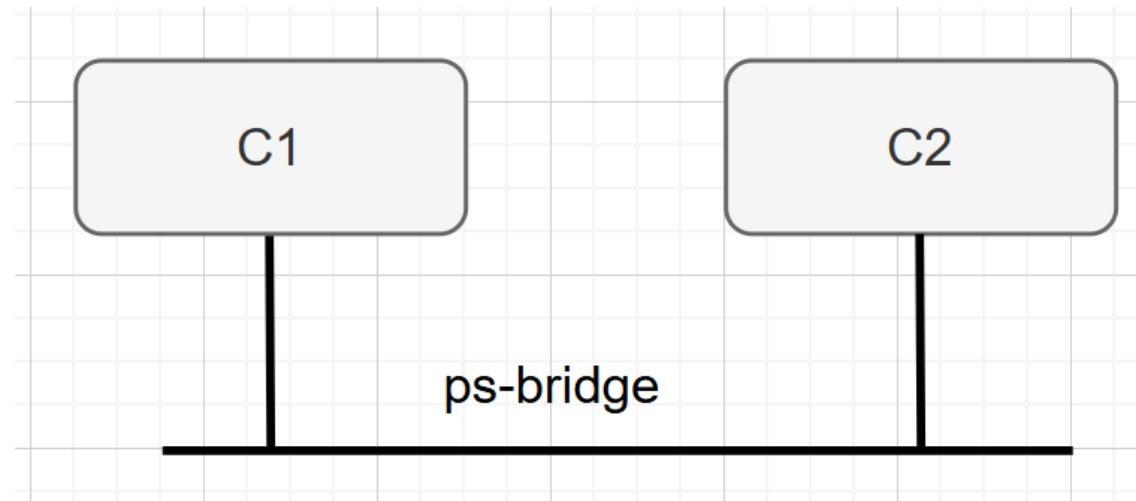
```
root@ubuntu:/home/modern# ip link list
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN mode DEFAULT group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: ens33: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP mode DEFAULT group default qlen 1000
    link/ether 00:0c:29:84:f3:ff brd ff:ff:ff:ff:ff:ff
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT group default
    link/ether 02:42:83:09:55:3e brd ff:ff:ff:ff:ff:ff
4: br-913ba9ee873a: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN mode DEFAULT group default
    link/ether 02:42:04:39:39:8f brd ff:ff:ff:ff:ff:ff
```

# Create container to connect

- C1, C2 container

```
root@ubuntu:/home/modern# docker run -dt --name c1 --network ps-bridge alpine sleep 1d
Unable to find image 'alpine:latest' locally
latest: Pulling from library/alpine
e6b0cf9c0882: Pull complete
Digest: sha256:2171658620155679240babee0a7714f6509fae66898db422ad803b951257db78
Status: Downloaded newer image for alpine:latest
859dab0048a0c8a4bae6c08d9b65d386868649775477dc18f6d3575e85165ba
root@ubuntu:/home/modern#
```

```
root@ubuntu:/home/modern# docker run -dt --name c2 --network ps-bridge alpine sleep 1d
c21ecdce9ad926eb6eda4c84cb7f66eb8ad2f0d72188a293008a7deab2e1ef38
```



# Inspect

- docker network inspect ps-bridge

```
"Containers": {  
    "859dab0048a0c8a4bae6c08d9b65d3868686849775477dc18f6d3575e85165ba": {  
        "Name": "c1",  
        "EndpointID": "ba183b618d72d649f0f4a22380a7ab8b5280cea8c2d67b64e74f874998a47e0a",  
        "MacAddress": "02:42:0a:00:00:02",  
        "IPv4Address": "10.0.0.2/24",  
        "IPv6Address": ""  
    },  
    "c21ecdce9ad926eb6eda4c84cb7f66eb8ad2f0d72188a293008a7deab2e1ef38": {  
        "Name": "c2",  
        "EndpointID": "d448349e9dfe0ea5279a3ac14e94565b8cace218df2412789aa96eb62ba36434",  
        "MacAddress": "02:42:0a:00:00:03",  
        "IPv4Address": "10.0.0.3/24",  
        "IPv6Address": ""  
    },  
},
```

# Bridge show

- Brctl show
  - Two virtual Ethernet (vethxxxx)

```
root@ubuntu:/home/modern# brctl show
bridge name      bridge id      STP enabled      interfaces
br-913ba9ee873a 8000.02420439398f    no          veth250bf5d
                                         vthe5aade0
docker0          8000.02428309553e    no
```

# Testing Connectivity

- Go inside container
- Check ip address
- Ping ->ip (can we ping by name ... testing ping c2)

```
root@ubuntu:/home/modern# docker exec -it c1 sh
/ # ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
5: eth0@if6: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
    link/ether 02:42:0a:00:02 brd ff:ff:ff:ff:ff:ff
    inet 10.0.0.2/24 brd 10.0.0.255 scope global eth0
        valid_lft forever preferred_lft forever
/ # ping 10.0.0.3
PING 10.0.0.3 (10.0.0.3): 56 data bytes
64 bytes from 10.0.0.3: seq=0 ttl=64 time=3.816 ms
64 bytes from 10.0.0.3: seq=1 ttl=64 time=0.085 ms
64 bytes from 10.0.0.3: seq=2 ttl=64 time=0.084 ms
^C
--- 10.0.0.3 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.084/1.328/3.816 ms
```

# FRRouting

- FRR is free software that implements and manages various IPv4 and IPv6 routing protocols. It runs on nearly all distributions of Linux and BSD and supports all modern CPU architectures.
- FRR currently supports the following protocols: BGP, OSPF, RIP, IS-IS, PIM, LDP, BFD, Babel, PBR, OpenFabric and VRRP, with alpha support for EIGRP and NHRP.



cloudscale.ch



LINDITEC  
creating value in network distribution



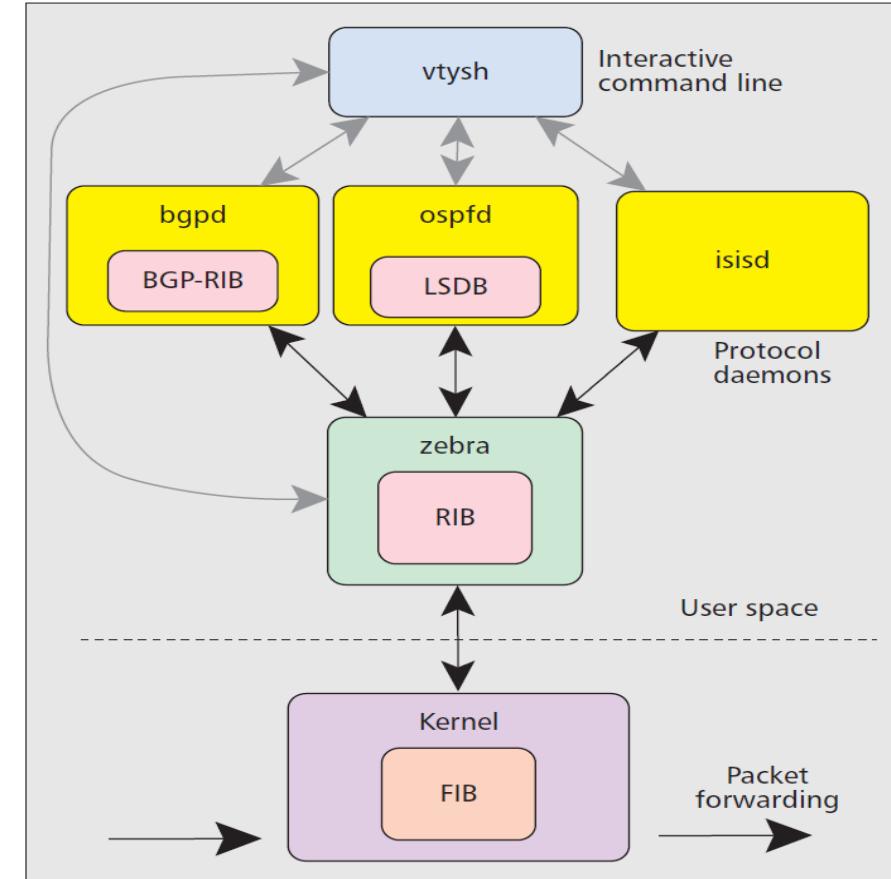
vmware®



# History of FRR

- ~1996 - Zebra development begins
- ~2002 - Quagga forked from Zebra
- 2016 - FRR forked from Quagga
- Jan. 2017 - First Release - FRR 2.0
- June 11, 2018 - Latest Stable - FRR 5.0

- FRR has its roots in the Quagga project.



Quagga architecture

Introduction to the Quagga Routing Suite, IEEE Network 2014

# Releases

- <https://frrouting.org/release/>
  - Latest Nov 9, 2025
- <https://docs.frrouting.org/en/latest/>

# What kernel do you need?

- VRF - Depends on l3mdev
  - 4.4 for Limited Functionality, 4.8 for full functionality
- BGP EVPN - Depends on NTF\_EXT\_LEARNED and ARP Suppression
  - 4.17 and 4.14
- PIM-SM - Depends on IGMPMSG\_WRVIFWHOLE and RTN\_MULTICAST netlink messages
  - 4.18
- These features are not currently available for non-Linux distributions

# Starting/Stopping FRR

- `systemctl <start|restart|reload|stop> frr`
- Config file: `/etc/frr/frr.conf`
- The interactive shell
  - `vtysh`

```
root@r1:/home/cumulus# systemctl start frr
root@r1:/home/cumulus# vtysh

Hello, this is FRRouting (version 4.0+cl3u2).
Copyright 1996-2005 Kunihiro Ishiguro, et al.

r1.rdu.cumulusnetworks.com# show ip route
Codes: K - kernel route, C - connected, S - static, R - RIP,
       0 - OSPF, I - IS-IS, B - BGP, E - EIGRP, N - NHRP,
       T - Table, v - VNC, V - VNC-Direct, A - Babel, D -
SHARP,
       F - PBR,
       > - selected route, * - FIB route

K>* 0.0.0.0/0 [0/0] via 10.50.11.1, eth0, 00:00:06
C>* 10.50.11.0/24 is directly connected, eth0, 00:00:06
C>* 192.168.210.0/24 is directly connected, swp1, 00:00:06
C>* 192.168.213.0/24 is directly connected, swp2, 00:00:06
C>* 192.168.214.0/24 is directly connected, swp3, 00:00:06
C>* 192.168.240.1/32 is directly connected, lo, 00:00:06
```

# Logging

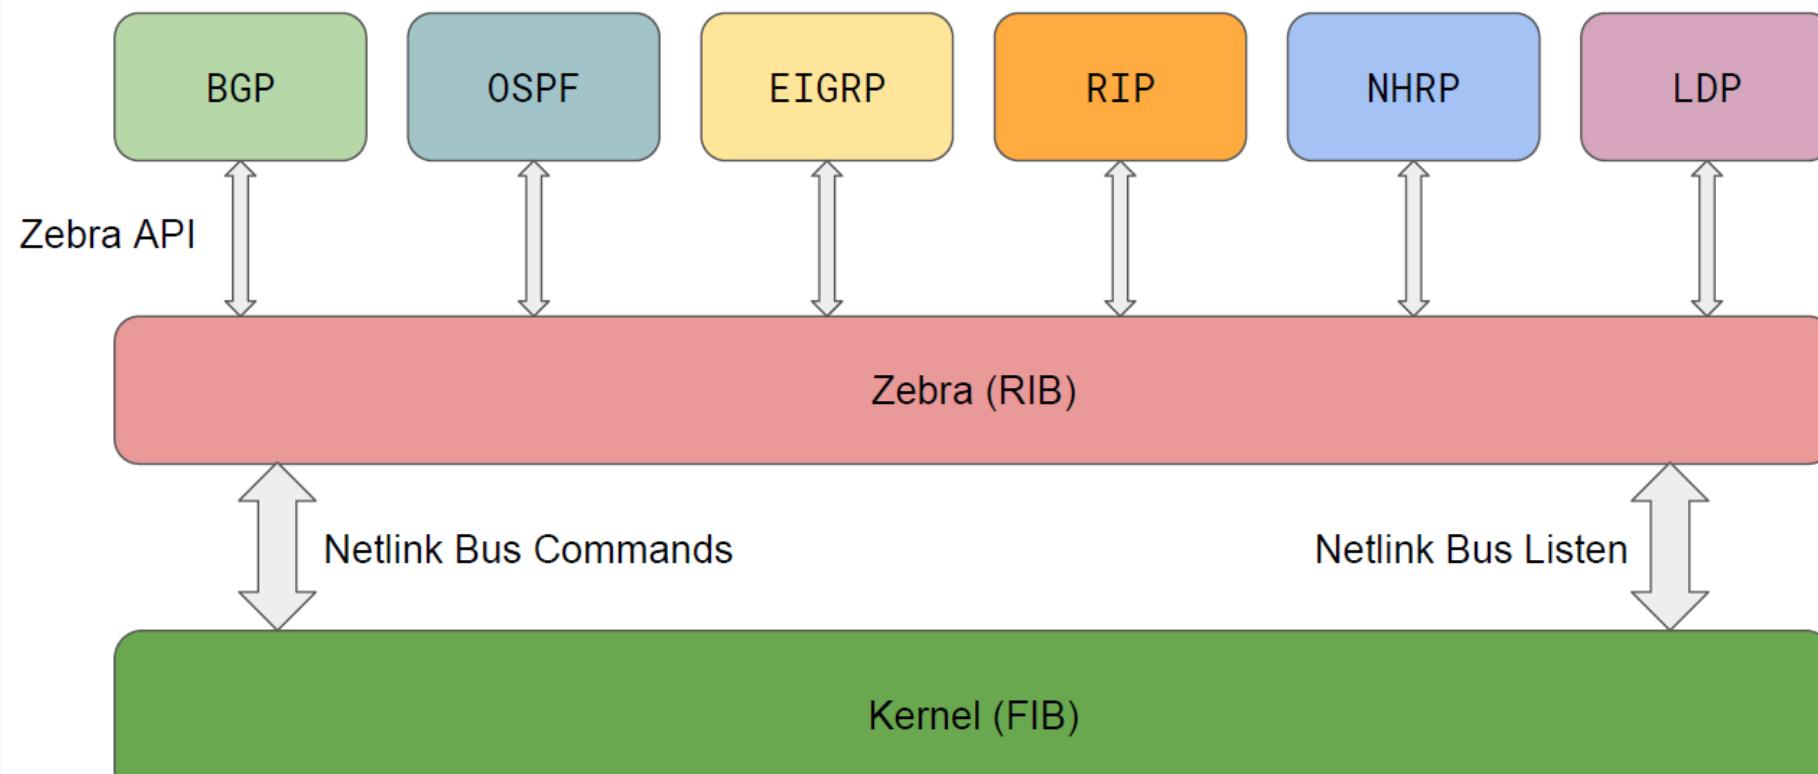
- Where do my logs go?
  - **show logging** in vtysh
  - 3 supported logging destinations:
    1. Syslog - **log syslog [level]**
    2. File - **log file [name]**
    3. Standard output - **log stdout**
- The log file has no output
- Designed to not be chatty unless debugs are turned on
- If you are not sure what is going on use debug commands liberally
- **debug [protocol] ?** if you are not sure
- what to turn on.

```
frr(config)# log file /tmp/frr.log
frr(config)# log syslog
frr(config)# log stdout
frr(config)#
frr# show logging

Logging configuration for zebra:
Syslog logging: level debugging, facility
daemon, ident zebra
Stdout logging: level debugging
Monitor logging: level debugging
File logging: disabled
Protocol name: ZEBRA
Record priority: disabled
Timestamp precision: 6

...
```

# FRR Architecture



# RIB -vs- FIB

- Routing Information Base
  - A.k.a ‘Control Plane’
  - This is in FRRouting
- Forwarding Information Base
  - A.k.a ‘Data plane’
  - This is in the Kernel

## FRR

```
robot# show ip route
Codes: K - kernel route, C - connected, S - static, R - RIP,
       O - OSPF, I - IS-IS, B - BGP, E - EIGRP, N - NHRP,
       T - Table, v - VNC, V - VNC-Direct, A - Babel, D - SHARP,
       F - PBR,
       > - selected route, * - FIB route

K * 0.0.0.0/0 [0/100] via 10.0.2.2, enp0s3 inactive, 01:24:15
D  4.3.2.1/32 [150/0] via 192.168.209.44, enp0s9, 00:00:02
K>* 4.3.2.1/32 [0/14000] is directly connected, enp0s8, 00:01:05
S  4.3.2.1/32 [1/0] is directly connected, enp0s10, 16:30:20
K * 4.3.2.1/32 [255/8192] is directly connected, enp0s9, 16:33:31
```

## Kernel

```
sharpd@robot ~> ip route show
default via 10.0.2.2 dev enp0s3 proto static metric 100
4.3.2.1 dev enp0s8 scope link metric 14000
4.3.2.1 dev enp0s9 scope link metric 4278198272
```

# OSPF Overview

- Considered an IGP (Interior Gateway Protocol)
  - Uses Link State Routing
  - Every router floods information, so that every router has all the information
  - Dijkstra's Algorithm
- Uses IP (Layer 3) to advertise routing information
- OSPFv2 -vs- OSPFv3
  - IPv4 -vs- IPv6
- Supports multiple areas (and types of areas)
  - If you do more than one area you must have area 0

# OSPF Basic Setup

```
# Put this in all FRR.conf
```

```
!
```

```
router ospf
```

```
network 0.0.0.0/0 area 0.0.0.0
```

```
!
```

# OSPF Debugging

- `show ip ospf interface`
  - What Interfaces is OSPF using?
- `show ip ospf neighbor`
  - What neighbors have I formed?
- `show ip ospf database`
  - Dump of the OSPF database
- `show ip ospf database router`
  - Dump of what each router sent you
- `debug ospf packet`
  - Look at data sent/received from peers, stored in log file
- `debug ospf nsm`
  - Look at neighbor state change events, stored in log file

# Create Linux Router

- ดาวน์โหลดอิมเมจ Alpine

```
docker pull alpine:latest
```

- Create 2 router containers: สร้างเร้าเตอร์จากนั้นกำหนดค่าต่าง ๆ ด้วยการสร้างเน็ตเวิร์กที่ชื่อว่า routernet และสร้าง router1 และ router2

```
docker network create --subnet=172.20.0.0/16 routernet
```

```
docker run -dit --name router1 --hostname router1 \
--network routernet --ip 172.20.0.2 \
--cap-add=NET_ADMIN --cap-add=NET_RAW \
alpine:latest
```

```
docker run -dit --name router2 --hostname router2 \
--network routernet --ip 172.20.0.3 \
--cap-add=NET_ADMIN --cap-add=NET_RAW \
alpine:latest
```

3. ติดตั้ง FRR ให้กับ router1 และ router2 ด้วยการเข้าไปใน shell ของเร้าเตอร์ทั้ง 2 ตัว

```
docker exec -it router1 sh  
  
apk update  
apk add frr frr-openrc ip6tables iptables
```

4. ติดตั้ง ospf ตั้งค่า /etc/frr/daemons เป็นรองรับการทำงานของ OSPF และ zebra (สามารถทำผ่าน nano) เพื่อกำหนดให้ค่า ospfd=yes และ zebra=yes

```
# Edit /etc/frr/daemons  
sed -i s/ospfd=no/ospfd=yes/ /etc/frr/daemons  
sed -i s/zebra=no/zebra=yes/ /etc/frr/daemons  
  
# Enable IP forwarding  
echo "net.ipv4.ip_forward=1" >> /etc/sysctl.conf  
sysctl -p
```

5. ติดตั้ง OSPF:

```
# Start FRR
/usr/lib/frr/frrinit.sh start

# Configure via vtysh
vtysh

configure terminal
router ospf
  network 172.20.0.0/16 area 0
exit
exit
write memory
```

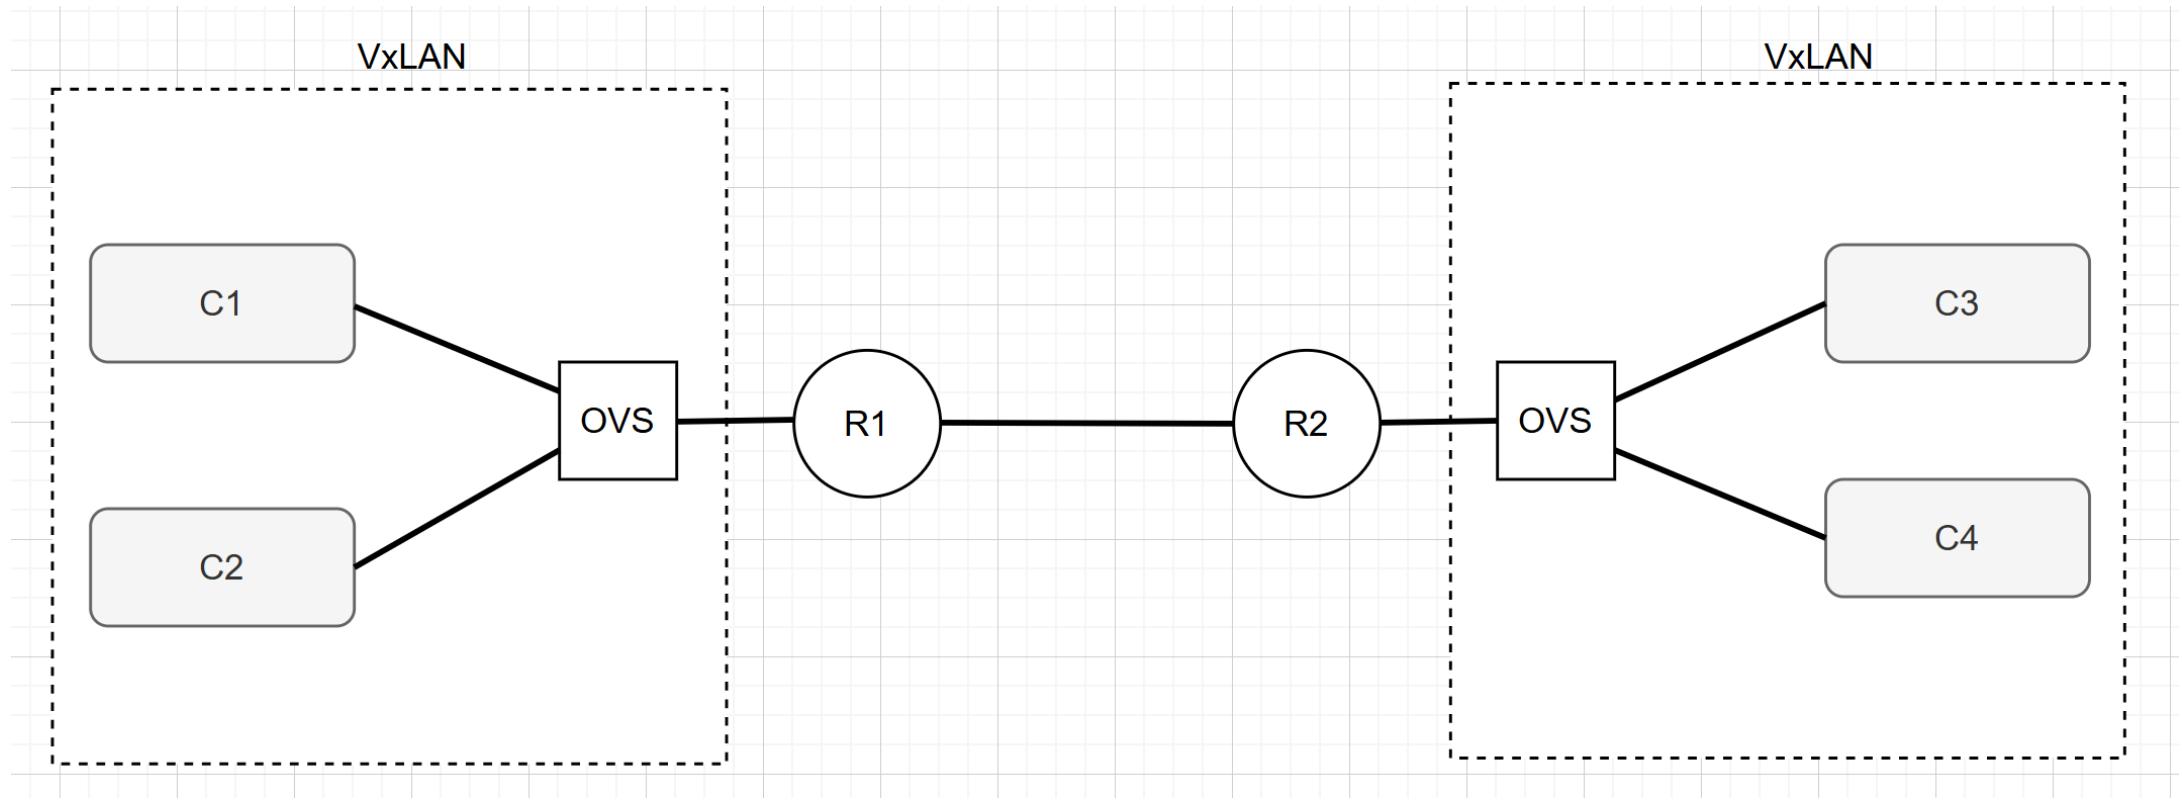
6. Verify:

```
# Check OSPF neighbors
vtysh -c "show ip ospf neighbor"

# Check routes
vtysh -c "show ip route"
```

กลับไปทำขั้นตอนที่ 3 - 5 ใหม่ที่ router 2 เพื่อให้ router ทั้งสองสามารถสื่อสารกันได้

# Create this network



สร้าง Tolopogy hosts -- OVS -- routers -- OVS -- hosts.

```
# Install OVS on host
apk add openvswitch
rc-service ovs-modules start
rc-service ovsdb-server start
rc-service ovs-vswitchd start

# Create 2 OVS bridges
ovs-vsctl add-br ovs1
ovs-vsctl add-br ovs2

# Create router containers with multiple networks
docker run -dit --name router1 --hostname router1 \
    --cap-add=NET_ADMIN --cap-add=NET_RAW --network none alpine:latest

docker run -dit --name router2 --hostname router2 \
    --cap-add=NET_ADMIN --cap-add=NET_RAW --network none alpine:latest

# Create host containers
docker run -dit --name host1 --network none alpine:latest
docker run -dit --name host2 --network none alpine:latest

# Connect using ovs-docker (utility script)
ovs-docker add-port ovs1 eth0 host1 --ipaddress=10.1.1.10/24
ovs-docker add-port ovs1 eth0 router1 --ipaddress=10.1.1.1/24
ovs-docker add-port ovs2 eth0 host2 --ipaddress=10.2.2.10/24
ovs-docker add-port ovs2 eth0 router2 --ipaddress=10.2.2.1/24

# Create router interconnect network
docker network create --subnet=192.168.0.0/30 router-link
docker network connect router-link router1 --ip=192.168.0.1
docker network connect router-link router2 --ip=192.168.0.2
```

# จาก Topology เดิม

- ทดสอบการสื่อสาร และสร้าง VxLAN ต่อไป