BGP

Mix of slides from: BRKENT-1179, APNIC, Kurose, Wu

Dynamic Routing Protocol

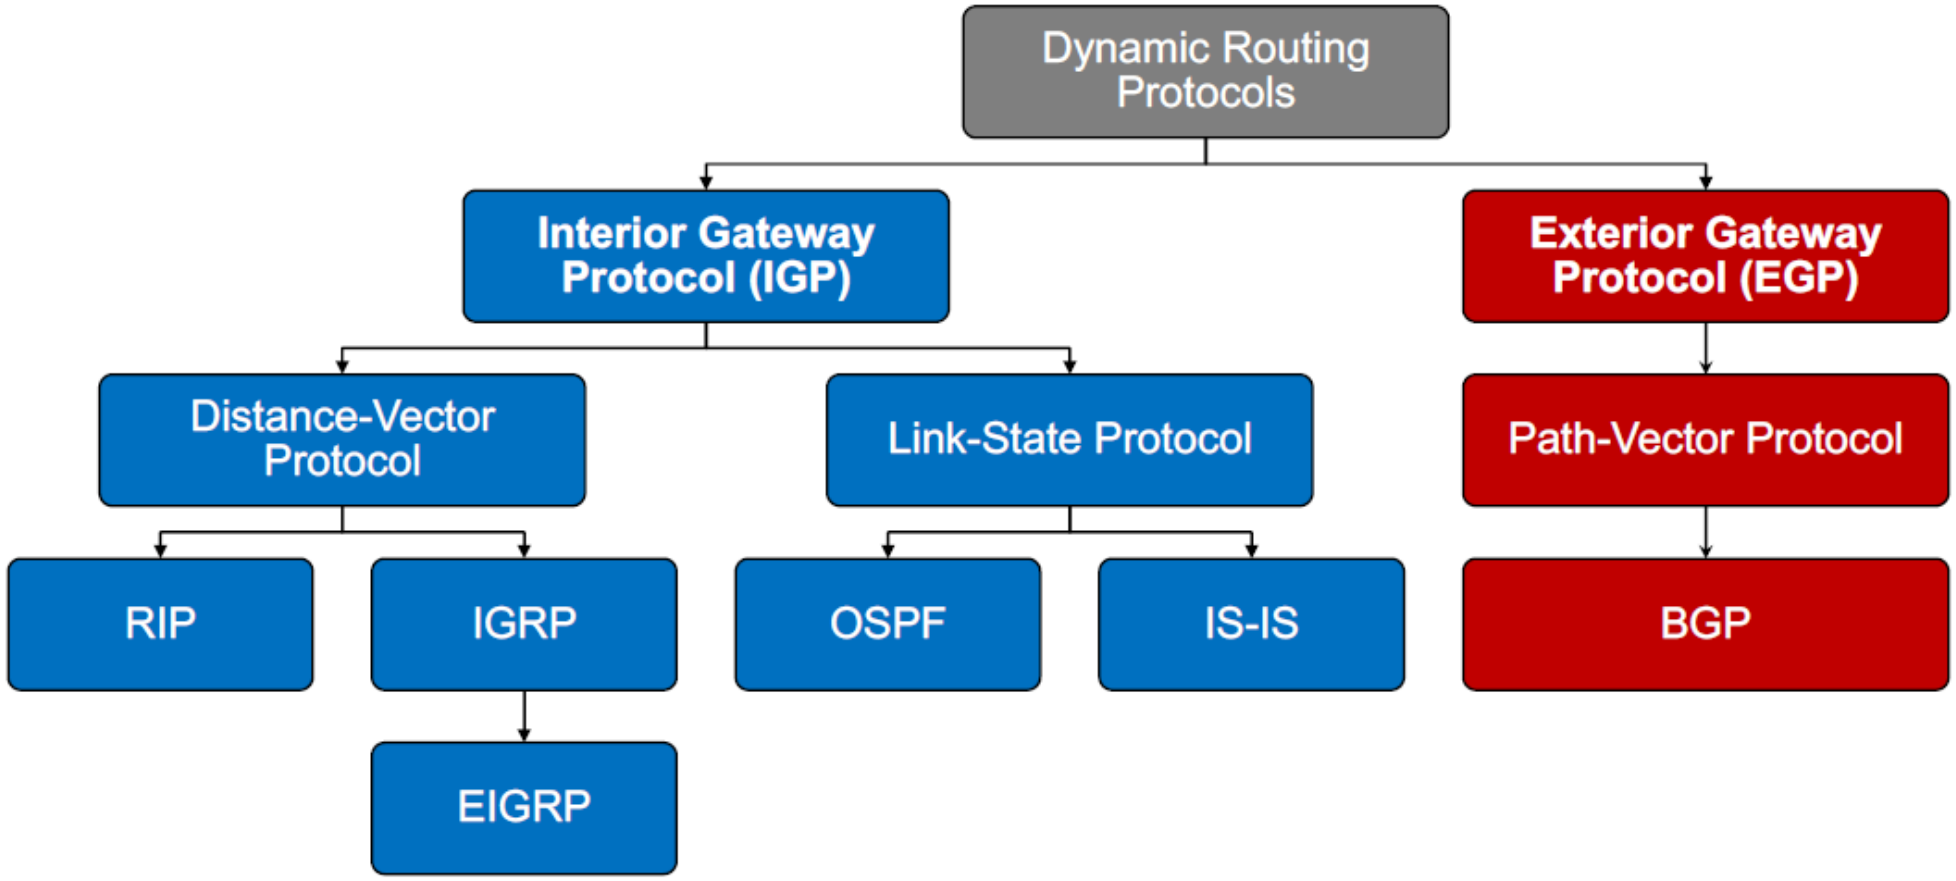

Autonomous Systems & Peering

Autonomous System

- A group of one or more IP prefixes (lists of IP addresses accessible on a network) run by one or more network operators that maintain a single, clearly-defined routing policy.

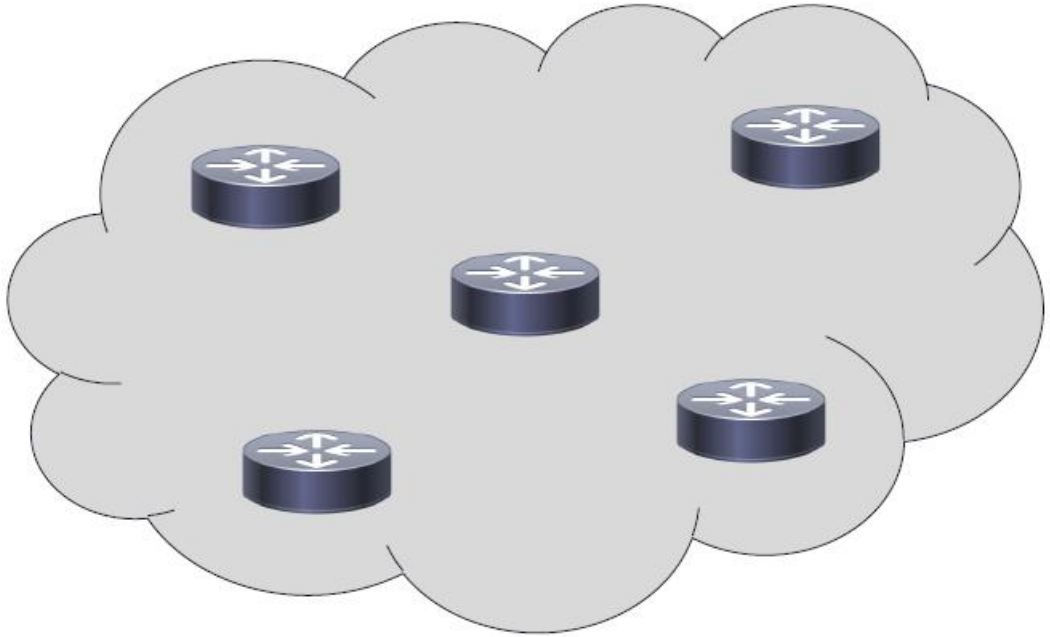

Autonomous System (AS)

- Types
  - Stub AS
  - Transit AS
- AS Number
  - Original scheme: 16 bits (BGP 2, RFC1105 (1989))
    - 0 to 65535 Private range 64512 to 65534
  - Extended number: 32 bits (BGP 4, RFC4893 (2007))
    - 0 to 4294967295 Additional private range 4200000000 to 4294967294

Stub AS

End customers; do not provide transit to others

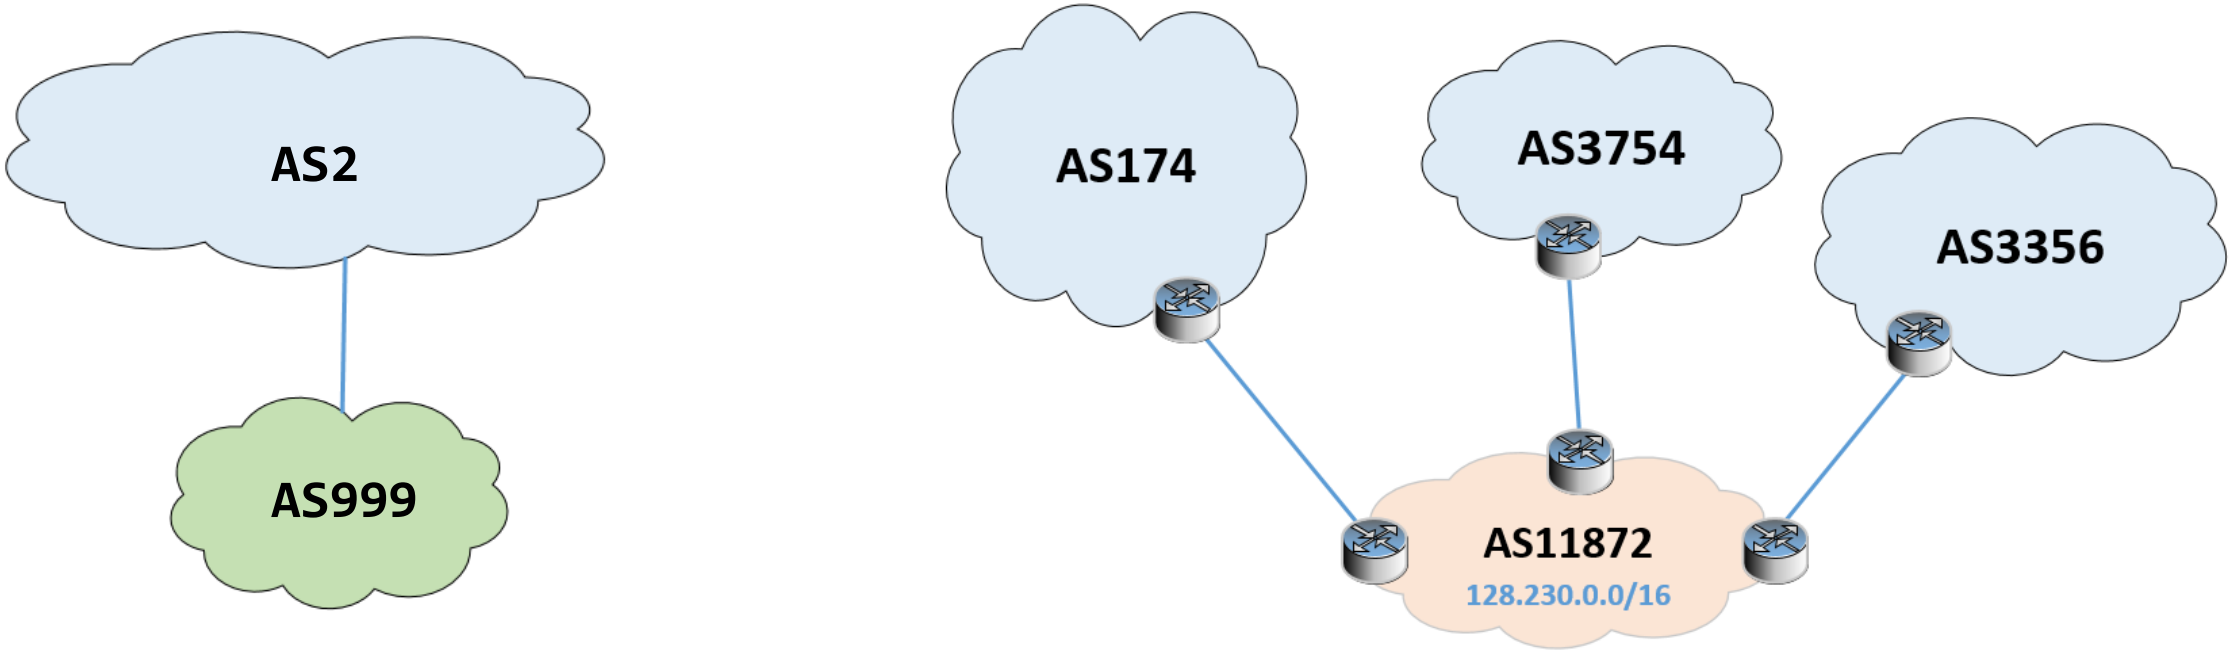

Transit AS

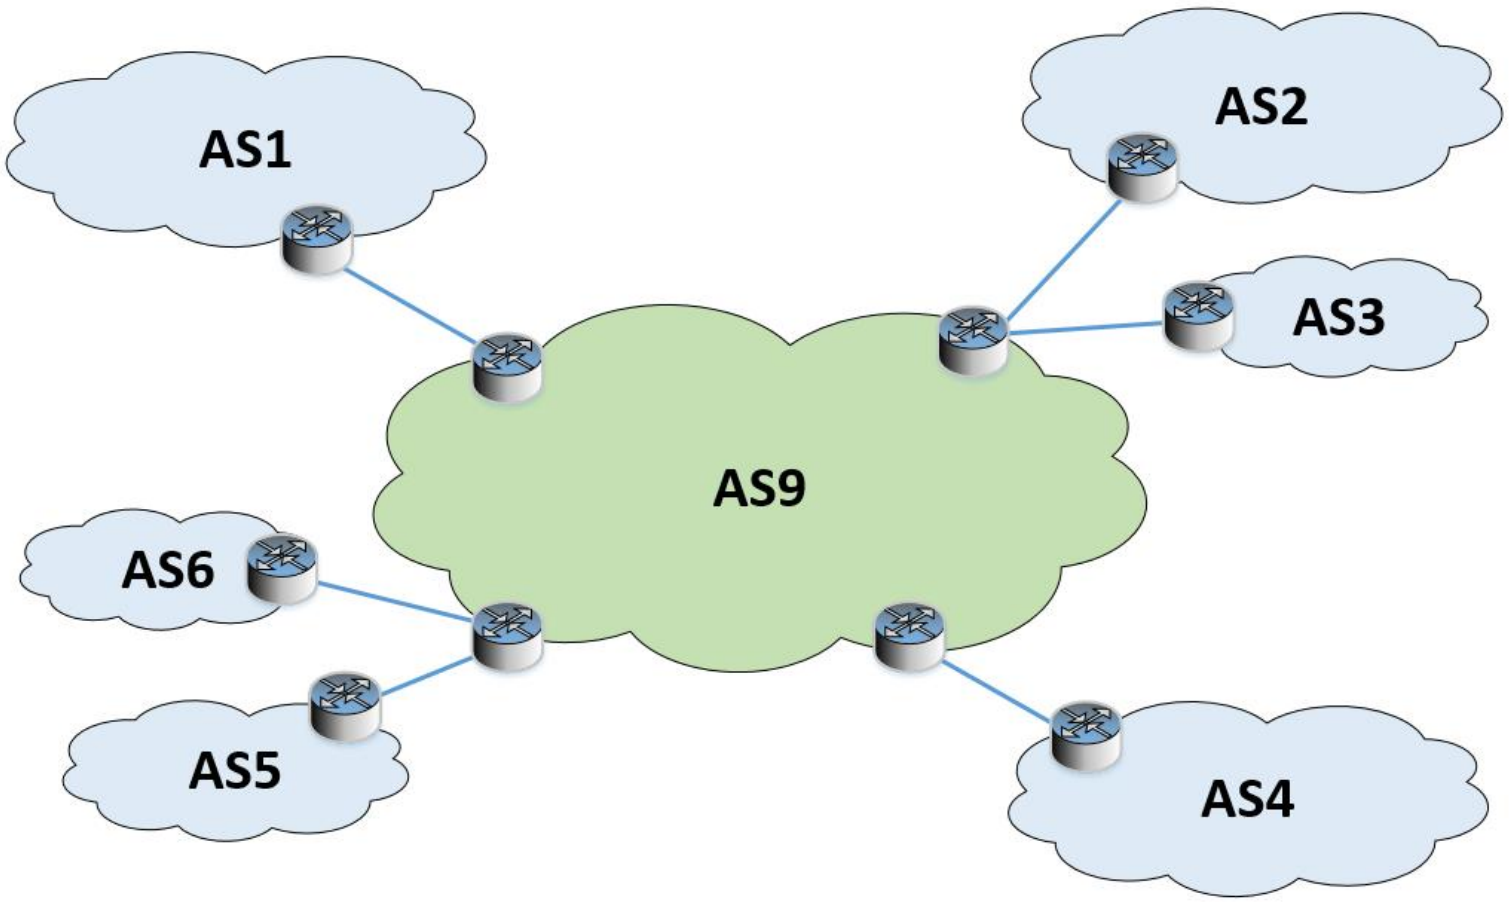

Provide transit to others

Peering

- Autonomous Systems peer with one another in an Internet Exchange or a data center

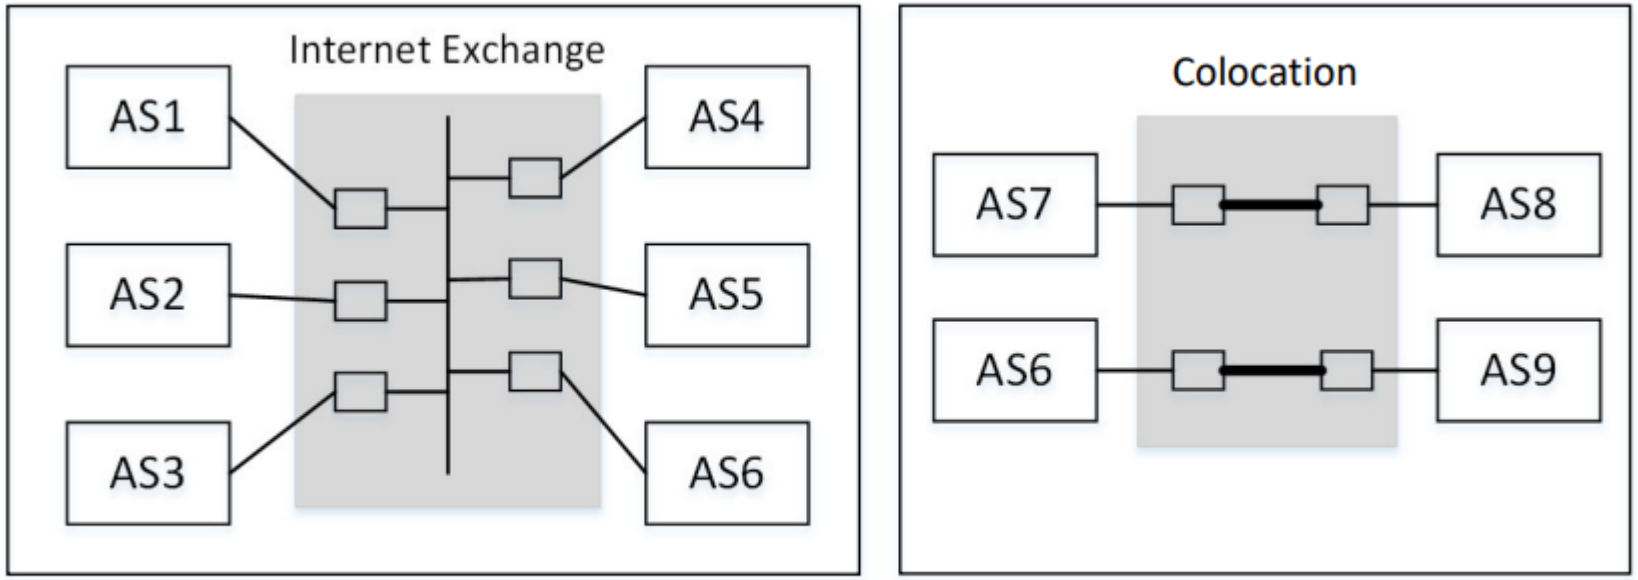

(A) Public peering

(B) Private peering

ASN and IP

- Find ASN from IP Address

```
$ whois -h whois.radb.net -- '-i origin AS11872' | grep route:
route:      128.230.0.0/16
route:      149.119.0.0/16
route:      128.230.0.0/17
route:      128.230.128.0/17
```

- Find IP prefix from ASN

```
$ whois -h whois.radb.net 31.13.78.3
route:      31.13.78.0/24
descr:      Facebook, Inc.
origin:     AS32934
mnt-by:     MAINT-AS32934
changed:    shaw@fb.com 20120423  #20:09:37Z
source:     RADB
```

Border Gateway Protocol

- Border Gateway Protocol Large scale, robust and stable routing protocol designed to operate between autonomous systems
- Based on TCP, listens on port 179
- Fundamentally a distance vector protocol
- Does not have the concept of a simple metric
- Strong control over advertised routes and their attributes

BGP Versions

- BGP was first described in 1989 and has been in use on the Internet since 1994
- There are four versions of BGP:

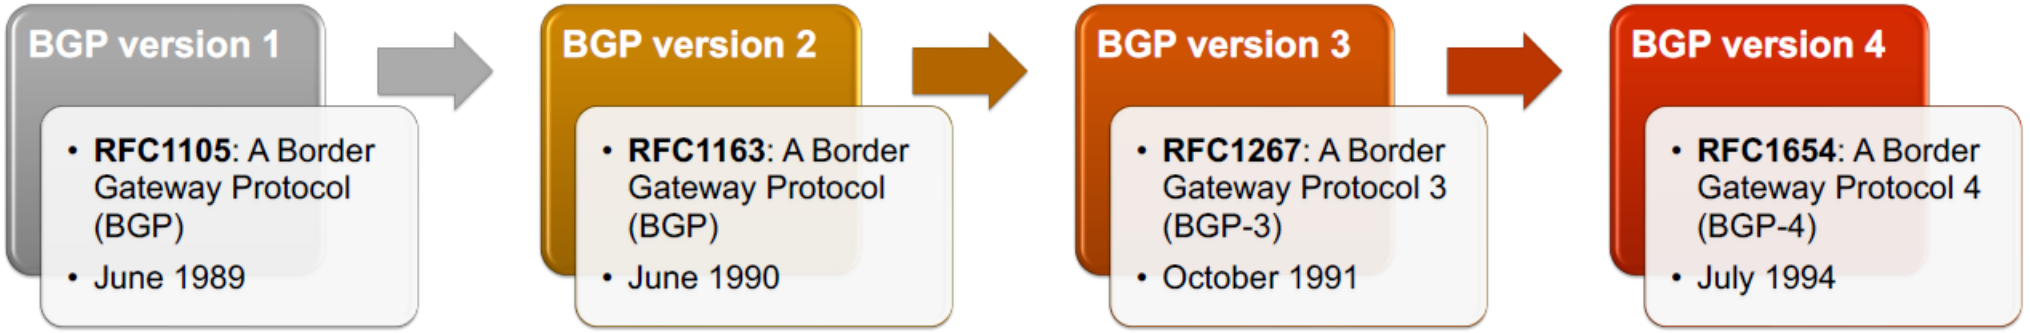

BGP Stability Considerations

- Network events burst.
- Fast reaction: better convergence, more churn.
- Delayed reaction: more stability, slower convergence.

- BGP prioritizes stability.
- Delays updates to reduce churn, combines changes.
- Sends only incremental updates.

- "Churn" reflects the percentage of customers who discontinue their use of a business's products or services over a certain period of time.

How BGP Works

Internet inter-AS routing: BGP

- BGP provides each AS a means to:
  - eBGP (External BGP): for BGPs from different AS
    - obtain subnet reachability information from neighboring ASes
  - iBGP (Internal BGP): for BGPs from the same AS
    - propagate reachability information to all AS-internal routers.

  - determine “good” routes to other networks based on reachability information and *policy*

IP Prefix Announcement

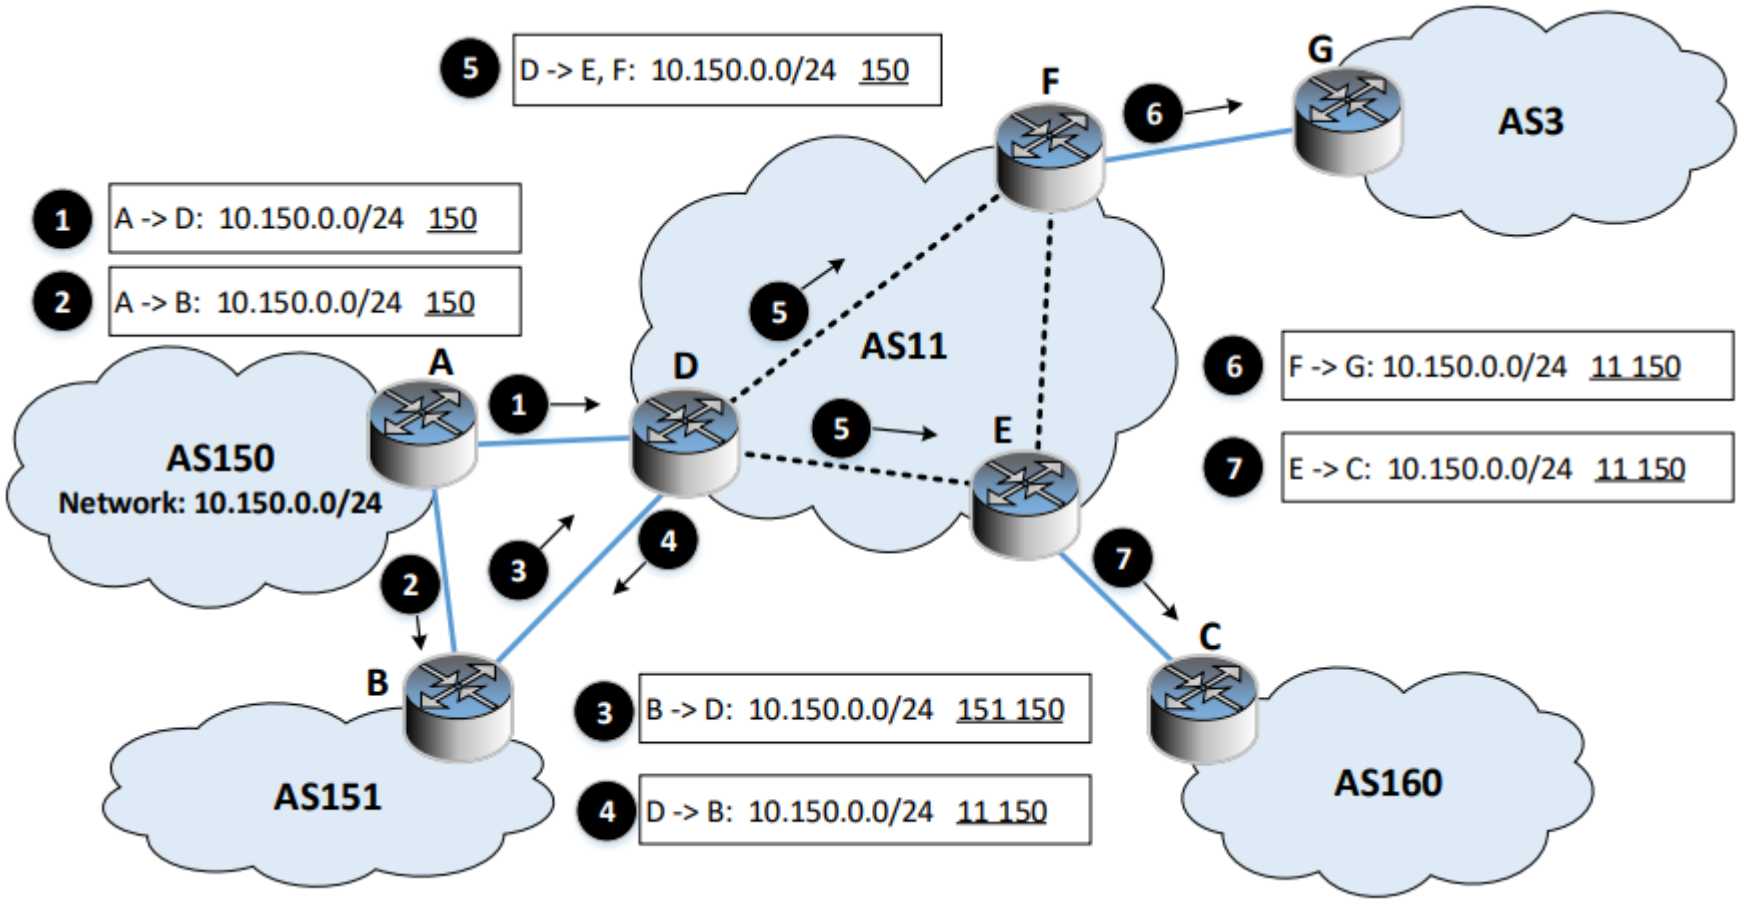

BGP Message Types

- BGP runs on TCP:
  - Byte stream-oriented
  - Unicast only
  - Connection-oriented and reliable
  - Provides flow and congestion control

BGP Message Types

- BGPv4 uses (only) 5 message types
  - OPEN
  - UPDATE
  - NOTIFICATION
  - KEEPALIVE
  - ROUTE-REFRESH

BGP OPEN Message

- BGP speakers use OPEN to advertise configuration and capabilities after TCP session starts:
  - Version
  - Autonomous System Number
  - Hold Time (advertise/negotiation)
  - BGP Router ID
  - Optional Capabilities (advertise/negotiation)

- Incompatible configurations terminate peering and close TCP session.

BGP OPEN Message

```
Border Gateway Protocol - OPEN Message
  Marker: ffffffffffffffffffffffffffffffff
  Length: 57
  Type: OPEN Message (1)
  Version: 4
  My AS: 64512
  Hold Time: 180
  BGP Identifier: 10.255.255.1
  Optional Parameters Length: 28
▾ Optional Parameters
  ▸ Optional Parameter: Capability
  ▸ Optional Parameter: Capability
  ▸ Optional Parameter: Capability
  ▸ Optional Parameter: Capability
  ▾ Optional Parameter: Capability
      Parameter Type: Capability (2)
      Parameter Length: 6
    ▸ Capability: Support for 4-octet AS number capability
```

BGP NOTIFICATION Message

- NOTIFICATION message sent for unrecoverable conditions to terminate peering.
- Sender closes session after NOTIFICATION.
- Contents useful for diagnostics.

```
Border Gateway Protocol - NOTIFICATION Message
 Marker: ffffffffffffffffffffffffffffffff
 Length: 21
 Type: NOTIFICATION Message (3)
 Major error Code: Cease (6)
 Minor error Code (Cease): Administratively Shutdown (2)
```

BGP KEEPALIVE Message

- BGP uses KEEPALIVE instead of TCP keepalives to show liveliness.
- KEEPALIVE is sent: Immediately after an agreeable OPEN message.
- Periodically, default is one-third of Hold Time.

```
Border Gateway Protocol - KEEPALIVE Message
 Marker: ffffffffffffffffffffffffffffffff
 Length: 19
 Type: KEEPALIVE Message (4)
```

BGP ROUTE-REFRESH Message

- Original BGP lacked a way to request prefix resend. Needed for inbound route policy changes.
- Vendors used "Soft Reconfiguration" to store unfiltered routes.

- RFC 2918 introduced ROUTE-REFRESH to request route resend for any address family

```
Border Gateway Protocol - ROUTE-REFRESH Message
 Marker: ffffffffffffffffffffffffffffffff
 Length: 23
 Type: ROUTE-REFRESH Message (5)
 Address family identifier (AFI): IPv4 (1)
 Subtype: Normal route refresh request [RFC2918] with/without ORF [RFC5291] (0)
 Subsequent address family identifier (SAFI): Unicast (1)
```

BGP UPDATE Message

- UPDATE message is BGP's workhorse.
- Advertises reachable (Network Layer Reachability Information) NLRIs with attributes.
- Withdraws unreachable NLRIs.
- Designed for maximum efficiency:
- Path attributes included once, followed by NLRIs sharing them.
- Each NLRI contains only the network prefix (with padding if needed).

BGP Update Message – New/Updated Routes

```
Border Gateway Protocol - UPDATE Message
 Marker: ffffffffffffffffffffffffffffffff
 Length: 67
 Type: UPDATE Message (2)
 Withdrawn Routes Length: 0
 Total Path Attribute Length: 28
▾Path attributes
 ▸Path Attribute - ORIGIN: IGP
 ▸Path Attribute - AS_PATH: empty
 ▸Path Attribute - NEXT_HOP: 10.255.255.1
 ▸Path Attribute - MULTI_EXIT_DISC: 1234
 ▸Path Attribute - LOCAL_PREF: 100
▾Network Layer Reachability Information (NLRI)
 ▸192.168.0.0/24
 ▸192.168.1.0/24
 ▸192.168.2.0/24
 ▸192.168.3.0/24
```

BGP Update Message – Withdrawn Routes

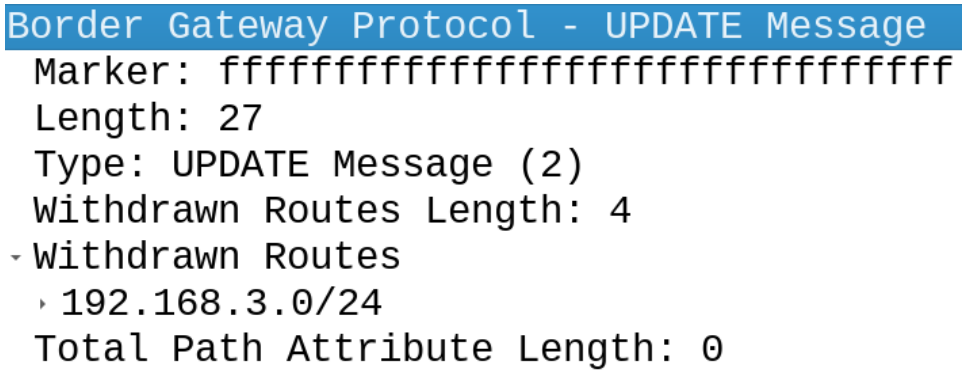

BGP Finite State Machine (FSM)

- BGP peer undergoes several state changes in its life cycle
  - IDLE
  - CONNECT
  - ACTIVE
  - OPENSENT
  - OPENCONFIRM
  - ESTABLISHED
- During each state, peers must send and receive messages, process message data, and initialize resources before proceeding to the next state.

BGP Finite State Machine (FSM)

- **IDLE**
  - Initializes resources
  - Resets ConnectRetryTimer
  - RFC4271 suggested 120 seconds as the default value for ConnectRetryTime
  - Initiates a TCP connection with its configured BGP peer, and listens for a TCP connection from its peer
- If no error, changes peer's state to CONNECT
- If error occurred, peer remains in IDLE state

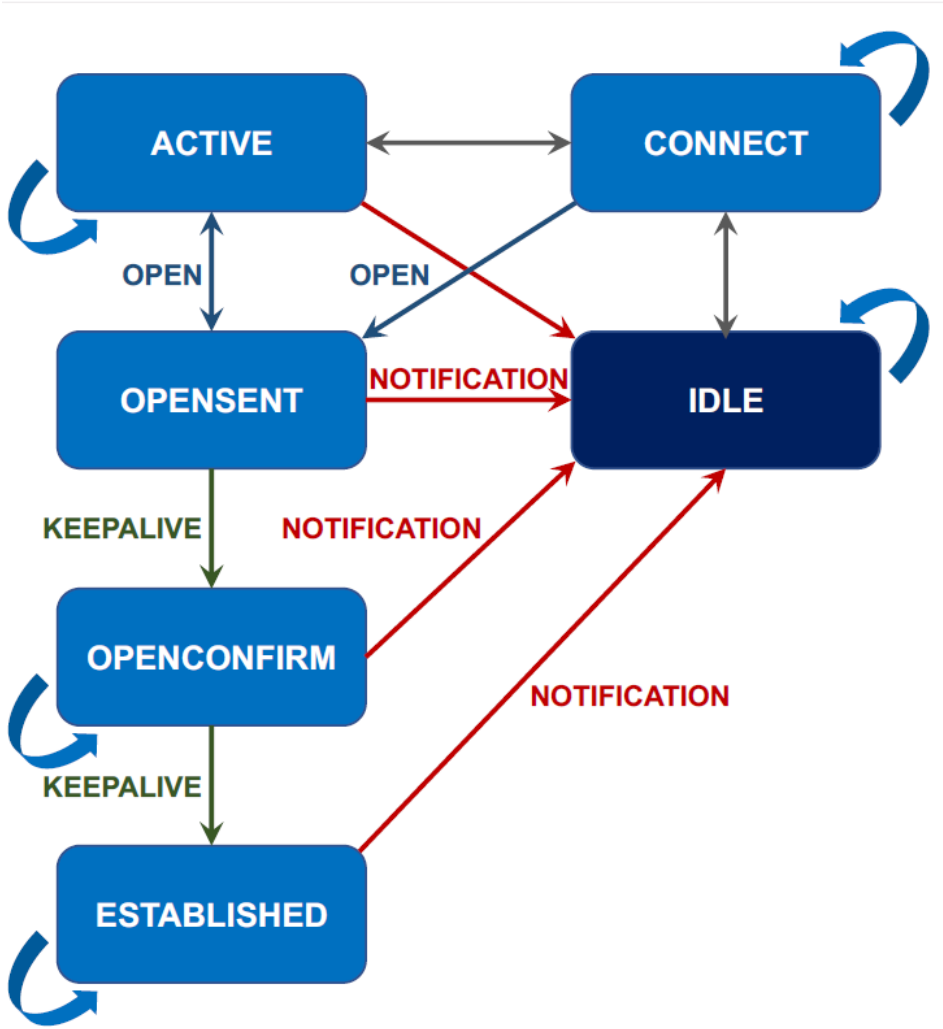

CONNECT

- Waits for successful TCP session
- Sends OPEN to peer
  - If no error, changes peer's state to OPENSENT
  - If error occurs, changes peer's state to ACTIVE
  - If ConnectRetryTimer expired, keeps peer in CONNECT state and resets ConnectRetryTimer, then tries a new TCP three-way handshake
- If something else happens, moves peer back to IDLE state

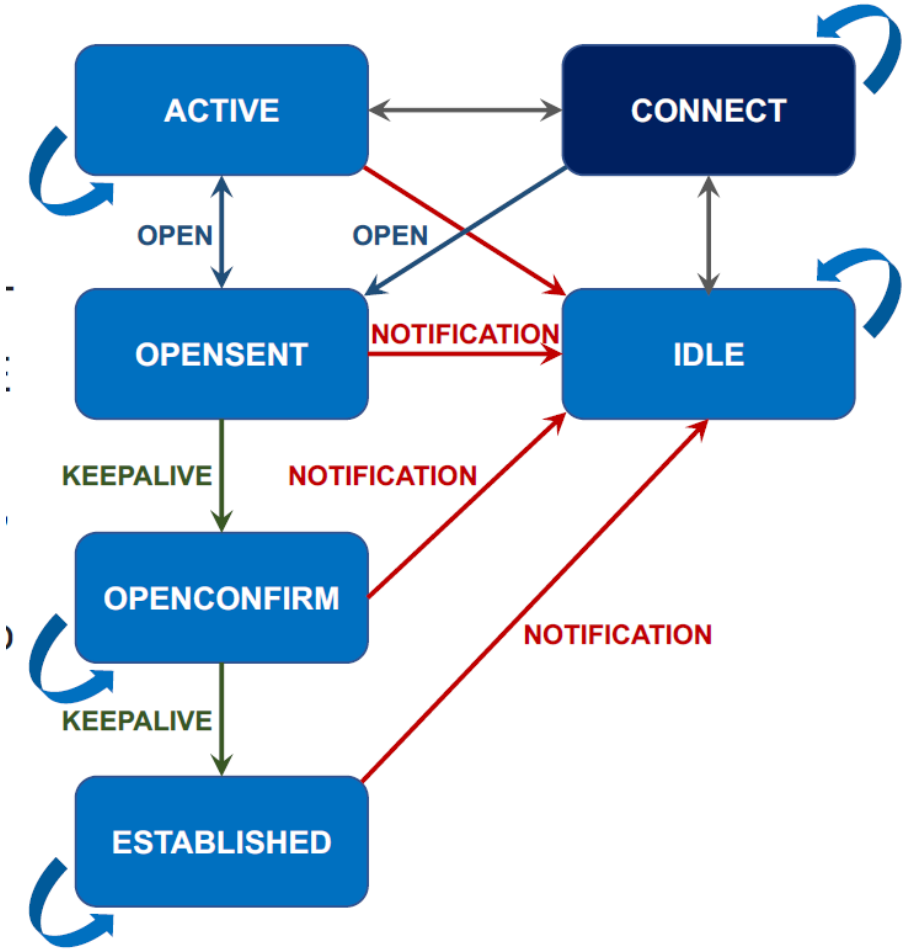

ACTIVE

- Unable to establish a successful TCP session
- Tries to restart another TCP session withthe peer
  - If successful, sends an OPEN to the peer and changes peer's state to OPENSENT
  - If unsuccessful, changes peer's state to IDLE
  - If ConnectRetryTimer expires, moves peer back to CONNECT state

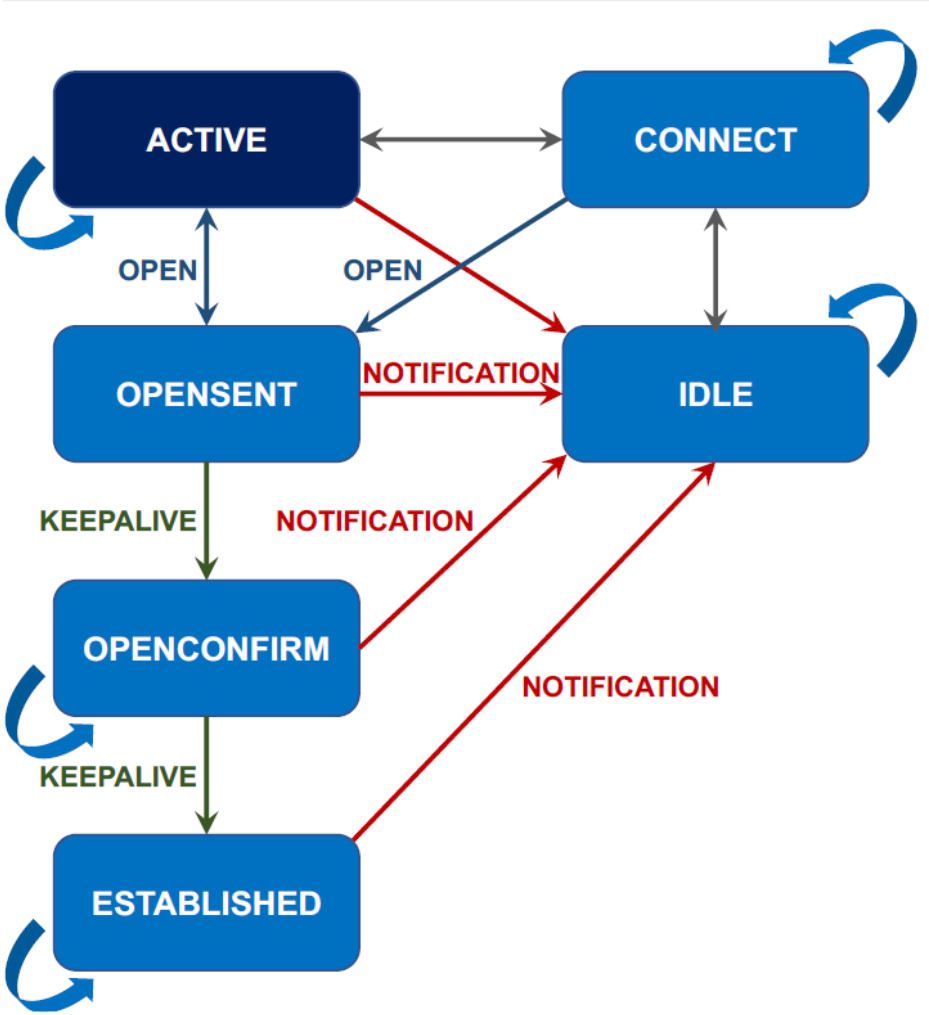

OPENSENT

- OPEN has been sent to peer
- Waits for OPEN from peer
- Checks validity of the received OPEN
  - If there is no error, sends KEEPALIVE message and changes peer's state to OPENCONFIRM
  - If error occurs due to mismatched OPEN between peers, sends NOTIFICATION and change peer's state to IDLE
- In case TCP session fails, moves peer back to ACTIVE state

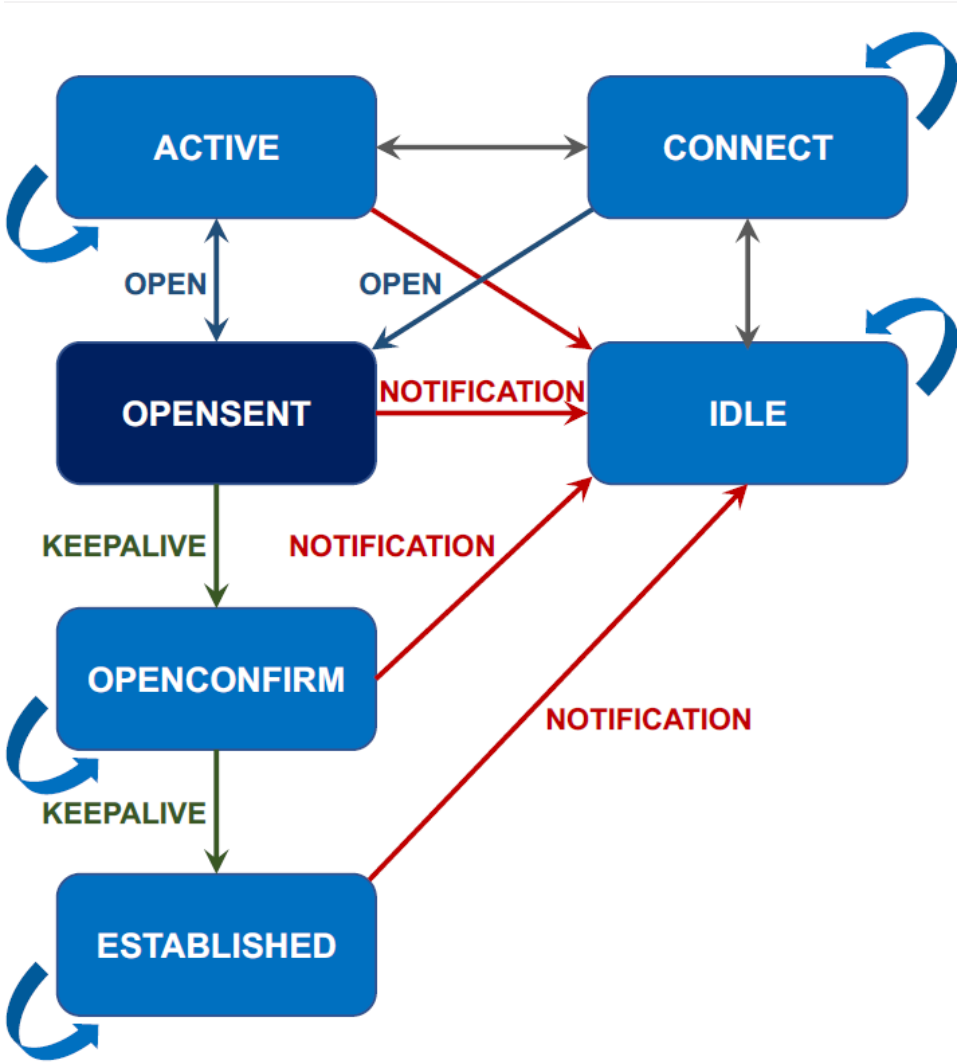

OPENCONFIRM

- Waits for a KEEPALIVE or NOTIFICATION from the peer
  - Upon receipt of peer's KEEPALIVE, changes peer's state to ESTABLISHED
  - If the HoldTimer expires or NOTIFICATION is received, changes peer's state to IDLE

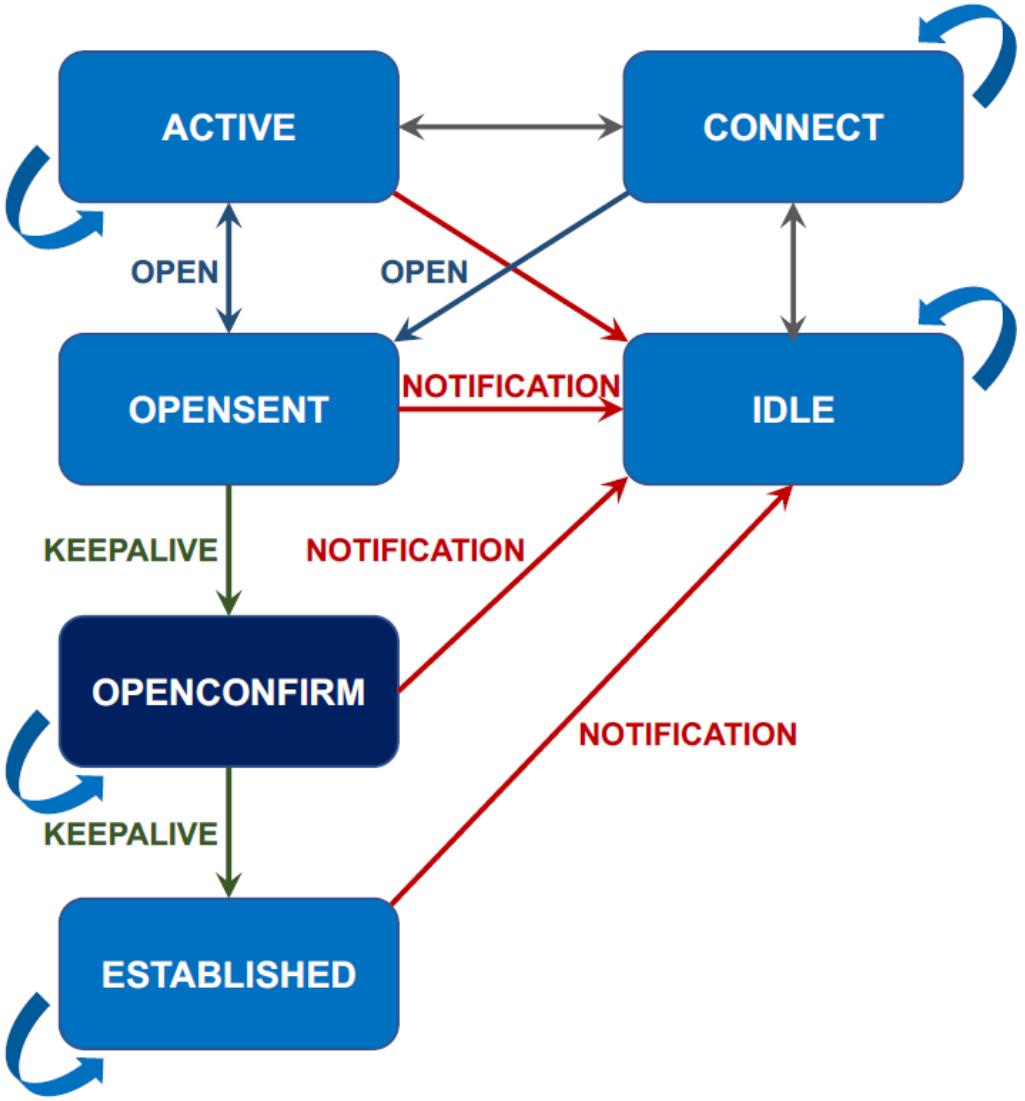

ESTABLISHED

- BGP peer adjacency is complete
- UPDATE is used for exchanging reachability information
  - Initial full routing table exchange
  - Incremental updates for later changes
- In case NOTIFICATION is received, changes peer's state back to IDLE

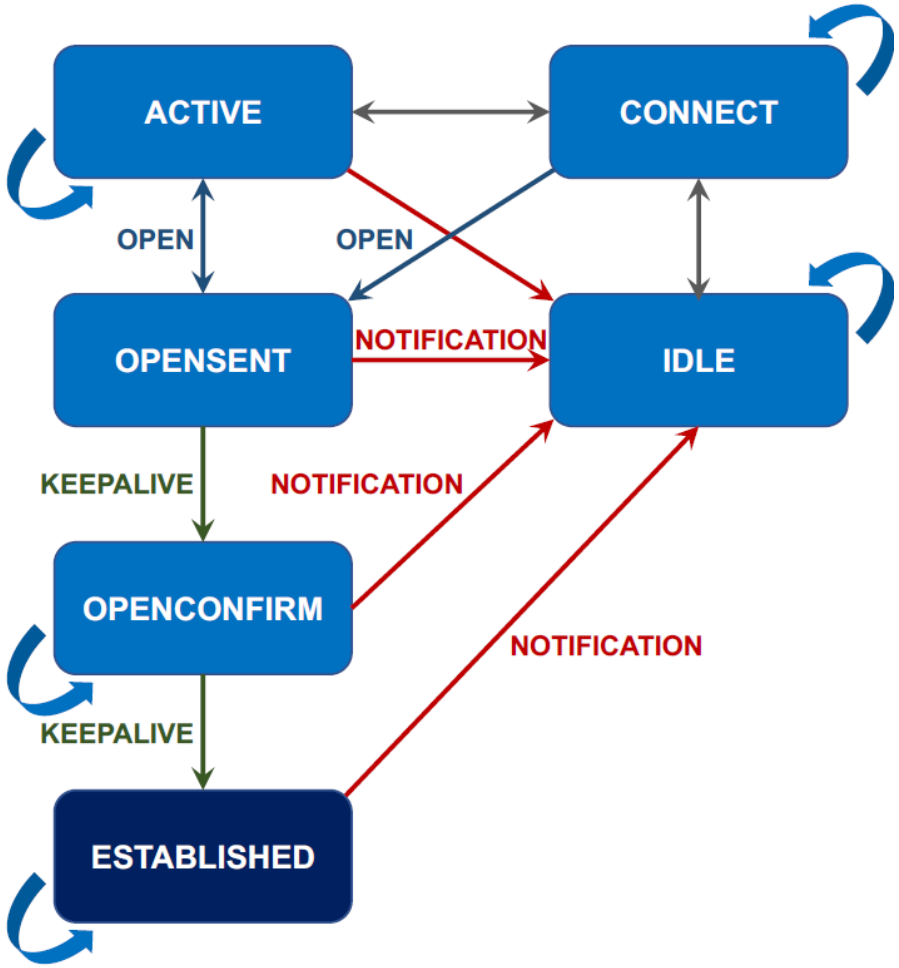

BGP Attributes

- An attribute is an additional piece of information accompanying an advertised NLRI

- BGP uses attributes in multiple ways
  - Prevents routing loops
  - Performs best path selection
  - Filters or sorts routes
  - … and many more

- Basic BGP specification recognizes only a handful of attributes
- Several new have been added over time for various applications and uses

BGP Attribute Types

- **Well-known: Every BGP implementation must support it**
  - Well-known mandatory: Must always be included with a NLRI
  - Well-known discretionary: May be included with a NLRI as needed

- **Optional:** BGP implementations do not need to support it
  - Optional transitive: When advertising a learned NLRI, keep the attribute with the NLRI even if not recognized
  - Optional non-transitive: When advertising a learned NLRI, remove the attribute from the NLRI if not recognized
- Note: All well-known attributes are transitive

BGP Attributes

- Well-known mandatory
  - AS_PATH
  - NEXT_HOP
  - ORIGIN
- Well-known discretionary
  - LOCAL_PREF
  - ATOMIC_AGGREGATE
- Optional transitive
  - AGGREGATOR
  - COMMUNITIES
  - EXTENDED_COMMUNITIES
- Optional non-transitive
  - MULTI_EXIT_DISC
  - CLUSTER_LIST

Internal vs External BGP

Internal vs External BGP

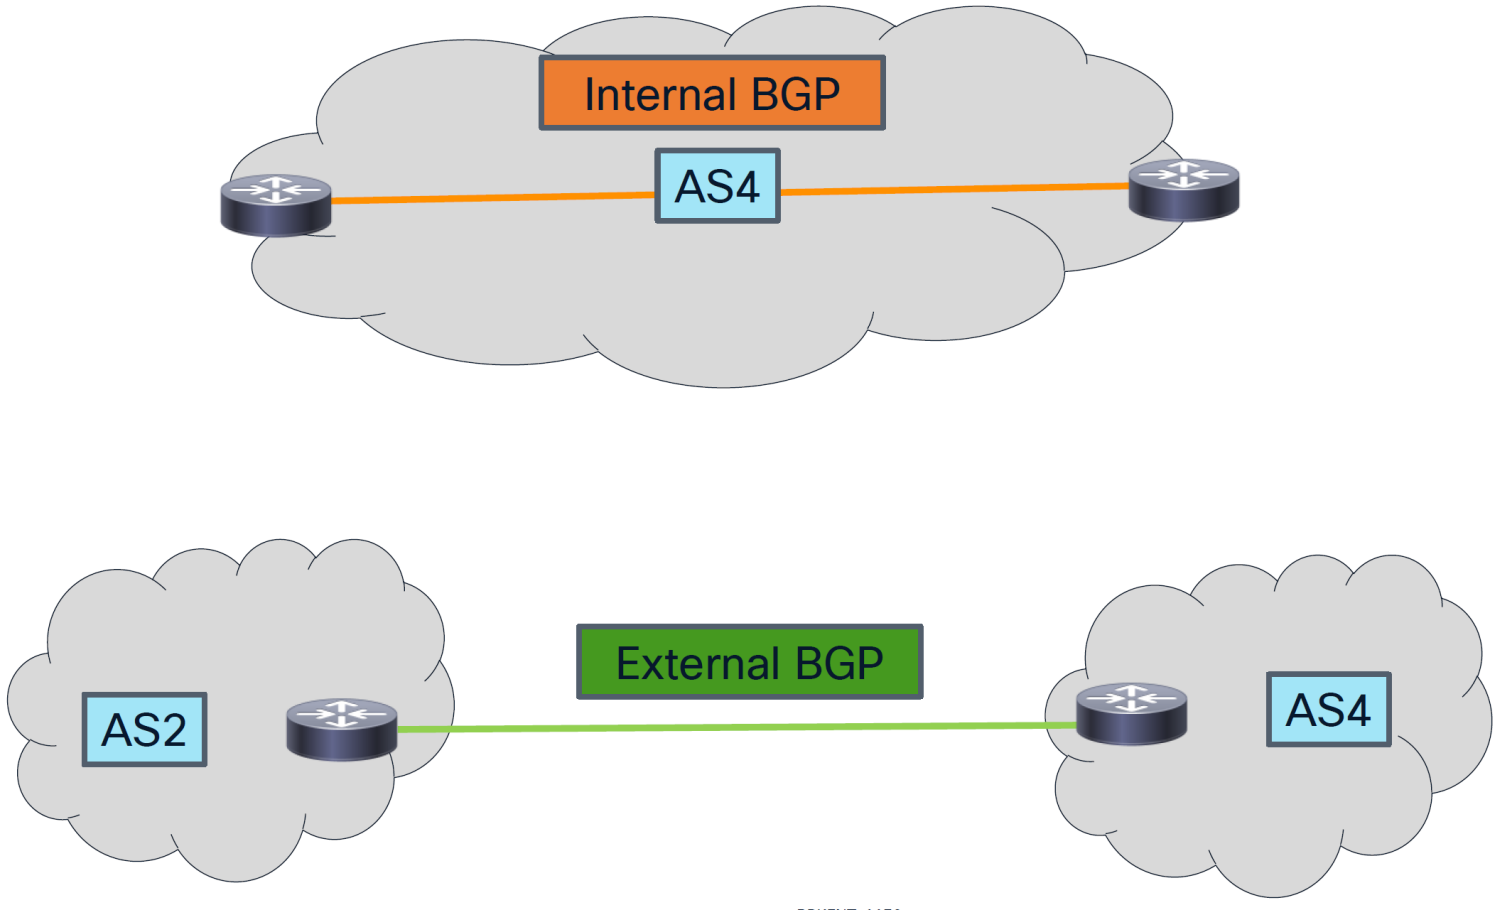

NEXT_HOP in eBGP

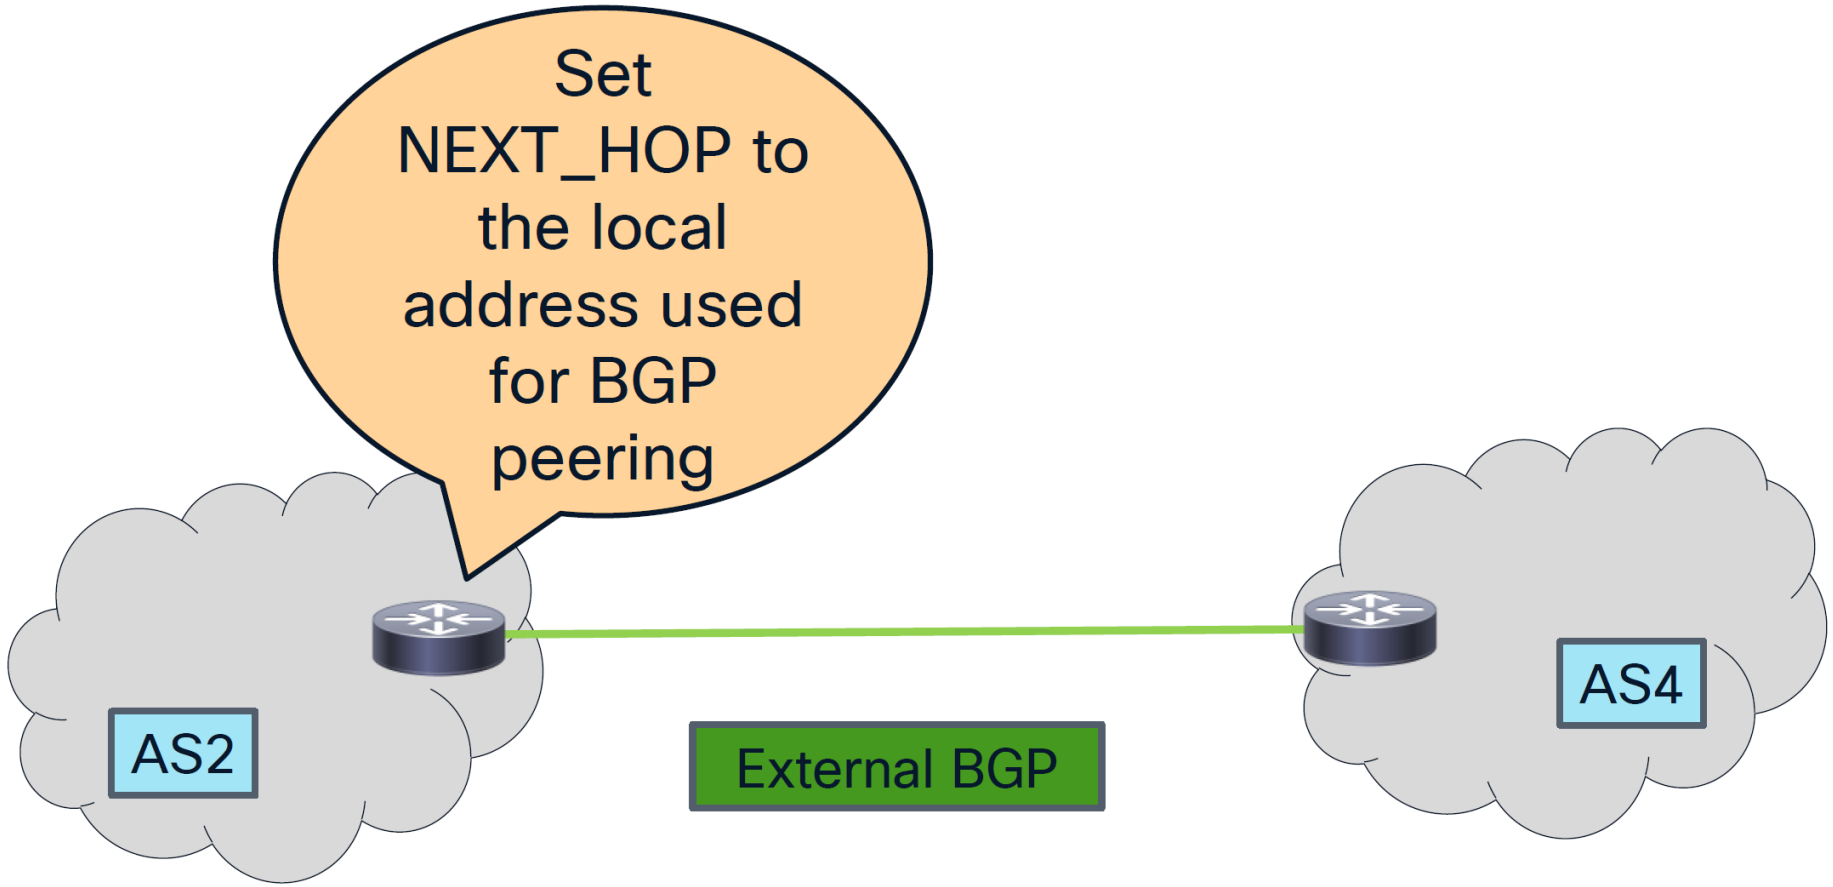

NEXT_HOP in iBGP

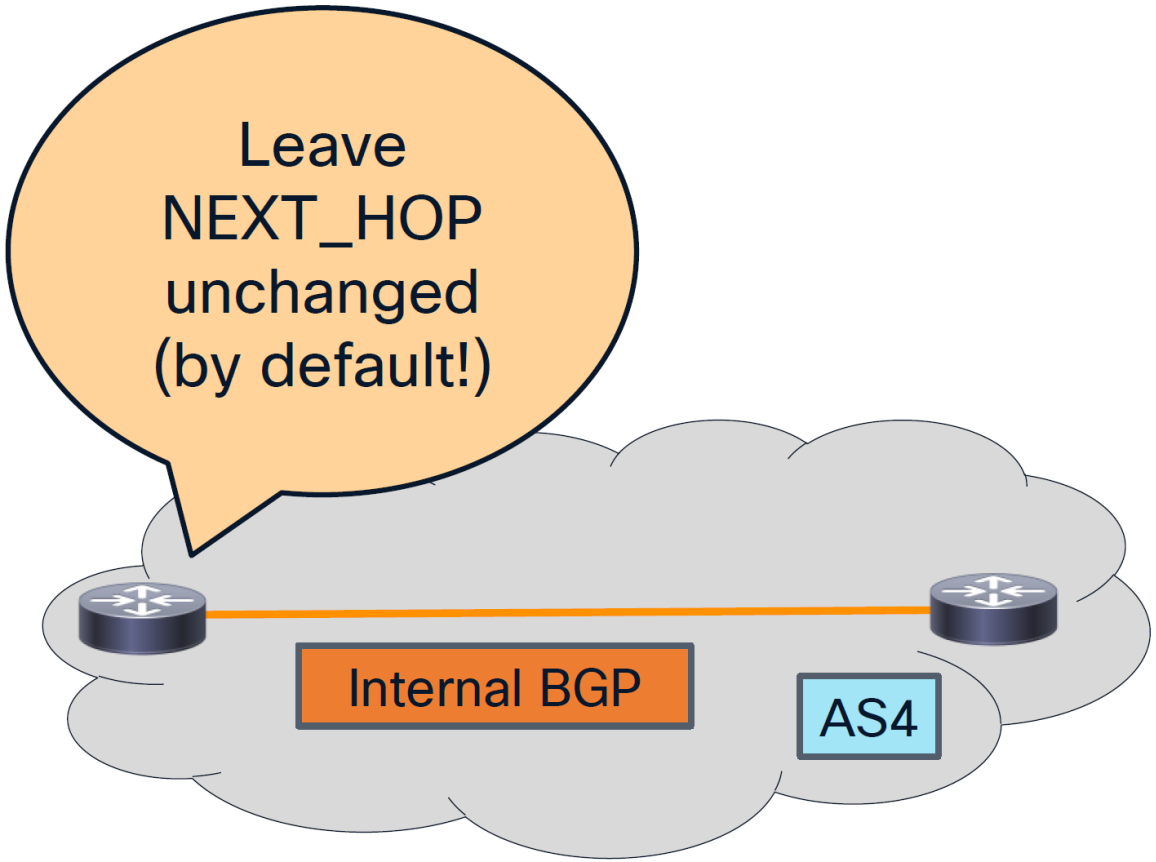

AS_PATH in eBGP

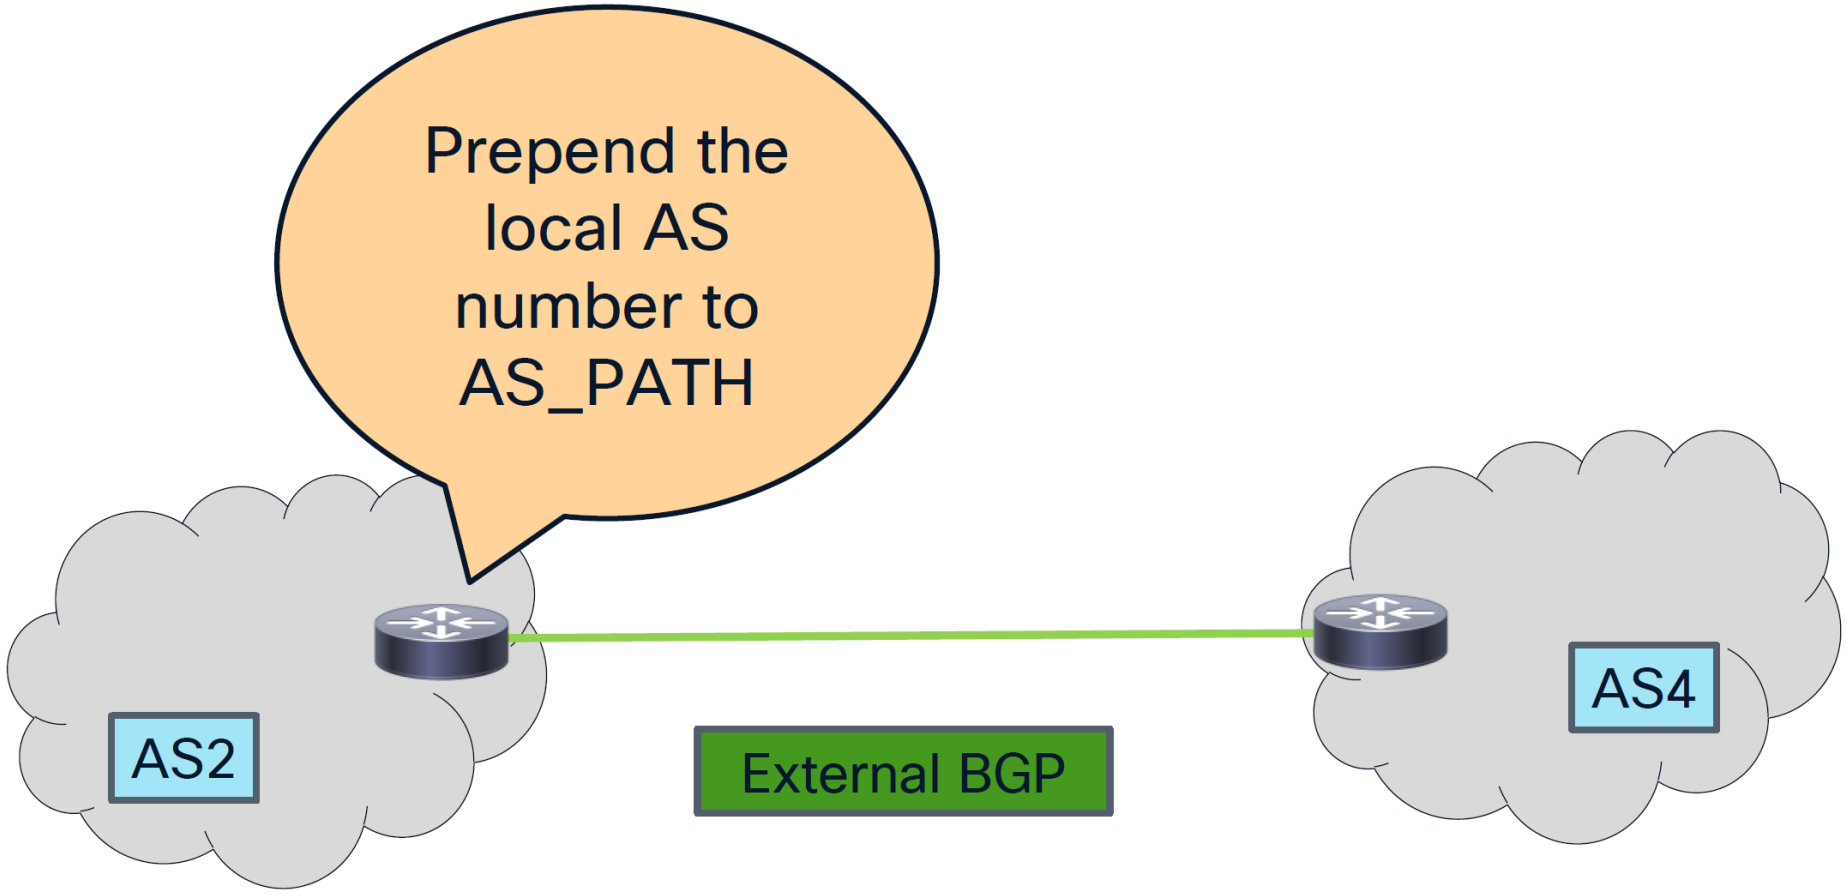

AS_PATH in iBGP

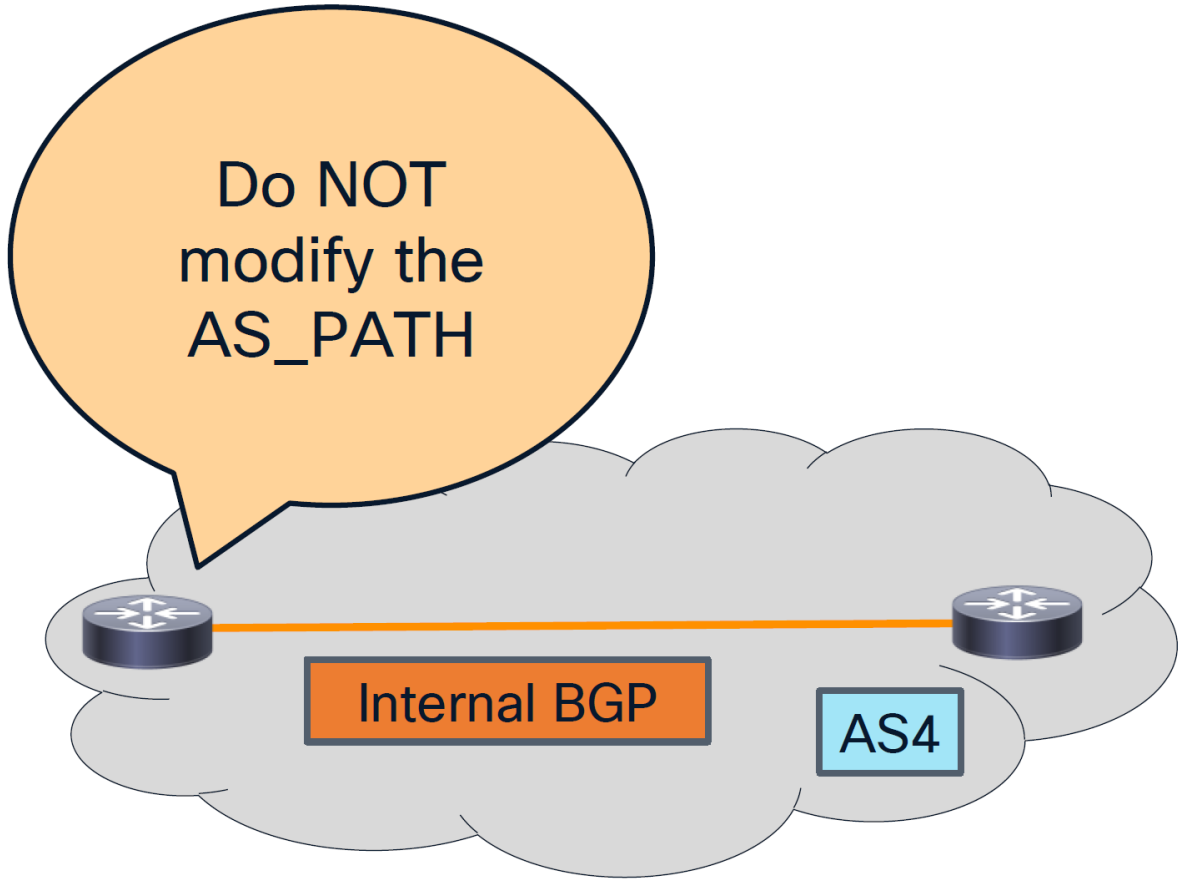

Updates in eBGP

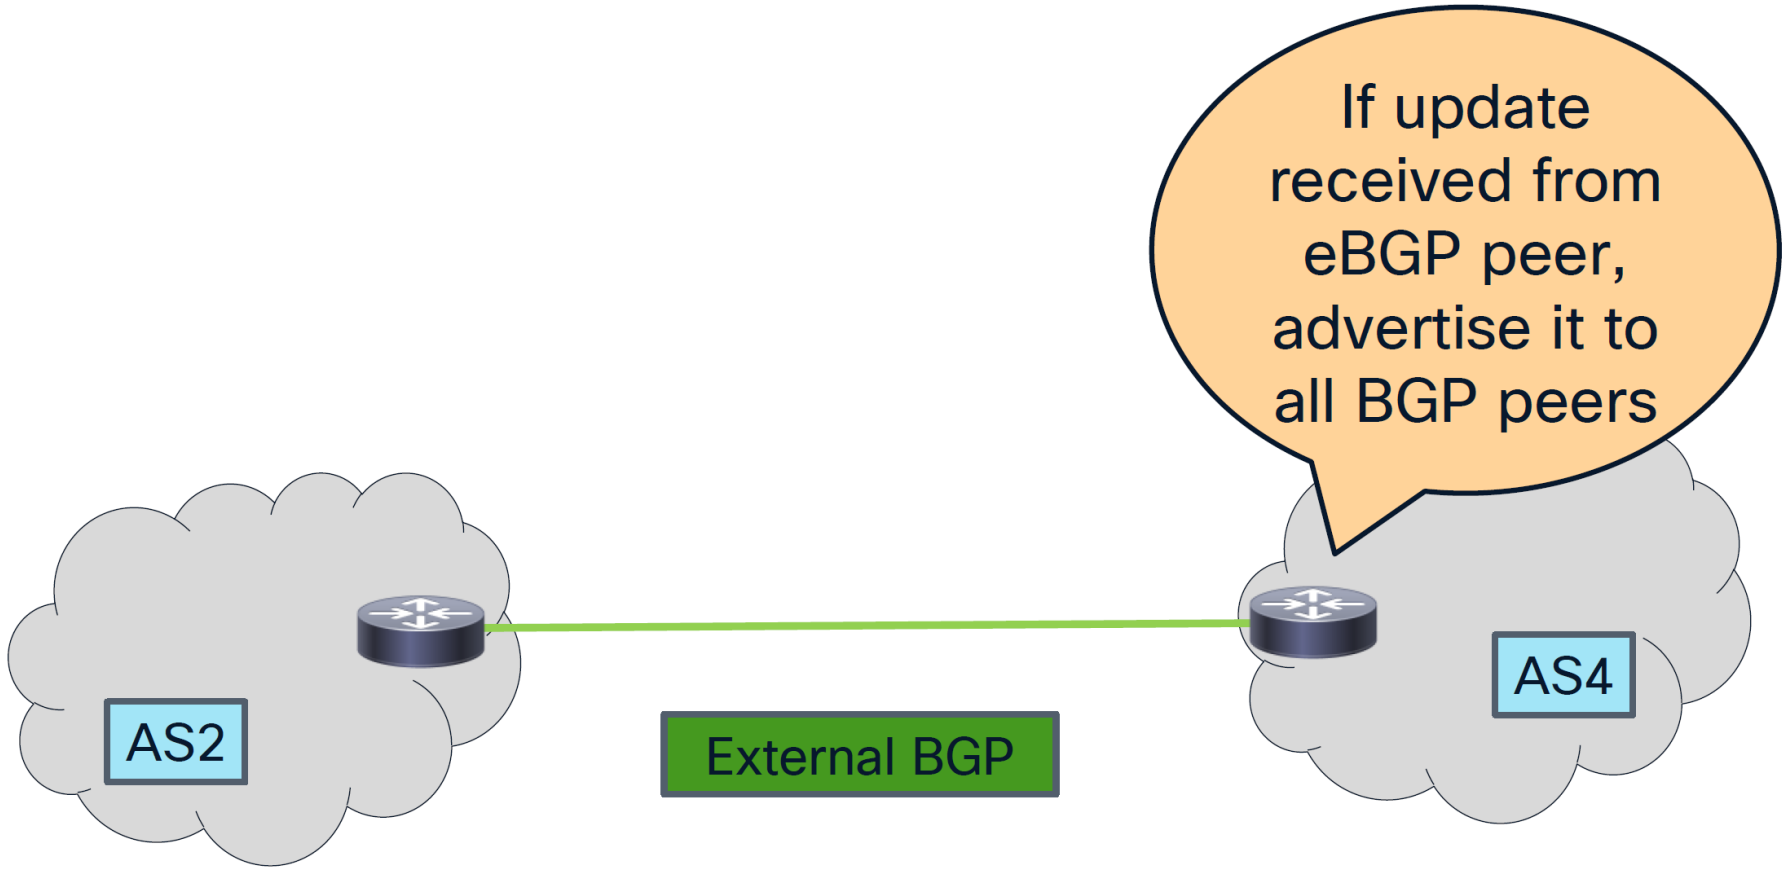

Updates in iBGP

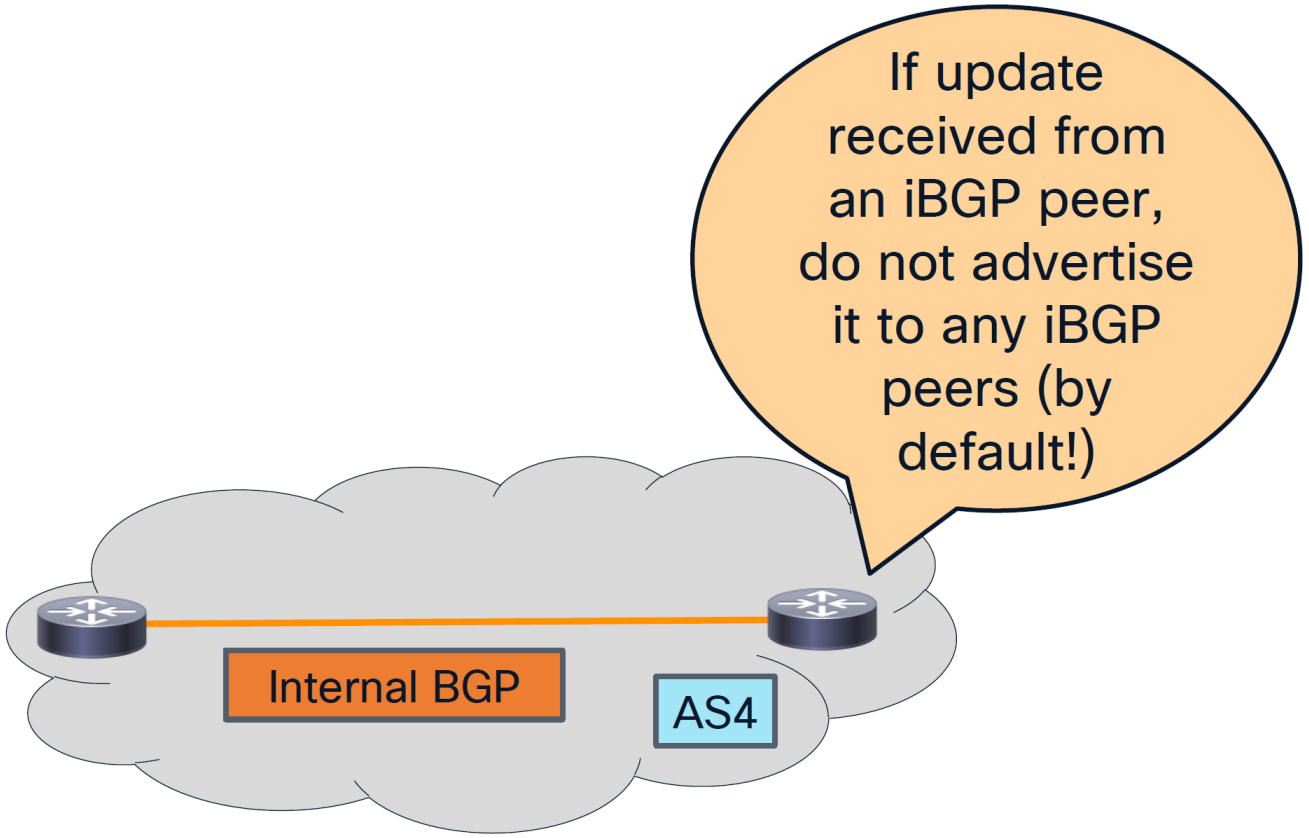

BGP Operations

BGP Transport

- BGP operates by exchanging Network Layer Reachability Information (NLRI).
  - NLRI includes a set of BGP path attributes and one or more prefixes which those attributes are associated
  - NLRI is encapsulated inside the BGP UPDATE message
- Does not have own transport protocol.
- Utilizes TCP and runs on TCP port 179.
- BGP messages are exchanged over the TCP session.

BGP Capabilities

- Capability codes indicate whether a BGP router is able to accommodate particular capabilities.
  - Advertised in OPEN message
- If received capability is not supported by remote peer, it sends back a NOTIFICATION message.
- BGP routers attempt to peer without the unsupported capability.
- Commonly implemented capabilities:
  - Route Refresh
  - Multi-protocol Extension
  - Support for 4-octet AS Number

BGP Router ID

- The BGP router identifier (ID) is a 4-byte field that is set to the highest IP address on the router.
  - (Cisco) Loopback interface addresses are considered before physical interface addresses because loopback interfaces are more stable than physical interfaces.
  - The BGP router ID is used in the BGP algorithm for determining the best path to a destination where the preference is for the BGP router with the lowest router ID.
  - It is possible to manually configure the BGP router ID using the bgp router-id command to influence the best path algorithm.

Router ID

- Cisco IOS: Highest Loopback IPv4 prefers than Highest active interface IPv4
- Juniper Junos OS: Lowest Loopback IPv4 prefers than Lowest physical interface IPv4
- MikroTik RouterOS: Lowest active interface IPv4

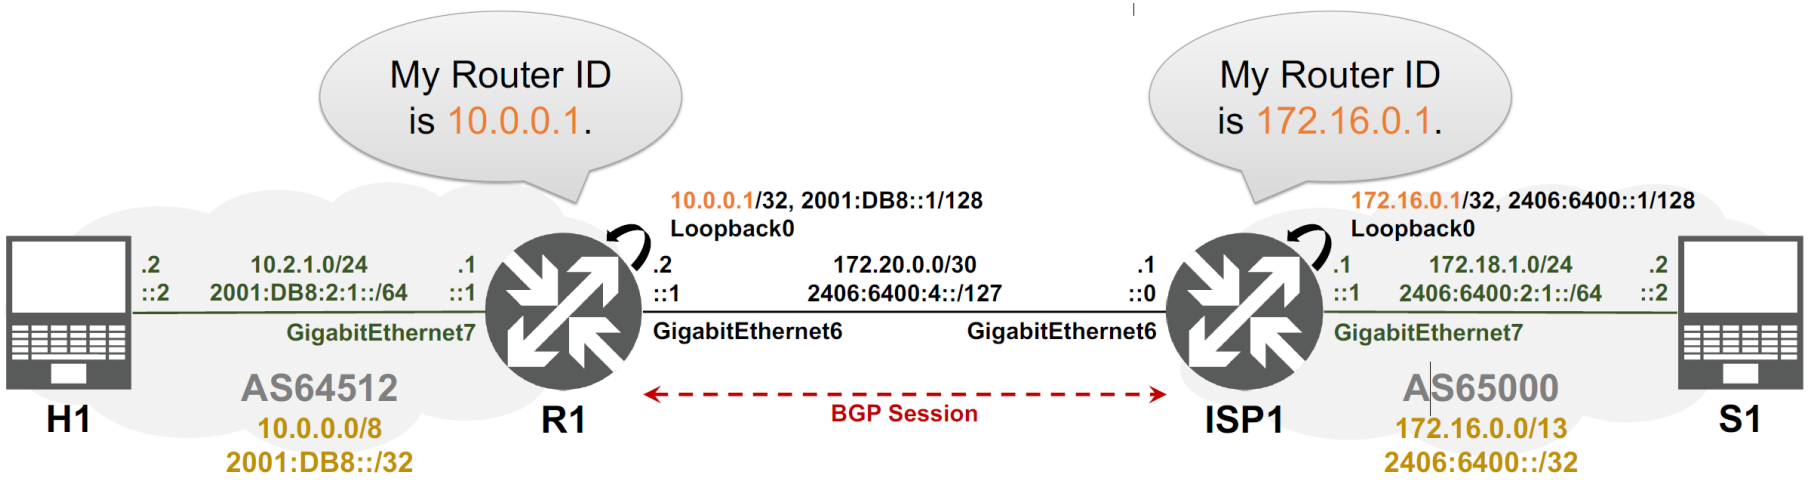

BGP Peer

- BGP does not perform auto-discovery for peers (neighbors).
- BGP peers are manually configured.
  - Local peer address and ASN
  - Remote peer address and ASN
  - Authentication (if necessary)

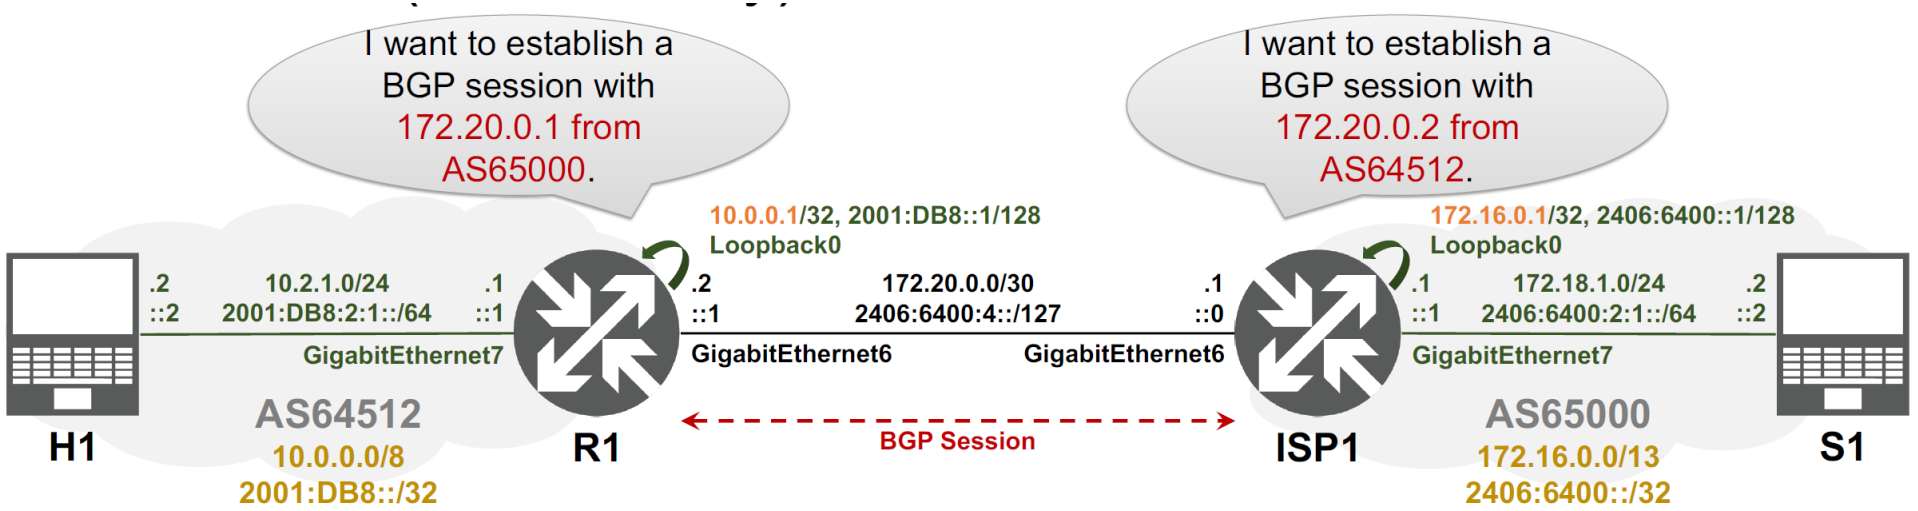

BGP Network

- Indicates a BGP prefix that should be originated by the router.
- By default, the prefix is advertised only if corresponding route is present in the routing table

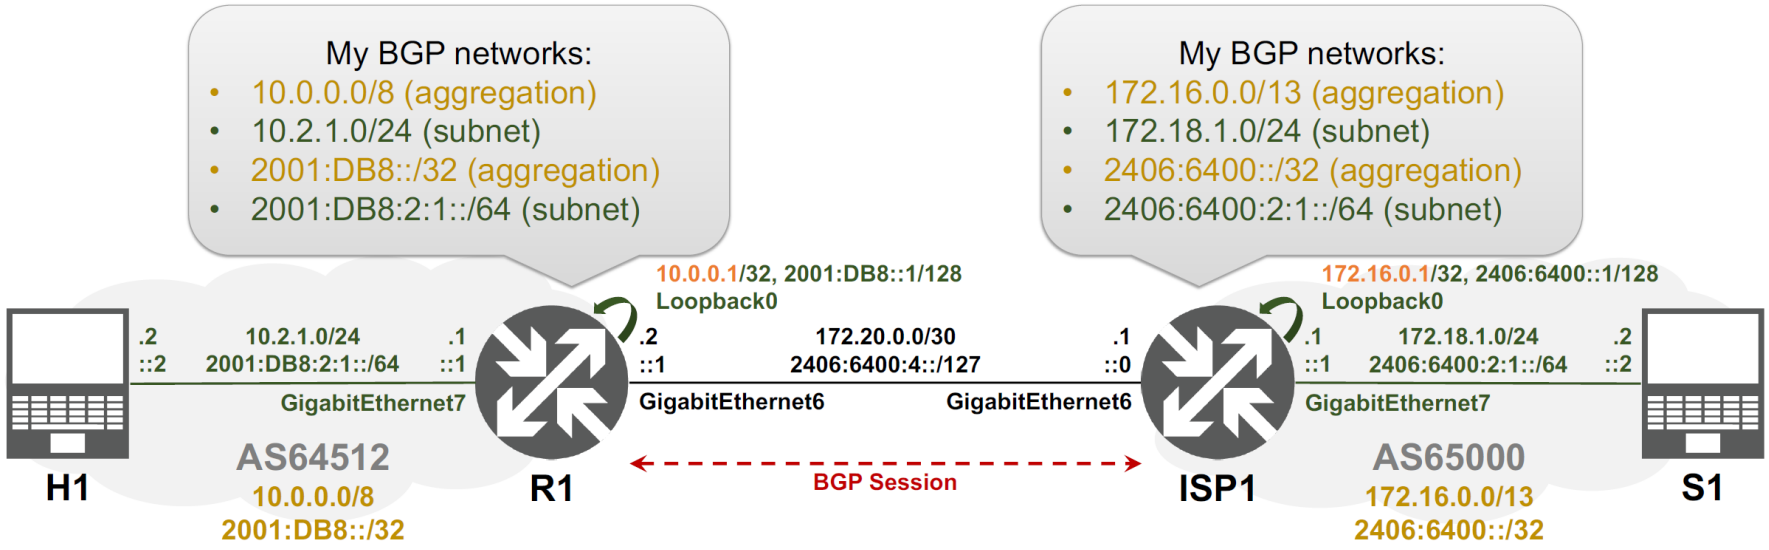

BGP Network

- Prefixes are usually subnet or aggregate routes instead of
- individual host routes.
- IPv4 prefixes longer than /24 and IPv6 prefixes longer than /48 won't generally be accepted on the Internet.

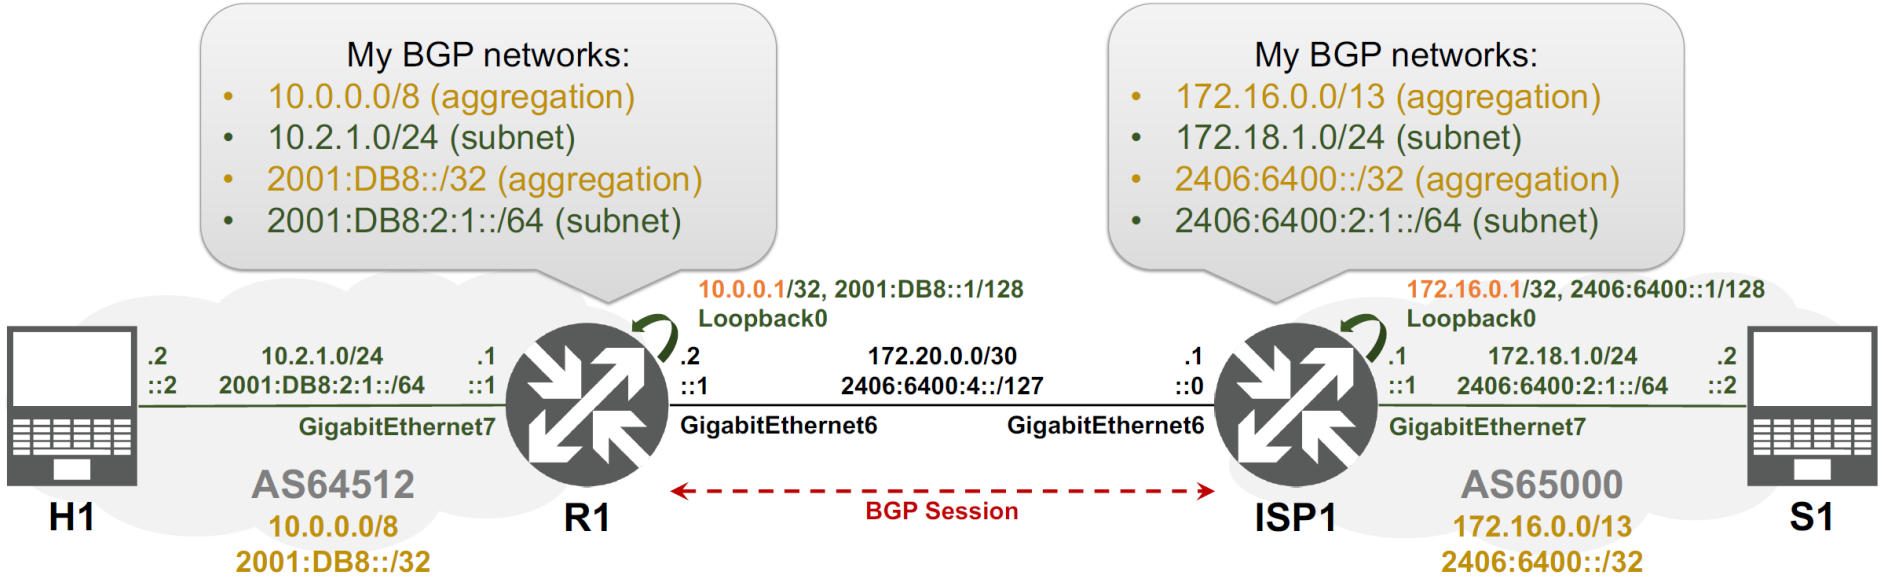

BGP Best Path

- Best path is the path that BGP selected to use in RIB.
- BGP uses path attributes to determine best path.
  - Administrator influence the selection process by routing policy
  - Best paths might not be the shortest path, but the most suitable path based on the routing policy
- By default, BGP installs single best path for each destination.
- BGP propagates only the best path to the peers.
- BGP Multipath is a feature that allows BGP to install multiple best paths when they have the same metrics.
  - For load sharing over multiple next hops

BGP Session Establishment

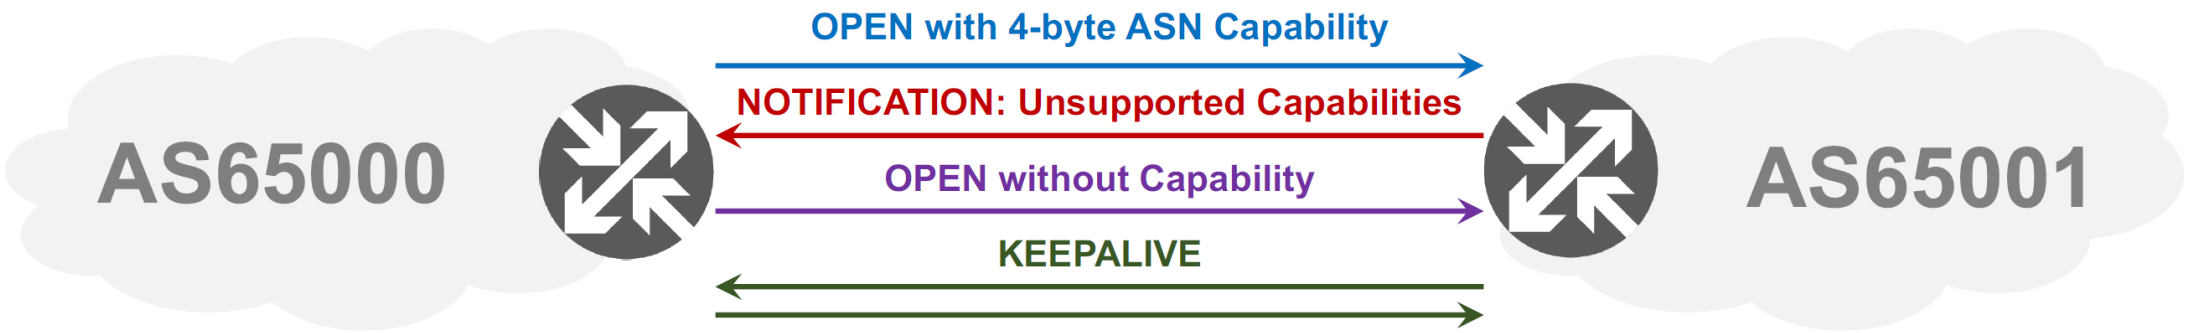

- Sends OPEN to peer after TCP three-way handshake.
- Peer replies NOTIFICATION if capabilities unsupported.
- Resends OPEN without unsupported capabilities.
- Peer replies KEEPALIVE if OPEN is acceptable.
- KEEPALIVE is sent periodically for maintaining the session.

BGP Updates

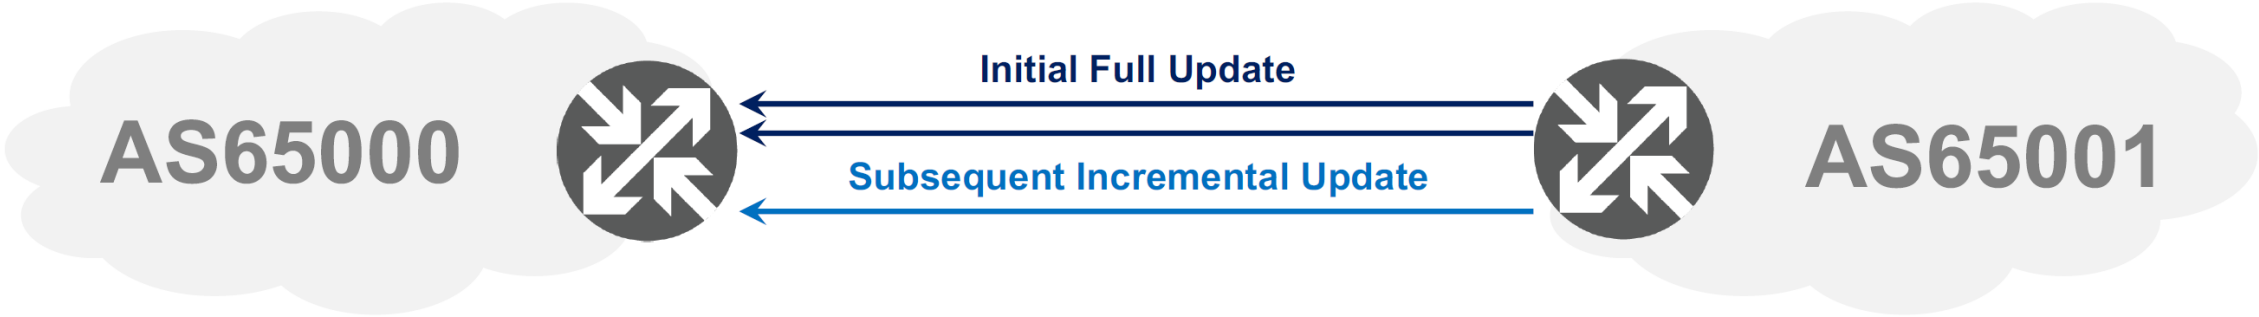

- Initial full update upon BGP session establishment.
- Subsequent incremental updates after initial full update.
  - When new prefixes are being advertised
  - When existing prefixes are being updated
  - When existing prefixes are being withdrawn

Route Refresh Capability

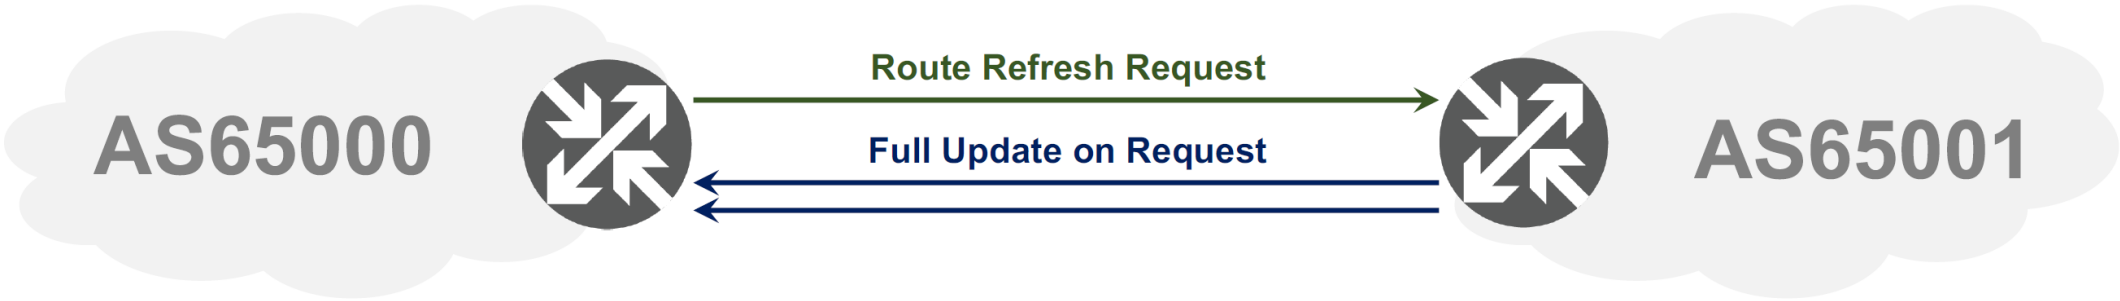

- Requests peer to resend full BGP update.
  - RFC2918: Route Refresh Capability for BGP-4

Path attributes and BGP routes

- BGP advertised route: prefix + attributes
  - prefix: destination being advertised
  - two important attributes:
    - AS-PATH: list of ASes through which prefix advertisement has passed
    - NEXT-HOP: indicates specific internal-AS router to next-hop AS
- policy-based routing:
  - gateway receiving route advertisement uses *import policy* to accept/decline path (e.g., never route through AS Y).
  - AS policy also determines whether to *advertise* path to other other neighboring ASes

BGP path advertisement

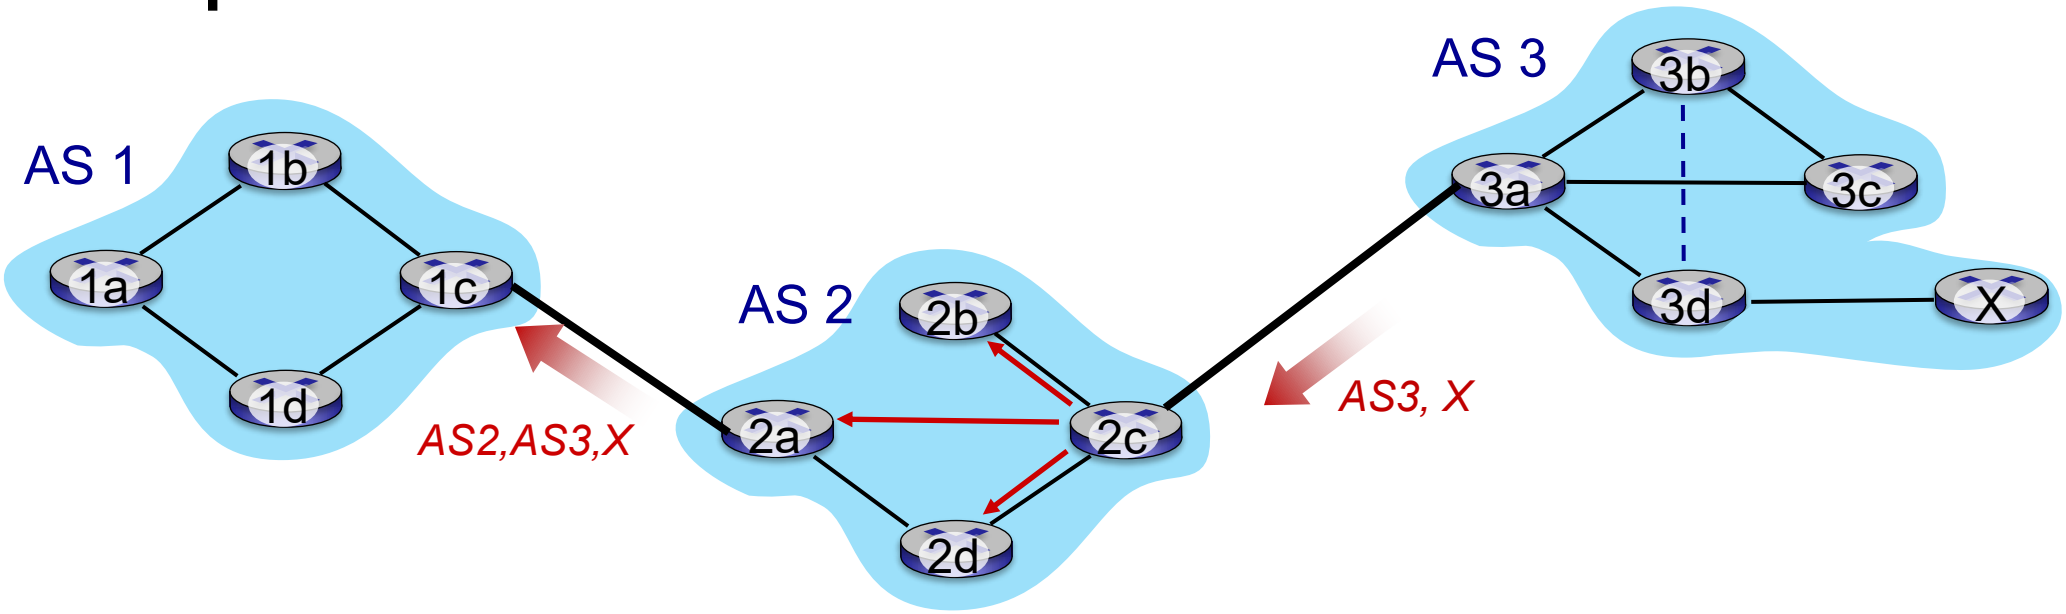

- AS2 router 2c receives path advertisement AS3,X (via eBGP) from AS3 router 3a
- based on AS2 policy, AS2 router 2c accepts path AS3,X, propagates (via iBGP) to all AS2 routers
- based on AS2 policy, AS2 router 2a advertises (via eBGP) path AS2, AS3, X to AS1 router 1c

BGP path advertisement (more)

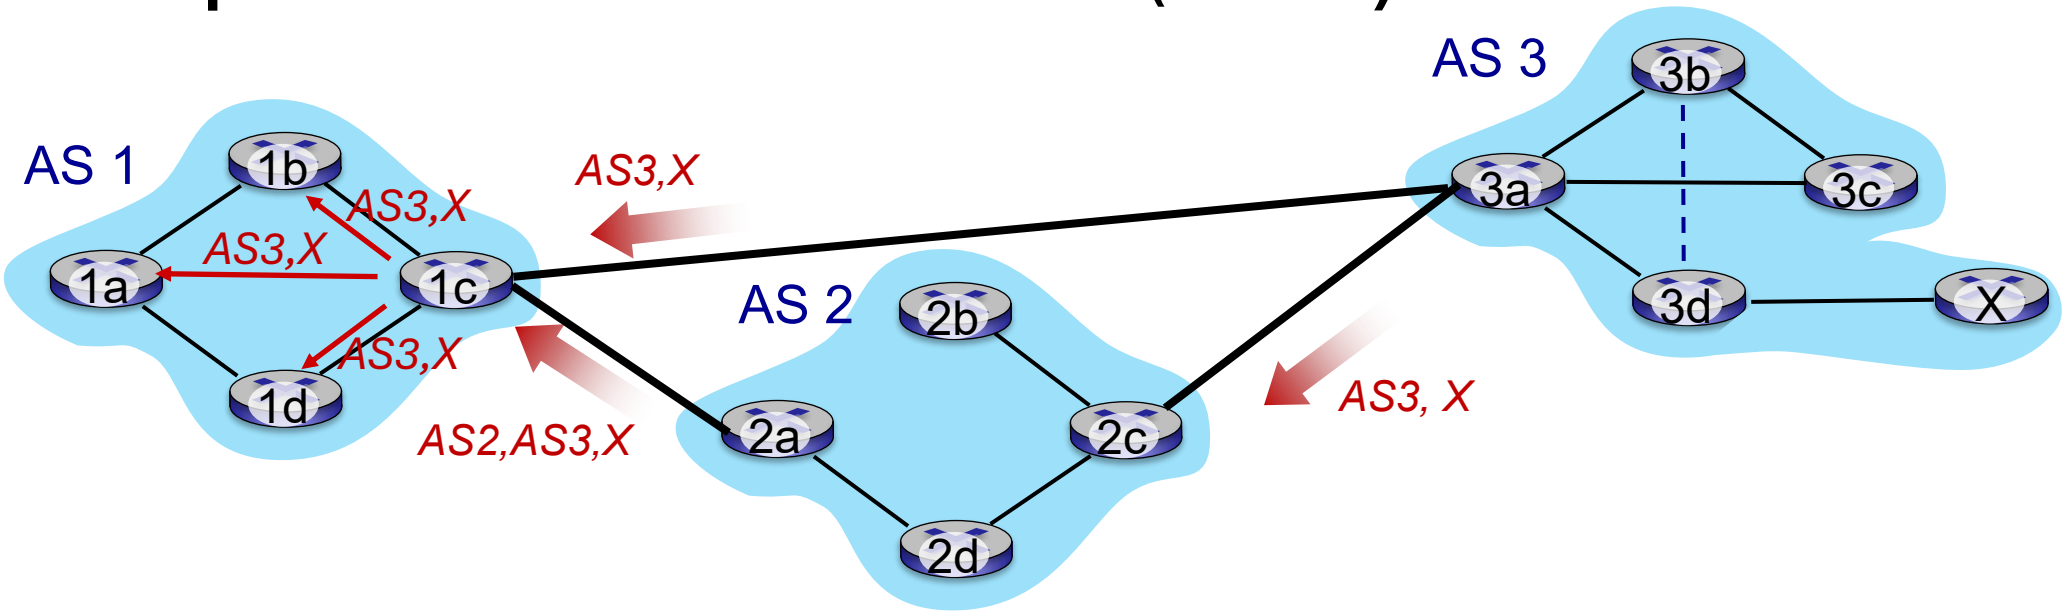

gateway router may learn about multiple paths to destination:

- AS1 gateway router 1c learns path *AS2,AS3,X* from 2a
- AS1 gateway router 1c learns path *AS3,X* from 3a
- based on *policy,* AS1 gateway router 1c chooses path *AS3,X* and advertises path within AS1 via iBGP

BGP path advertisement

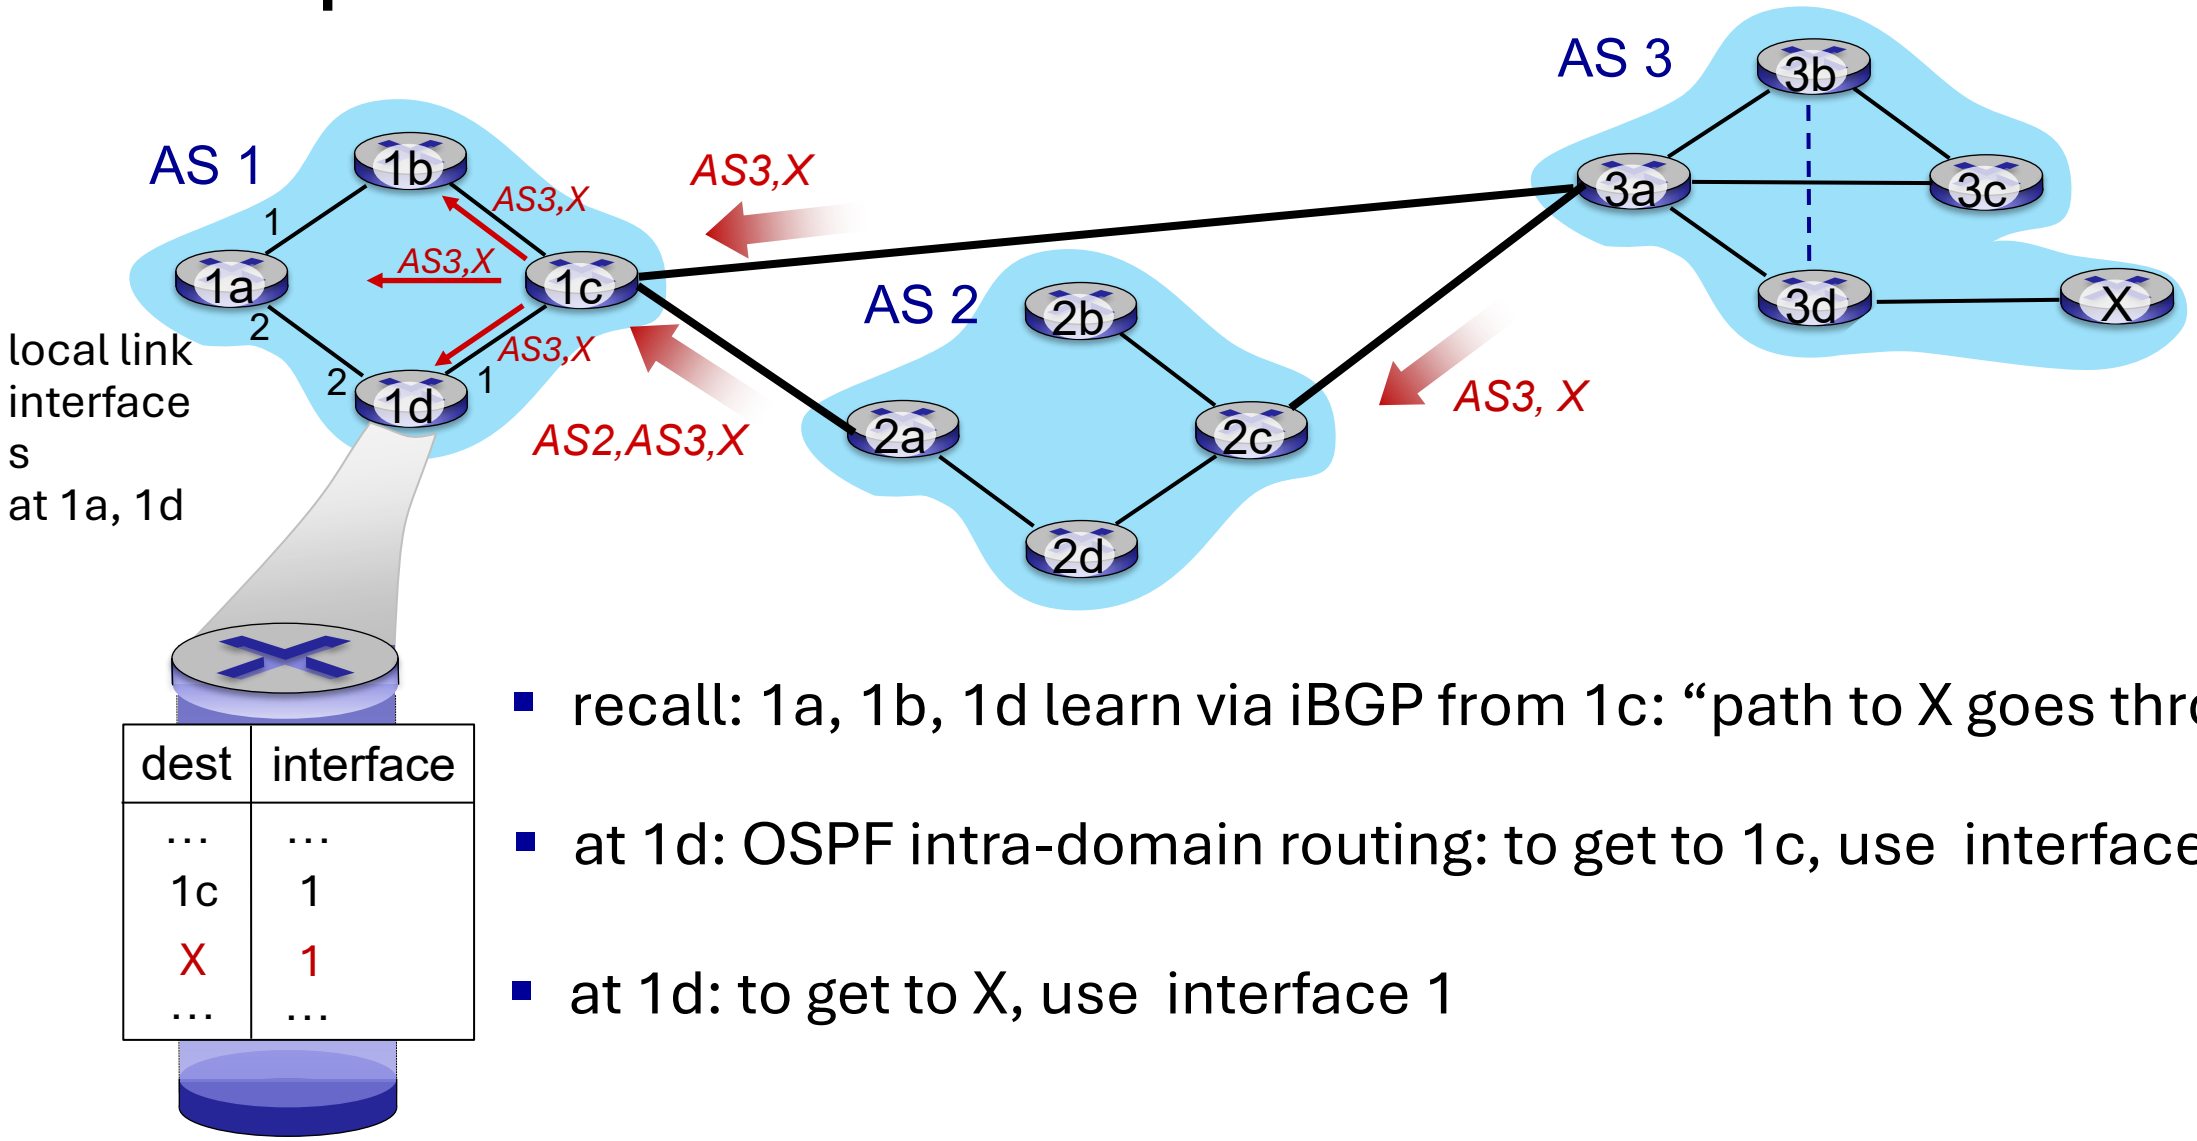

- recall: 1a, 1b, 1d learn via iBGP from 1c: “path to X goes through 1c”
- at 1d: OSPF intra-domain routing: to get to 1c, use interface 1
- at 1d: to get to X, use interface 1

BGP path advertisement

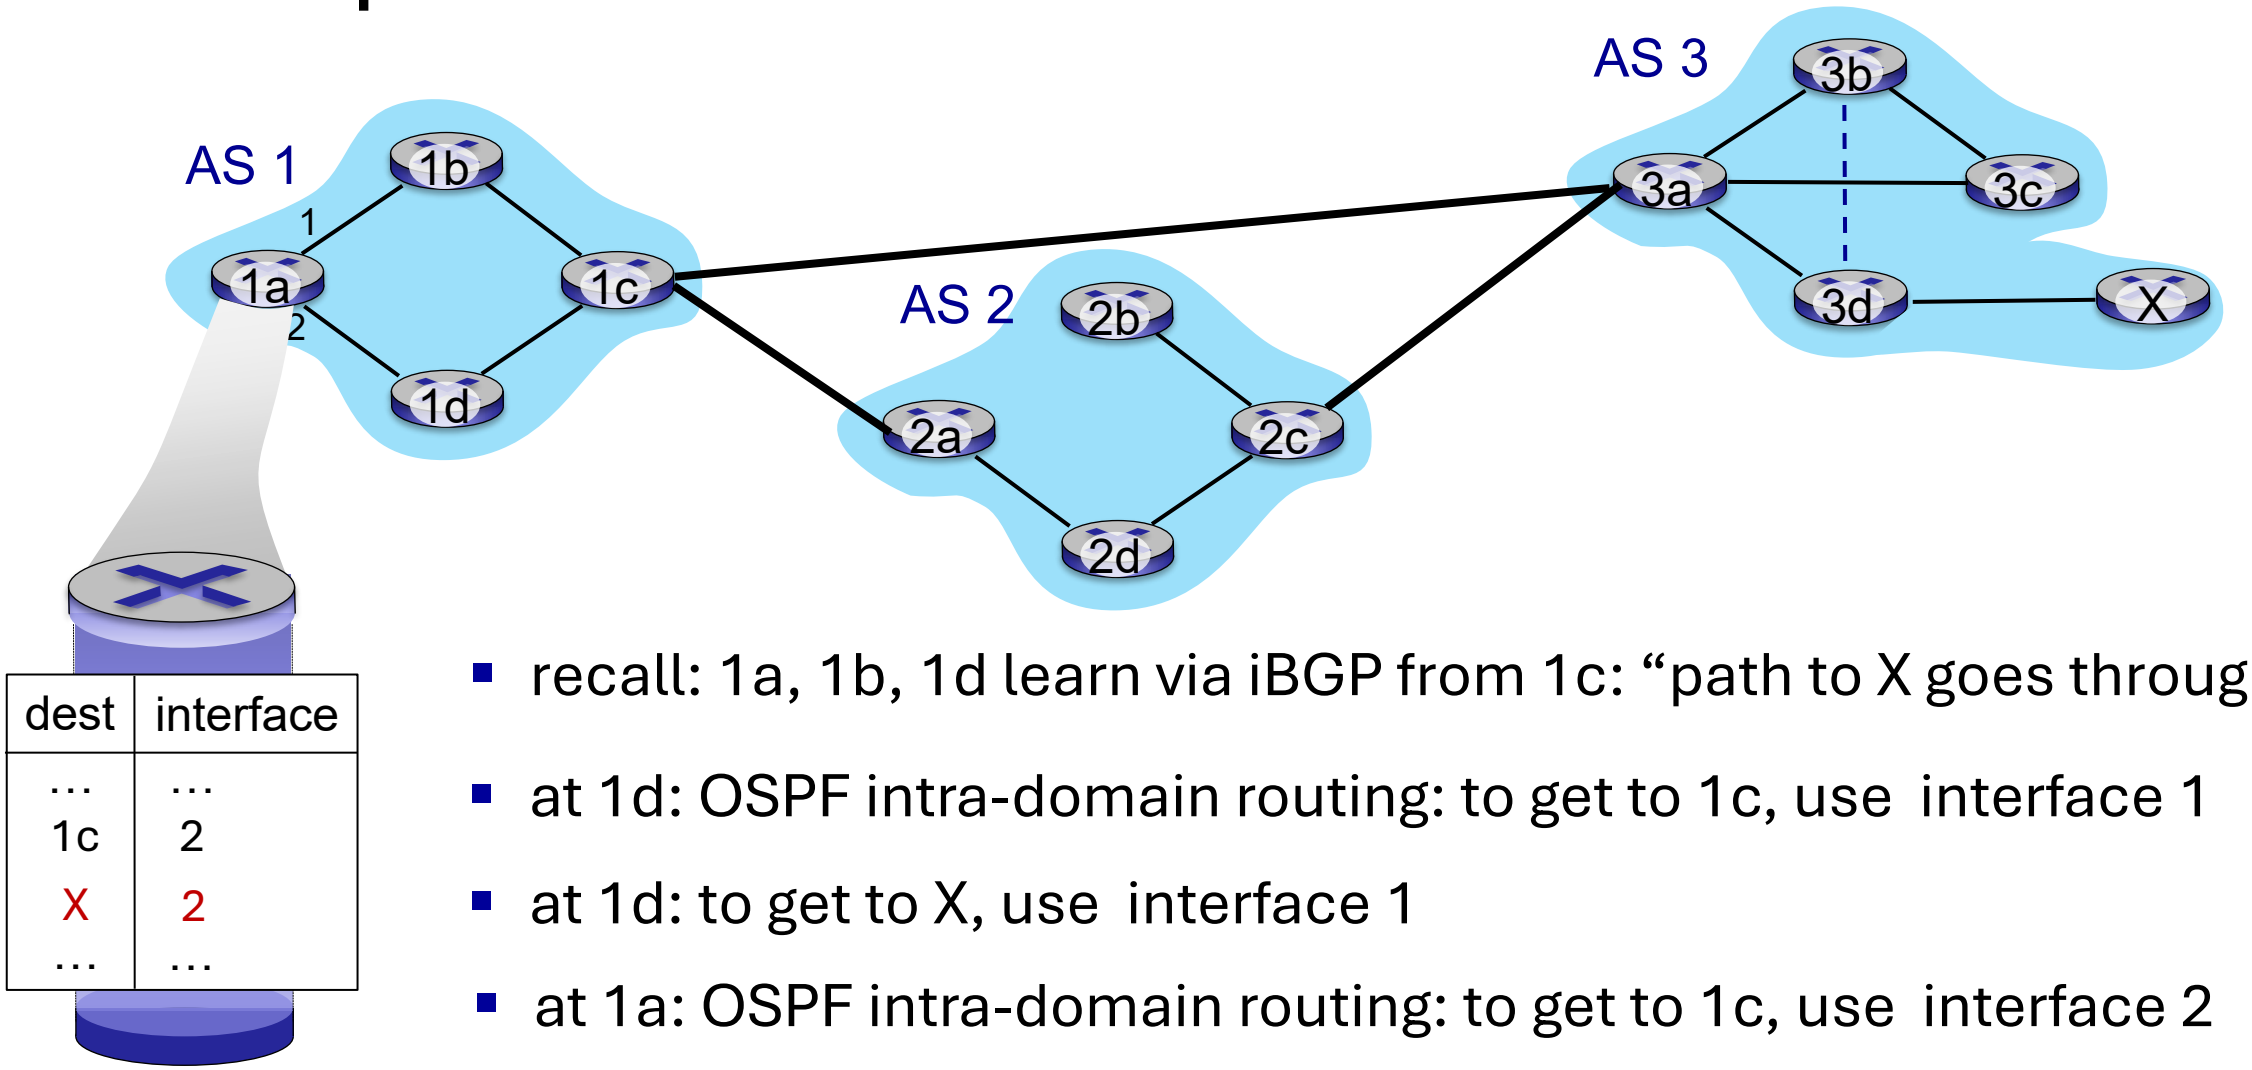

- recall: 1a, 1b, 1d learn via iBGP from 1c: “path to X goes through 1c”
- at 1d: OSPF intra-domain routing: to get to 1c, use interface 1
- at 1d: to get to X, use interface 1
- at 1a: OSPF intra-domain routing: to get to 1c, use interface 2
- at 1a: to get to X, use interface 2

Why different Intra-, Inter-AS routing ?

policy:

- inter-AS: admin wants control over how its traffic routed, who routes through its network
- intra-AS: single admin, so policy less of an issue

scale:

- hierarchical routing saves table size, reduced update traffic

performance:

- intra-AS: can focus on performance
- inter-AS: policy dominates over performance

Hot potato routing

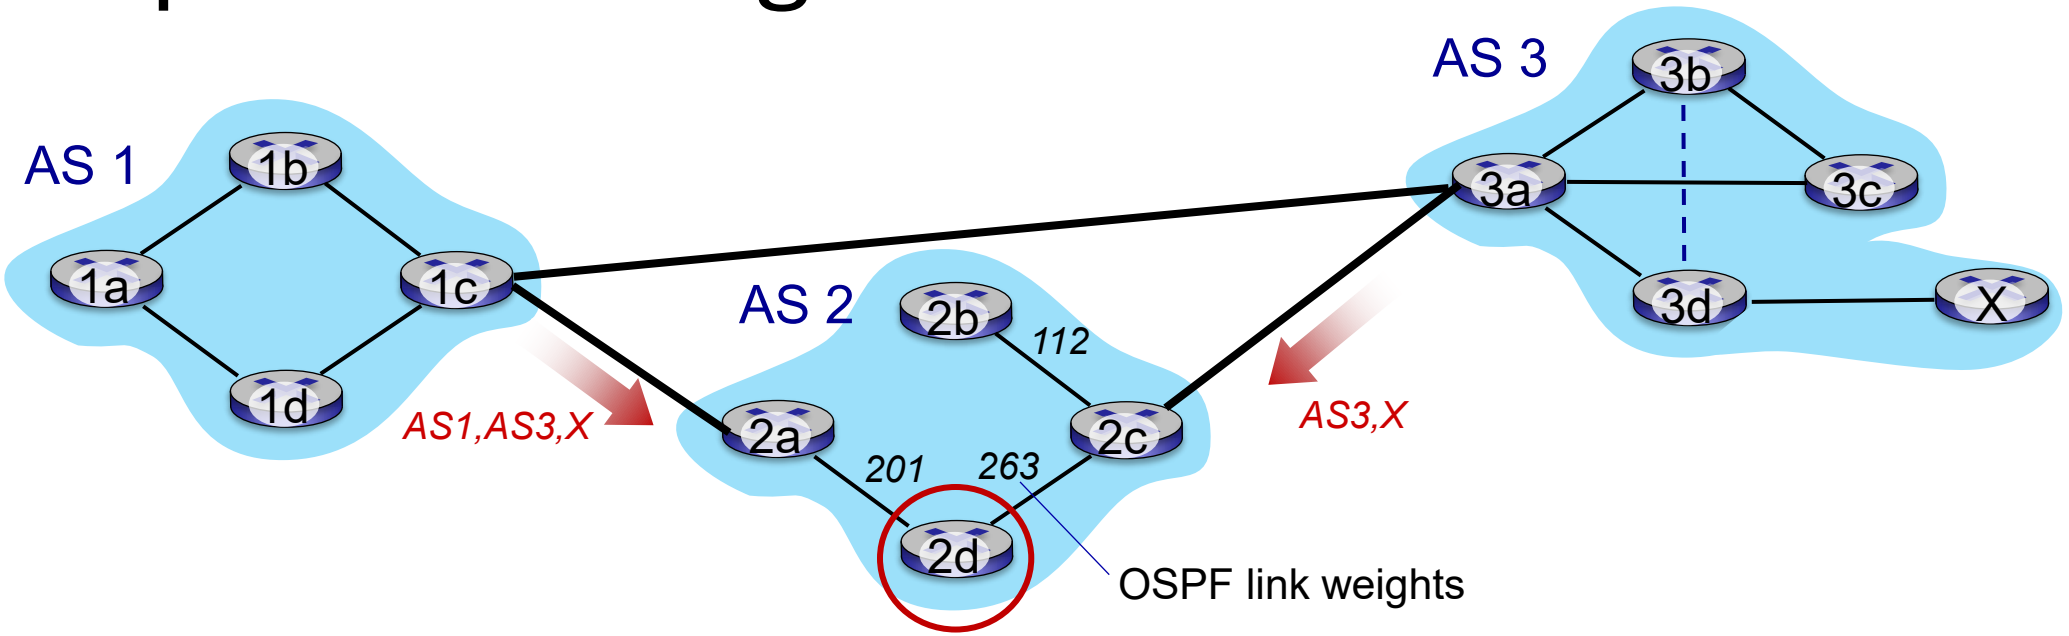

- 2d learns (via iBGP) it can route to X via 2a or 2c
- hot potato routing: choose local gateway that has least *intra-domain* cost (e.g., 2d chooses 2a, even though more AS hops to *X*): don't worry about inter-domain cost!

BGP: achieving policy via advertisements

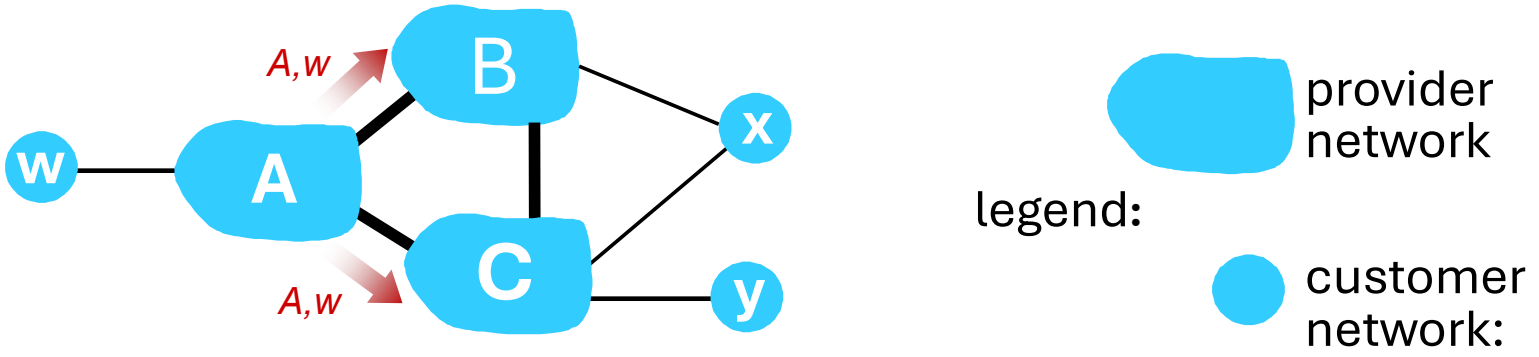

ISP only wants to route traffic to/from its customer networks (does not want to carry transit traffic between other ISPs – a typical "real world" policy)

- A advertises path Aw to B and to C
- B *chooses not to advertise* BAw to C!
  - B gets no "revenue" for routing CBAw, since none of C, A, w are B's customers
  - C does *not* learn about CBAw path
- C will route CAw (not using B) to get to w

BGP: achieving policy via advertisements (more)

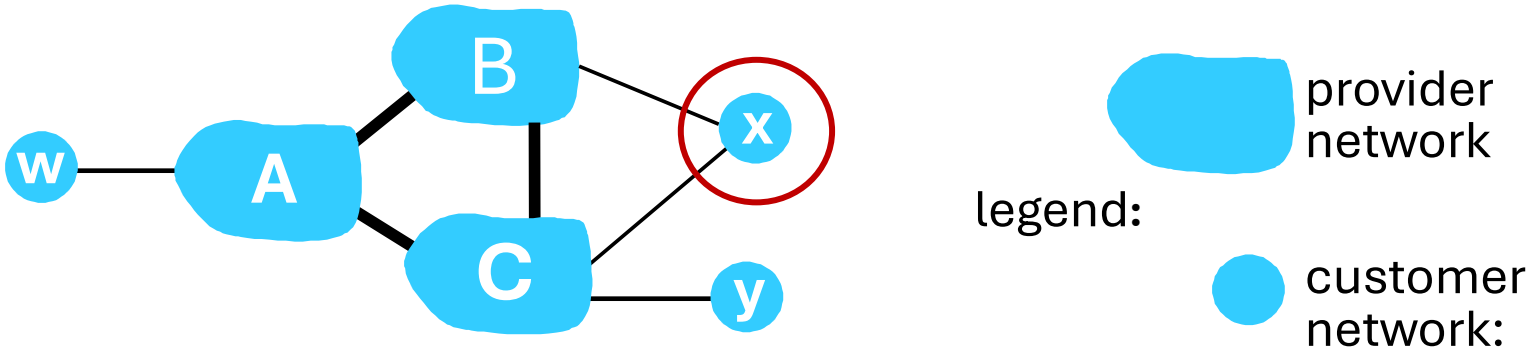

ISP only wants to route traffic to/from its customer networks (does not want to carry transit traffic between other ISPs – a typical "real world" policy)

- A,B,C are provider networks
- x,w,y are customer (of provider networks)
- x is dual-homed: attached to two networks
- *policy to enforce:* x does not want to route from B to C via x
  - .. so x will not advertise to B a route to C

BGP route selection

- router may learn about more than one route to destination AS, selects route based on:
  1. local preference value attribute: policy decision
  2. shortest AS-PATH
  3. closest NEXT-HOP router: hot potato routing
  4. additional criteria